Search Smooths Discontinuities in Discrete/Continuous Problems

Christopher D. Carroll[*] JHU
Fedor Iskhakov[†] Stony Brook University
Shujaat Khan[‡] IMF
Adrian Monninger[§] Texas A&M

October 15, 2026

Abstract

Discrete-continuous problems solved by the endogenous gridpoint method produce kinks in the value function and jumps in the decision rule. Those jumps make the Euler equation necessary but not sufficient, so the method returns suboptimal points. Extreme Value Type I taste shocks remove the jumps [**?**]. This paper shows that directed search together with a rejectable price draw can match that smoothness, using objects that are part of the agent's problem. The agent chooses effort in the unit interval at a convex utility cost, and effort raises the probability of receiving an option to switch. Conditional on receiving the option, the agent draws a price and may reject it. A retirement model illustrates the mechanism. In a keep-versus-adjust model of durables, taste shocks, price dispersion, and search all smooth the (S, s) band, and their steady-state effects on the adjustment probability have opposite signs. Those signs identify the search cost and the price dispersion separately. Matching impulse responses of consumption and of the adjustment probability remains open.

Keywords: Consumption, Computational Economics, Discrete-Continuous Choice, Endogenous Gridpoint Method, Search, Price Dispersion
JEL codes: C13, C63

Repo: https://github.com/AMonninger/DCEGM_Search
PDF: https://github.com/AMonninger/DCEGM_Search/Text/dcegmSearch.pdf

The computational results in this paper were constructed using tools in the Econ-ARK/HARK toolkit. The toolkit can be cited by its digital object identifier 10.5281/zenodo.1001068, as is done in the paper's own references as **?**. Thanks to Matthew White for initial discussions of this subject.

[*]Carroll: Department of Economics, Johns Hopkins University, http://econ.jhu.edu/people/ccarroll/, ccarroll@jhu.edu.
[†]Iskhakov: Stony Brook University.
[‡]Khan: IMF.
[§]Monninger: Department of Economics, Texas A&M University, https://sites.google.com/view/adrianmonninger, amonninger@tamu.edu.

1 Introduction

This paper studies dynamic problems with both a discrete and a continuous choice. Applied to these problems, the endogenous gridpoint method of **?** returns suboptimal points as well as the optimum. Extreme Value Type I taste shocks remove the jumps that produce those points [**?**]. Directed search, together with a rejectable price draw, can match that smoothness. Both objects are part of the agent's problem. A retirement model illustrates the mechanism. A keep-versus-adjust model of durables is where the three devices can be told apart: they move the steady-state adjustment probability in opposite directions.

For dynamic stochastic optimization problems with continuous choice variables, the Endogenous Gridpoint Method (EGM) is significantly more efficient compared to standard root-finding methods in finding the optimal decision rules. In these problems, the value functions are typically concave and the Euler equation is a necessary and sufficient condition for the optimal decision rule. In contrast, when the choice variables include both continuous and discrete choices (DC models), the Euler equation has multiple solutions for the continuous choice. Therefore, the Euler equation is only a necessary, but not sufficient, condition for the optimal decision rule.

The complication in solving a dynamic optimization problem arises due to the kinks in the value function and discontinuities in the decision rule at the points in the state space where discrete choices are made. These discontinuities make the decision rule non-monotonic, which leads to the Euler equation producing suboptimal points. In addition, each period's kinks and discontinuities are propagated back in time as the solution is solved by backward iteration, thus, exacerbating the problem. **?**, **?**, **?**, and **?** provide methods for finding and discarding these suboptimal points from the final solution.

In order to mitigate the complications arising from the accumulation of the discontinuities, **?** introduce Extreme Value Type I taste shocks that affect the likelihood of the discrete choice outcomes. These taste shocks have the advantage of smoothing out the expected value function and the expected marginal utility function. With larger taste shocks, the value function can be 'concavified' to a greater degree. Such taste shocks, however, are in fact 'behavioral' shocks that add a degree of randomization to the discrete choice outcomes.

The agent chooses search effort in the unit interval. The cost of effort is convex and is subtracted from utility. Effort raises the probability of receiving an option: the chance to retire, or a match at which a purchase price is drawn. Optimal effort is continuous in the value gap between the two alternatives, so the combined value is smooth where that gap is positive. Where the gap is not positive, optimal effort is zero, and search alone leaves that side of the threshold unsmoothed. Conditional on receiving the option, the agent draws a realization and may reject it. In the retirement model the draw is an account factor ϕ, and retirement wealth is ϕ times market resources. In the durable model the draw is the purchase price of the new stock. Rejecting a bad draw softens the extensive margin on the side of the threshold where effort is zero. The option value of the draw raises effort. Together, the two objects can match the smoothness of the Extreme Value shock.

The common object in the two applications is effort in $[0, 1]$, a convex utility cost, and a success probability equal to effort. In both, the household knows the distribution of the draw, chooses effort before the realization, and may reject the draw if search succeeds. The remaining difference is the date of the benefit. Retirement effort chosen at the end of period t is the probability that an option arrives at $t+1$, so β enters the effort rule and effort is irrelevant at the terminal date. Durable effort is the probability of a match in the same period. The match draws a purchase price, and the household accepts or keeps the current stock. The probability of buying is effort times the probability that the draw is acceptable. Both rules are continuous in a value gap and flat where the gap is not positive, so either date delivers that smoothness. The probabilities remain different objects. The probability

of retiring is next period's arrival times the probability that claiming is selected. The adjustment probability in Section 3 is this period's effort times the acceptance probability, and the steady-state comparison disciplines that object. A search duration is the expected wait until an option arrives in the retirement model, and the probability of a match this period in the durable model.

Section 2 illustrates the mechanism in a consumption-saving model with a work-versus-retire choice, in the style of **?**. The comparison is Extreme Value Type I taste shocks against search combined with the account factor. The section shows how those objects smooth the consumption function. It is not the paper's quantitative result.

Section 3 applies the same pair to a keep-versus-adjust model of durables in the (S, s) tradition of **?** and **?**. The durable stock is continuous, so the problem is a genuine discrete-continuous choice. Extreme Value taste shocks, price dispersion, and search all smooth the (S, s) band. In steady state their effects on the adjustment probability have opposite signs, and that opposition is the result. A larger taste scale raises the adjustment probability, and durables and assets fall. Wider price dispersion also raises the adjustment probability, because better draws are accepted; durables rise and assets fall. A higher search cost lowers the adjustment probability; durables fall and assets rise. Search and price dispersion can therefore be disciplined separately, and the model can be asked to match impulse responses of consumption and of the adjustment probability. That matching exercise is in progress. Section 4 records what a fuller quantitative model still requires.

Search has three advantages over an Extreme Value taste shock. First, it acts near the (S, s) thresholds, so the Euler equation still describes agents far from indifference. Taste shocks put probability on the dominated branch at every state. **?** find that three quarters of adjusters in an estimated model move because of match-quality shocks. Search can be asked to match the same adjustment share while leaving agents far from the band on their Euler equation. Second, search intensity maps into an expected duration, which is observable. Where that intensity is the probability an option arrives next period, a value of one half is an expected wait of two periods. The scale of an Extreme Value shock has no such reading. It is typically chosen to match the adjustment share of agents far from the band, as in **?**, or to match intertemporal marginal propensities to consume, as in **?**. Third, a model with permanent and transitory income can be normalized by permanent income [**?**]. The search cost in this paper is a utility cost. Specified so that it scales with permanent income, it adds no state to the normalized problem. An Extreme Value shock of the same scale for every household does not scale in that way, and permanent income remains a state variable.

The rest of the paper proceeds as follows. Section 2 presents the retirement illustration. Section 3 compares taste shocks, search, and price dispersion in the basic durable model. Section 4 states what the quantitative model still requires. Section 5 concludes.

2 Retirement Model

In this section, we build a consumption-saving model with an option to retire. We first state the problem with a hard maximum over work and retirement, and no device that smooths that margin. We then summarize the Extreme Value Type I taste shocks of **?**, which are the usual smoother. We then replace those shocks with directed search and an endogenous price of retirement wealth, and derive that problem in detail. Labor income while at work is a constant y. The illustration abstracts from income risk.

2.1 The problem without smoothing

Agents derive utility from consuming non-durable goods c_t and get disutility from working. The disutility parameter is denoted by δ. Denoting the choice to retire and the choice to work by $d_t = 0$ and $d_t = 1$, respectively, and assuming logarithmic utility from consumption, the utility function is

$$u(c_t, d_t) = \log(c_t) - \delta d_t. \tag{1}$$

At the end of each period, households save their end-of-period liquid assets A_t at a risk-free rate R. Working in the current period determines whether the agent receives income in the following period, so beginning-of-period resources are $M_t = RA_{t-1} + y\, d_{t-1}$.

Retirement is self-absorbing, so the retiree's problem is the ordinary consumption-saving problem. In Bellman form,

$$V_t^{\text{Ret}}(M_t) = \max_{0 \le c_t \le M_t} \left\{ \log(c_t) + \beta V_{t+1}^{\text{Ret}}(M_{t+1}) \right\}, \tag{2}$$

where $M_{t+1} = A_t R$ and $A_t = M_t - c_t$. Using the gothic script V ($\mathcal{V}$) for the value of ending the period with assets A_t,

$$\mathcal{V}_t^{\text{Ret}}(A_t) \equiv V_{t+1}^{\text{Ret}}(M_{t+1}(A_t)) = V_{t+1}^{\text{Ret}}(A_t R), \tag{3}$$

and the retiree's problem in terms of end-of-period assets is

$$\mathcal{V}_{t-1}^{\text{Ret}}(A_{t-1}) = \max_{0 \le c_t \le A_{t-1} R} \left\{ \log(c_t) + \beta \mathcal{V}_t^{\text{Ret}}(A_t) \right\}. \tag{4}$$

A worker who has not yet retired has the option to retire this period. The discrete choice is a hard maximum of the two conditional values, both evaluated at book resources M_t,

$$V_t(M_t) = \max \left\{ V_t^{\text{Ret}}(M_t);\ V_t^{W}(M_t) \right\}. \tag{5}$$

The worker who continues solves

$$V_t^{W}(M_t) = \max_{0 \le c_t \le M_t} \left\{ \log(c_t) - \delta + \beta\, V_{t+1}(M_{t+1}) \right\}, \tag{6}$$

with $A_t = M_t - c_t$ and $M_{t+1} = RA_t + y$. Next period's value is again the hard maximum, so a kink at the wealth level where the household is indifferent between retiring and working is inherited by the continuation value. The worker's consumption function therefore jumps at each such threshold, and those jumps propagate backward. That sawtooth is the object the rest of this section smooths.

2.2 Extreme Value Type I taste shocks

The usual smoother puts i.i.d. Extreme Value Type I shocks on the two discrete alternatives. After observing M_t, the household draws $(\zeta_t^{\text{Ret}}, \zeta_t^{W})$ and chooses

$$V_t(M_t) = \max \left\{ V_t^{\text{Ret}}(M_t) + \sigma_\zeta \zeta_t^{\text{Ret}},\ V_t^{W}(M_t) + \sigma_\zeta \zeta_t^{W} \right\}, \tag{7}$$

where $\sigma_\zeta > 0$ is the taste-shock scale. The shocks do not enter the continuous consumption choice inside either branch. Because the draws are i.i.d. Extreme Value Type I, the expected value has the closed form

$$V_t(M_t) = \sigma_\zeta \log \left(\exp\left(V_t^{\text{Ret}}(M_t)/\sigma_\zeta\right) + \exp\left(V_t^{W}(M_t)/\sigma_\zeta\right) \right), \tag{8}$$

and the probability of retiring is the logit weight on V_t^{Ret}. As $\sigma_\zeta \to 0$ this collapses to the hard maximum above. A larger scale puts weight on the inferior branch near indifference, so the worker's consumption function is smooth in M_t. The results below show this case with search and account dispersion shut down, so the household always faces the retirement choice.

2.3 Search and an endogenous account price

We now replace the taste shocks with two mechanisms that enter together. Each period there is a possibility of receiving an option to retire. Search effort ε_{t-1} exerted in the previous period is the probability of receiving that option. Zero effort means the option does not arrive. Search is costly, with

$$\theta(\varepsilon_t) = \sigma_s \left(\varepsilon_t + (1-\varepsilon_t)\log(1-\varepsilon_t)\right), \tag{9}$$

so the first units of effort are cheap and further effort is increasingly costly. Utility is then

$$u(c_t, d_t, \varepsilon_t) = \log(c_t) - \delta d_t - \theta(\varepsilon_t). \tag{10}$$

We write ϕ_t for the account realization. Contingent on receiving the option, the household observes ϕ_t and decides whether to take the draw—retire, absorbing, with that wealth—or to deny it and work another period. The realization is the endogenous price of the nest egg at the moment of claiming. It follows a mean-one log-normal distribution,

$$\log \phi_t \sim \mathcal{N}\Big(-\frac{\sigma_\phi^2}{2}, \sigma_\phi^2\Big), \qquad \mathbb{E}[\phi_t] = 1, \tag{11}$$

independent across periods. If the household claims, retirement begins with wealth

$$\tilde{M}_t = \phi_t M_t, \tag{12}$$

and from that point the agent is an ordinary retiree. If the household denies the draw, it remains a worker with book resources M_t and the standard budget $A_t = M_t - c_t$; the realization ϕ_t is discarded. The draw is a one-shot, rejectable mark-to-market of the nest egg—a stochastic liquidation value, annuity conversion factor, or defined-contribution account realization—and not an ongoing price of consumption in retirement. Once the draw is accepted, ϕ_t is locked into the level of wealth and is not redrawn. A high realization of ϕ_t raises the value of retiring and, conditional on claiming, the resources out of which the retiree consumes. A low realization has the opposite effect and can be rejected.

The retiree's problem is the one stated above. The value of claiming is that function evaluated at scaled wealth,

$$V_t^{\mathrm{Ret}}(\tilde{M}_t) = V_t^{\mathrm{Ret}}(\phi_t M_t), \tag{13}$$

with claiming-period consumption $c_t^{\mathrm{Ret}}(\phi_t M_t)$ read off the same policy. The retiree's decision rule is unchanged. Only the wealth at which it is evaluated depends on ϕ_t. There is no separate price-contingent retiree problem, and no need to re-solve the retiree at each quadrature node of ϕ: the intensive margin is movement along a single consumption function.

2.4 Worker's sub-problem

Each period, a worker either receives the option to retire or has to continue working. The search effort is made at the end of a period and is the probability of receiving the option to retire in the following period. When period t starts, beginning-of-period book resources M_t are known. The household then learns whether the option arrived, which was determined by last period's effort. If it did not, the agent must work, and ϕ_t is irrelevant, because there is no account to claim. If it did, the agent observes ϕ_t and decides whether to claim and retire or to deny the draw and continue working. The discrete comparison is a hard maximum, contingent on the realized account price.

Having made the discrete choice, the agent makes a consumption decision. If the agent chooses to retire, they stay as a retiree, because retirement is self-absorbing. If the agent chooses to continue working, they exert search effort, which is the probability of receiving a new draw in the next period. Fig. 1 outlines the timeline of the realization of the shocks and the decisions in the worker's problem.

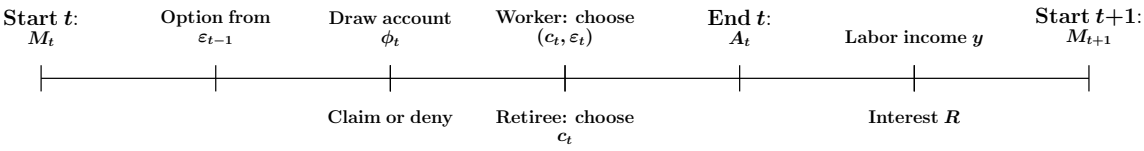


Figure 1: Timeline of the worker's problem

Let the set of possible choice sets for the worker's discrete choice be given by $\mathcal{D}_t \in \{\{\text{Retire}, \text{Work}\}, \{\text{Work}\}\} \equiv \{\{d_t = 0, d_t = 1\}, \{d_t = 1\}\}$. The first element of this set is the case when the agent has the option to switch and the second element is the case when the agent does not have the option to switch and must continue working. Write $P_1(\varepsilon_t) = \varepsilon_t$ for the probability of receiving the option next period and $P_2(\varepsilon_t) = 1 - \varepsilon_t$ for the probability of not receiving it. For the case when $\mathcal{D}_t = \{W\}$, i.e. the worker does not have the option to switch, the Bellman equation is given by:

$$V_t^W(M_t) = \max_{\substack{0 \le c_t \le M_t \\ 0 \le \varepsilon_t \le 1}} \left\{ u(c_t, 1, \varepsilon_t) + \beta P_1(\varepsilon_t) \mathcal{V}_t^1(A_t) + \beta P_2(\varepsilon_t) \mathcal{V}_t^2(A_t) \right\}, \tag{14}$$

where $A_t = M_t - c_t$, $\mathcal{V}_t^1(A_t)$ is the value of saving A_t when having the option to retire in $t+1$, and $\mathcal{V}_t^2(A_t)$ is the value of saving A_t when not having the option to retire in $t+1$. Current ϕ_t does not enter (14). A worker who did not claim does not cash the account this period.

For the case when $\mathcal{D}_t = \{R, W\}$, i.e. the case when the worker has the option to switch and has observed the account price, the Bellman equation is given by:

$$V_t(M_t, \{R, W\}, \phi_t) = \max \left\{ V_t^{\text{Ret}}(\phi_t M_t);\ V_t^W(M_t) \right\}, \tag{15}$$

or

$$\begin{aligned} V_t(M_t, \{R, W\}, \phi_t) = \max \Bigg[& \max_{0 \le c_t \le \phi_t M_t} \left\{ u(c_t, 0, 0) + \beta \mathcal{V}_t^{\text{Ret}}(A_t) \right\}; \\ & \max_{\substack{0 \le c_t \le M_t \\ 0 \le \varepsilon_t \le 1}} \left\{ u(c_t, 1, \varepsilon_t) + \beta P_1(\varepsilon_t) \mathcal{V}_t^1(A_t) + \beta P_2(\varepsilon_t) \mathcal{V}_t^2(A_t) \right\} \Bigg]. \end{aligned} \tag{16}$$

A high ϕ_t makes claiming more attractive and, if accepted, delivers higher retirement consumption along $c_t^{\text{Ret}}(\cdot)$. A low ϕ_t makes claiming less attractive and is denied whenever $V_t^{\text{Ret}}(\phi_t M_t) < V_t^W(M_t)$.

The value of saving A_t when having the option to retire in $t+1$, i.e. $\mathcal{V}_t^1(A_t)$, integrates over next period's account draw. Next period's labor income is the constant y, so book resources $M_{t+1} = A_t R + y$ are certain:

$$\mathcal{V}_t^1(A_t) = \int_{\phi_{t+1}} V_{t+1}\big(M_{t+1}(A_t), \{R, W\}, \phi_{t+1}\big)\, dF(\phi_{t+1}), \tag{17}$$

where $V_{t+1}(M_{t+1}, \{R, W\}, \phi_{t+1})$ is the hard maximum (16) one period ahead. The value of saving A_t when not having the option to retire in $t+1$, i.e. $\mathcal{V}_t^2(A_t)$, is

$$\mathcal{V}_t^2(A_t) = V_{t+1}^W\big(M_{t+1}(A_t)\big). \tag{18}$$

This continuation does not integrate over ϕ_{t+1}, because a worker who must work cannot claim. Future draws still sit inside V^W through $\mathcal{V}^1$ and $\mathcal{V}^2$ one period ahead. Equivalently, the beginning-of-period value of a worker who will have the option, after integrating out the account price, is

$$\bar{V}_t^{R,W}(M_t) \equiv \mathbb{E}_{\phi_t}\Big[\max\left\{V_t^{\text{Ret}}(\phi_t M_t),\ V_t^W(M_t)\right\}\Big], \tag{19}$$

so that $\mathcal{V}_{t-1}^1(A_{t-1}) = \bar{V}_t^{R,W}(RA_{t-1}+y)$. Integrating (15) over ϕ_t replaces a single wealth threshold with a family of thresholds, one for each realization, so the ex-ante claiming probability is interior.

2.5 Optimal effort and consumption functions

Since we assume that the consumption decision is made before the effort decision, (14) can be split into two sequential sub-period problems and rewritten as:

$$V_t^W(M_t) = \max_{0\le c_t\le M_t}\left[\mathcal{V}_t(A_t)\right] \tag{20}$$

where

$$\mathcal{V}_t(A_t) = \max_{0\le\varepsilon_t\le 1}\left\{u(c_t,1,\varepsilon_t) + \beta P_1(\varepsilon_t)\mathcal{V}_t^1(A_t) + \beta P_2(\varepsilon_t)\mathcal{V}_t^2(A_t)\right\}. \tag{21}$$

Search effort is chosen only on the working branch and does not depend on the current draw. The current ϕ_t is already sunk from the point of view of next period's option. Because the arrival probability equals effort, $\partial P_1/\partial\varepsilon_t = 1$ and $\partial P_2/\partial\varepsilon_t = -1$. The first order condition gives the optimal effort as a function of the end-of-period assets that are determined after the optimal consumption decision has been made in the previous sub-period:

$$\frac{\partial\theta(\varepsilon_t)}{\partial\varepsilon_t} = \beta\left(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\right) \tag{22}$$

$$-\sigma_s\log(1-\varepsilon_t^*) = \beta\left(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\right) \tag{23}$$

$$\varepsilon_t^* = \max\left\{0,\ 1-\exp\left(-\frac{\beta}{\sigma_s}\left(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\right)\right)\right\}. \tag{24}$$

Optimal effort is zero when receiving the option is not strictly better than not receiving it, and it is increasing in that excess value when the gap is positive. Account-value dispersion raises $\mathcal{V}_t^1$ relative to $\mathcal{V}_t^2$, because the option to retire is also an option on ϕ_{t+1}: the household claims when the draw is attractive and otherwise continues working. That option value increases the excess value of receiving the option, and therefore increases search effort. The two frictions are complements. Search governs whether a draw arrives; ϕ_t is the content of the draw, and the hard maximum (16) is the decision to take it or to work another period. Together they smooth both the arrival of the option and the terms on which it is exercised.

As $\sigma_\phi \to 0$, $\phi_t \to 1$ almost surely and (14) is the search problem with a hard maximum at book resources. Shutting search down as well, so that the retirement choice is simply available, recovers the single wealth threshold of the unsmoothed problem.

Next, we derive the expressions for the optimal consumption functions for the retiree and the worker. Since retirement is self-absorbing, and ϕ_t is locked in at the moment of claiming, the retiree's consumption function is the standard one, evaluated at claimed wealth:

$$u'(c_t,0,0) = R\beta u'(c_{t+1},0,0) \tag{25}$$

or

$$c_t^{\text{Ret}}(\phi_t M_t) = \left(R\beta\left(c_{t+1}^{\text{Ret}}(A_t R)\right)^{-1}\right)^{-1}, \tag{26}$$

where $A_t = \phi_t M_t - c_t^{\text{Ret}}(\phi_t M_t)$. To derive the consumption function for the worker, we start with the solution to the problem of a worker who does not have the option to retire, i.e. $\mathcal{D}_t = \{W\}$. The first-order condition for (14), with respect to c_t, is given by

$$u'(c_t, 1, \varepsilon_t^*) = \beta \left[P_1(\varepsilon_t^*) \frac{\partial \mathcal{V}_t^1(A_t)}{\partial A_t} + P_2(\varepsilon_t^*) \frac{\partial \mathcal{V}_t^2(A_t)}{\partial A_t} \right], \tag{27}$$

where ε_t^* is the optimal effort function in (24), and the derivatives of the end-of-period values are given by

$$\frac{\partial \mathcal{V}_t^1(A_t)}{\partial A_t} = R\, \mathbb{E}_{\phi_{t+1}} \left\{ \frac{\partial V_{t+1}(M_{t+1}, \{R, W\}, \phi_{t+1})}{\partial M_{t+1}} \right\} \tag{28}$$

and

$$\frac{\partial \mathcal{V}_t^2(A_t)}{\partial A_t} = R\, \frac{\partial V_{t+1}^W(M_{t+1})}{\partial M_{t+1}}, \tag{29}$$

where the expectation in the first line is over next period's account price and $M_{t+1} = A_t R + y$. Substituting these expressions yields

$$u'(c_t, 1, \varepsilon_t^*) = R\beta \left[P_1(\varepsilon_t^*) \mathbb{E}_{\phi_{t+1}} \left\{ \frac{\partial V_{t+1}(M_{t+1}, \{R, W\}, \phi_{t+1})}{\partial M_{t+1}} \right\} + P_2(\varepsilon_t^*) \frac{\partial V_{t+1}^W(M_{t+1})}{\partial M_{t+1}} \right]. \tag{30}$$

The envelope condition on the working branch is

$$\frac{\partial V_t^W(M_t)}{\partial M_t} = u'(c_t, 1, \varepsilon_t^*), \tag{31}$$

which can be iterated forward to yield

$$\frac{\partial V_{t+1}^W(M_{t+1})}{\partial M_{t+1}} = u'(c_{t+1}, 1, \varepsilon_{t+1}^*). \tag{32}$$

For the worker who has the option to retire, the value function is the maximum of the value of claiming and the value of continuing to work,

$$V_t(M_t, \{R, W\}, \phi_t) = \max \left\{ V_t^{\text{Ret}}(\phi_t M_t);\ V_t^W(M_t) \right\}. \tag{33}$$

An extra unit of book wealth is worth ϕ_t units of retirement wealth if the household claims. Therefore,

$$\frac{\partial V_{t+1}(M_{t+1}, \{R, W\}, \phi_{t+1})}{\partial M_{t+1}} = \mathbb{1}_{t+1}^0 \cdot \phi_{t+1} \cdot u'(c_{t+1}, 0, 0) + \mathbb{1}_{t+1}^1 \cdot u'(c_{t+1}, 1, \varepsilon_{t+1}^*), \tag{34}$$

where the retiree marginal utility is evaluated at claimed wealth $\phi_{t+1} M_{t+1}$, $\mathbb{1}_{t+1}^0 \equiv \mathbb{1}_{t+1}^0(M_{t+1}, \phi_{t+1}) = 1$ if $V_{t+1}^{\text{Ret}}(\phi_{t+1} M_{t+1}) \geq V_{t+1}^W(M_{t+1})$ and 0 otherwise, and $\mathbb{1}_{t+1}^1 = 1 - \mathbb{1}_{t+1}^0$. Substituting the two envelopes into the first-order condition yields the consumption function for the worker who continues to work

$$\begin{aligned} u'(c_t, 1, \varepsilon_t^*) = R\beta \Big[& P_1(\varepsilon_t^*) \int_{\phi_{t+1}} \left[\mathbb{1}_{t+1}^0 \cdot \phi_{t+1} \cdot u'\big(c_{t+1}^{\text{Ret}}(\phi_{t+1} M_{t+1}), 0, 0\big) + \mathbb{1}_{t+1}^1 \cdot u'\big(c_{t+1}^W(M_{t+1}), 1, \varepsilon_{t+1}^*\big) \right] dF(\phi_{t+1}) \\ & + P_2(\varepsilon_t^*)\, u'\big(c_{t+1}^W(M_{t+1}), 1, \varepsilon_{t+1}^*\big) \Big], \end{aligned} \tag{35}$$

or, with logarithmic utility,

$$c_t^W = \left(R\beta \Big[P_1(\varepsilon_t^*) \int_{\phi_{t+1}} \Big[\mathbb{1}_{t+1}^0 \, \phi_{t+1} \left(c_{t+1}^{\mathrm{Ret}}(\phi_{t+1} M_{t+1}) \right)^{-1} + \mathbb{1}_{t+1}^1 \left(c_{t+1}^W(M_{t+1}) \right)^{-1} \Big] dF(\phi_{t+1}) \right.$$
$$\left. + P_2(\varepsilon_t^*) \left(c_{t+1}^W(M_{t+1}) \right)^{-1} \Big] \right)^{-1}, \tag{36}$$

where $M_{t+1} = A_t R + y$. A worker who has received the option observes ϕ_t before the claim-or-work choice, so the realized policy is the hard maximum at that draw. The ex-ante consumption function is the ϕ-weighted average of those choices. Working-branch consumption c_t^W and effort ε_t^* do not depend on the current draw. The envelope conditions and the EGM algorithm for this baseline are derived in Appendix A.3.

2.6 Results

Having established the solution of the retirement problem, we show how the consumption rule of the worker who continues to work responds to Extreme Value Type I taste shocks, to search, and to dispersion in the retirement account. Every figure is drawn at period $t = 15$ of a $T = 20$ problem, with $R = 1$, $\beta = 0.98$, and $y = 20$. The object plotted is the worker's consumption function $C(m_t)$.

With search and account dispersion shut down, so that the household always faces the retirement choice, Fig. 2 varies the Extreme Value scale σ_ζ. At $\sigma_\zeta = 0$ the hard maximum leaves a sequence of sharp discontinuities. Raising the scale to 0.05 rounds the peaks. At 0.10 and 0.15 the kinks are essentially gone. Larger taste shocks also shift consumption down through the middle of the resource range and up at high resources.

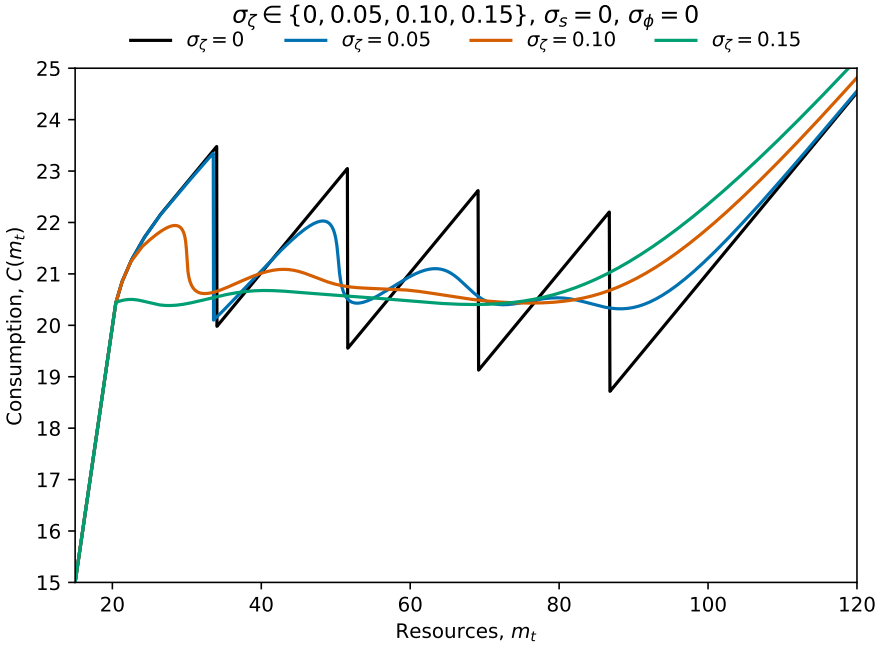


Figure 2: Worker consumption at $t = 15$ under Extreme Value Type I taste shocks, $\sigma_\zeta \in \{0, 0.05, 0.10, 0.15\}$. Search and account dispersion are shut down ($\sigma_s = 0$, $\sigma_\phi = 0$), so the household always faces the retirement choice.

Fig. 3 shuts the taste shocks and the account price down. The option arrives only through search. At $\sigma_s = 0$ the discontinuities remain. A small search cost, $\sigma_s = 0.15$, trims them from the right. At $\sigma_s = 0.50$ and $\sigma_s = 1$ the rule is much smoother, and the smoothing sits to the right of each kink: workers with low liquid resources do not search, so the left side of each discontinuity is left intact. Higher search costs also shift consumption up at high resources.

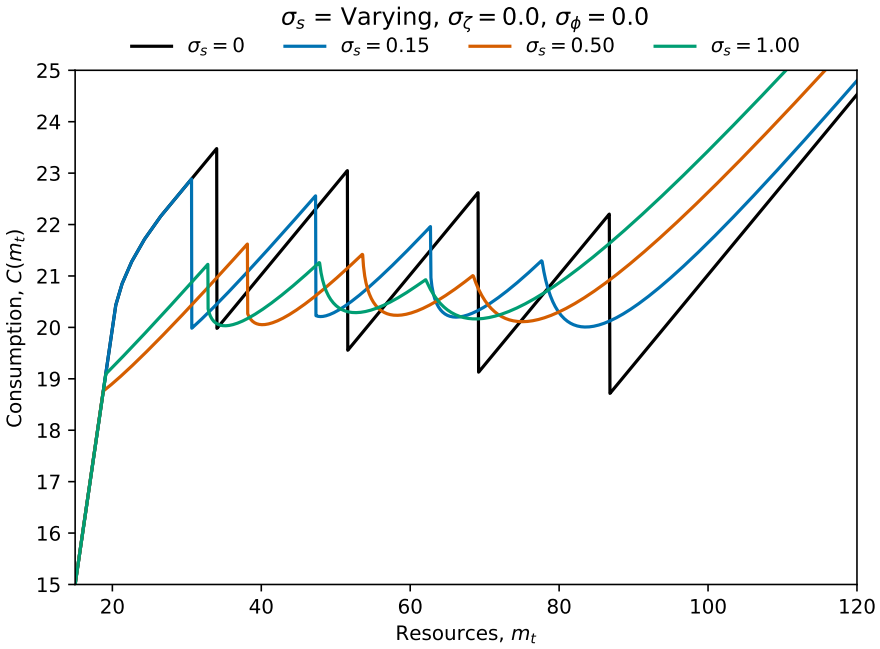


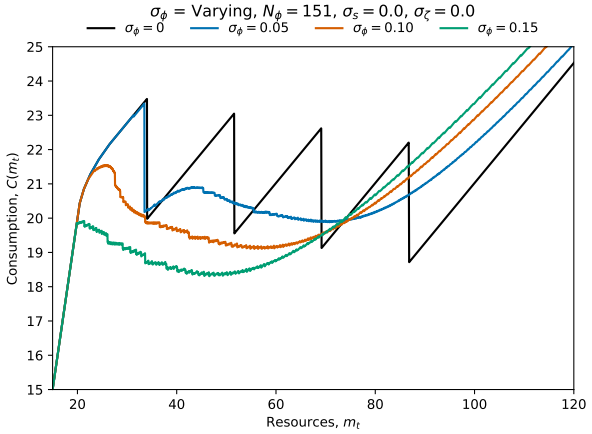
Figure 3: Worker consumption at $t = 15$ under search alone. Taste shocks and account dispersion are shut down ($\sigma_\zeta = 0$, $\sigma_\phi = 0$). The option arrives only through search.

Account dispersion on its own is shown in Fig. 4. Search is shut down, so a draw of ϕ_t arrives with certainty. In panel (a), with $N_\phi = 151$ quadrature nodes, $\sigma_\phi = 0$ reproduces the sawtooth. A small dispersion, $\sigma_\phi = 0.05$, removes every kink but the first. At $\sigma_\phi = 0.10$ and 0.15 the rule is smooth, and it lies below the hard-maximum rule through the middle of the grid. Panel (b) holds $\sigma_\phi = 0.1$ and varies the number of nodes. A single node is the hard maximum. Nine nodes follow the smooth rule but leave a sawtooth of their own. At $N_\phi = 151$ and $N_\phi = 351$ the two curves lie on top of each other, so further nodes do not change the consumption function.

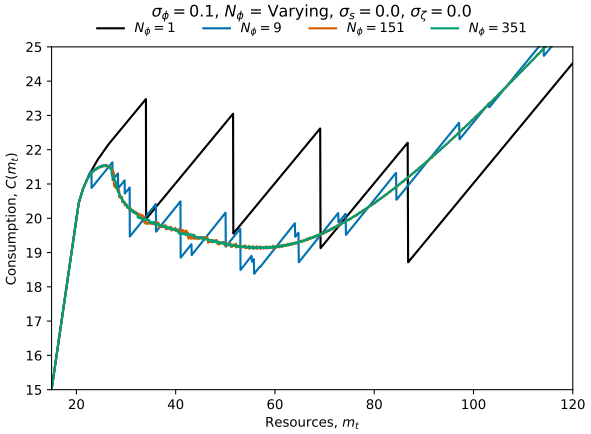
Search and the account price together are shown in Fig. 5, with no taste shocks. The option arrives only through search. In panel (a), $\sigma_\phi = 0.1$ and $N_\phi = 151$ are held fixed while σ_s rises from 0 to 1. The price draw has already removed most of the sawtooth; raising the search cost removes the remaining bump at low resources and shifts the rule up. In panel (b), $\sigma_s = 0.5$ and $N_\phi = 151$ are held fixed while σ_ϕ rises. At $\sigma_\phi = 0$ the one-sided kinks of search alone are still visible. They shrink at $\sigma_\phi = 0.05$ and are gone at 0.10 and 0.15. Used together, the two mechanisms produce a smooth consumption function at scales at which either one, used alone, still leaves kinks.

2.7 Direct comparison

The four figures above vary one device at a time. Fig. 6 puts an Extreme Value consumption function and a search-and-account consumption function on the same axes. The target is the Extreme Value rule at $\sigma_\zeta = 0.10$. The search cost and the account dispersion are chosen to minimize the root mean

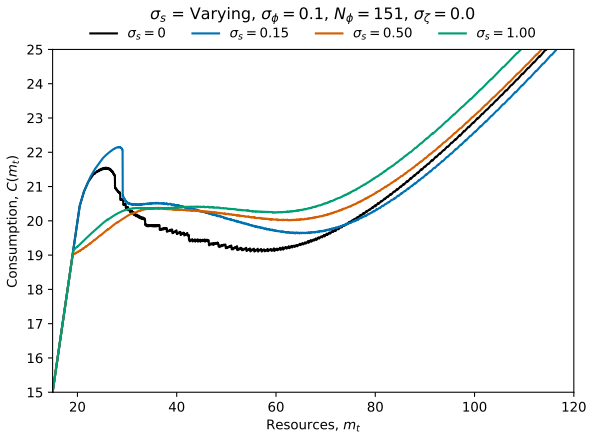


(a) Varying account dispersion, $N_\phi = 151$.

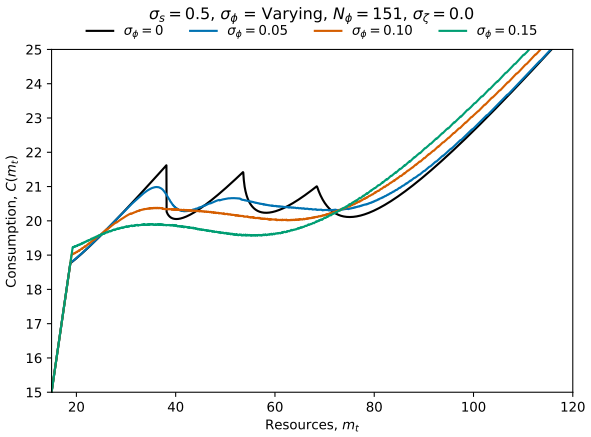


(b) Varying the number of account-price nodes, $\sigma_\phi = 0.1$.

Figure 4: Worker consumption at $t = 15$ with account dispersion alone. Search and taste shocks are shut down ($\sigma_s = 0$, $\sigma_\zeta = 0$), so the account draw arrives with certainty.



(a) Varying search, $\sigma_\phi = 0.1$, $N_\phi = 151$.



(b) Varying account dispersion, $\sigma_s = 0.5$, $N_\phi = 151$.

Figure 5: Worker consumption at $t = 15$ with search and account dispersion together. Taste shocks are shut down ($\sigma_\zeta = 0$). The option arrives only through search.

squared gap between the two rules on resources from 15 to 120. The minimum is at $\sigma_s = 0.171$ and $\sigma_\phi = 0.046$. The remaining gap is 0.23 consumption units, so the search-and-account rule tracks the Extreme Value rule and stays a short distance from it.

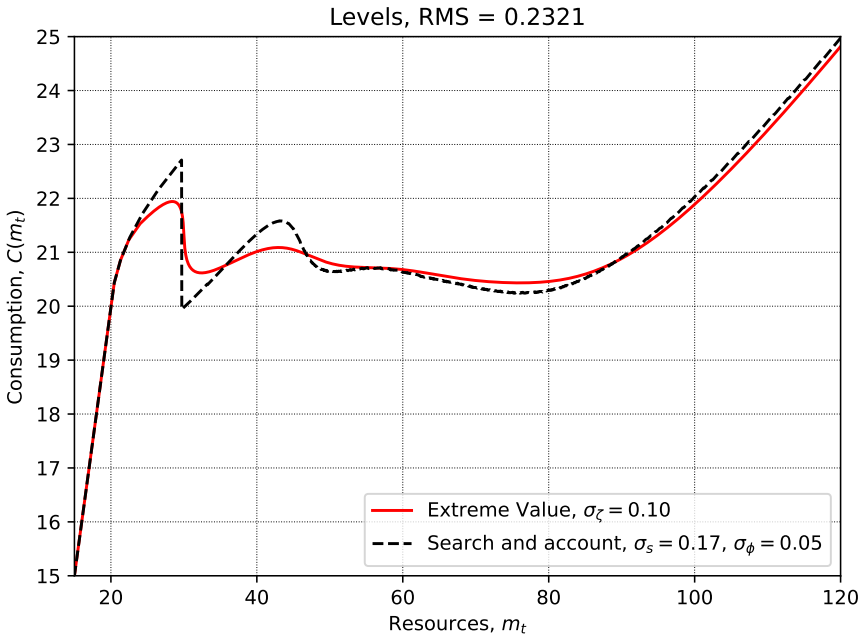


Figure 6: Worker consumption at $t = 15$. The Extreme Value rule uses $\sigma_\zeta = 0.10$, with search and account dispersion shut down. The search-and-account rule uses $\sigma_s = 0.171$ and $\sigma_\phi = 0.046$, the pair that minimizes the root mean squared gap between the two rules on resources from 15 to 120. Taste shocks are shut down on that rule, and the option arrives only through search. The minimized gap is 0.23.

2.8 Algorithm

This section outlines the algorithm for solving the problem of the worker who continues to work in the current period, when search and account dispersion are both active. Effort is the probability of receiving the option next period. When search is shut down, the retirement choice or the account draw is available with certainty and this effort step is not used. The retiree is solved first, by the standard consumption-saving EGM, and is evaluated later at claimed wealth ϕM. Taste shocks are shut down, so the option comparison is a hard maximum. Values are stored in the inverse form $V^T = -1/V$ and recovered by $V = -1/V^T$ before that comparison. After the EGM step, the upper envelope of **?** drops any decreasing segments of the endogenous grid.

2.9 Other smoothing mechanisms

Three devices smooth the same margin and are not part of the comparison above. The specifications they had in this model are recorded here so the later draft does not have to reconstruct them.

Transitory income. Labor income while working was $y_{t+1} = y\eta_{t+1}$, with η_{t+1} mean-one log-normal, $\log \eta_{t+1} \sim \mathcal{N}(-\sigma_\eta^2/2, \sigma_\eta^2)$. Next period's book resources $M_{t+1} = RA_t + y\eta_{t+1}$ were random,

Algorithm 1: EGM-step with search and account dispersion for the worker who continues to work

Let $\vec{\phi} = \{\phi_1, \ldots, \phi_K\}$ be quadrature nodes for the account price, with weights $\vec{\mu} = \{\mu_1, \ldots, \mu_K\}$.

Form an ascending grid over end-of-period wealth, $\vec{A}_t = \{A_t^1, \ldots, A_t^I\}$.

for $i = 1, \ldots, I$ **do**

- Compute book resources $M_{t+1}(A_i) = RA_i + y$.
- Interpolate the worker at book resources: $c_{t+1}^W(M_{t+1})$ and $V_{t+1}^W = -1/V_{t+1}^{T,W}(M_{t+1})$. The worker does not depend on ϕ.
- **for** $k = 1, \ldots, K$ **do**
  - Compute claimed wealth $\tilde{M} = \phi_k M_{t+1}(A_i)$.
  - Interpolate the retiree at claimed wealth: $c_{t+1}^{\text{Ret}}(\tilde{M})$ and $V_{t+1}^{\text{Ret}} = -1/V_{t+1}^{T,\text{Ret}}(\tilde{M})$.
  - Set $V_{t+1}^{RW}(M_{t+1}, \phi_k) = \max\{V_{t+1}^{\text{Ret}}(\tilde{M}),\, V_{t+1}^W(M_{t+1})\}$.
  - Set $\mathbb{1}_{t+1}^0 = 1$ if $V_{t+1}^{\text{Ret}}(\tilde{M}) \geq V_{t+1}^W(M_{t+1})$, and 0 otherwise; set $\mathbb{1}_{t+1}^1 = 1 - \mathbb{1}_{t+1}^0$.
  - Set $\text{RHS}_1^{\text{pre}}(M_{t+1}, \phi_k) = \mathbb{1}_{t+1}^0 \cdot \phi_k \cdot u'(c_{t+1}^{\text{Ret}}(\tilde{M})) + \mathbb{1}_{t+1}^1 \cdot u'(c_{t+1}^W(M_{t+1}))$.
- Average over the account price: $\mathcal{V}_t^1(A_i) = \sum_k \mu_k V_{t+1}^{RW}(M_{t+1}, \phi_k)$ and $\text{RHS}_1(A_i) = \sum_k \mu_k \text{RHS}_1^{\text{pre}}(M_{t+1}, \phi_k)$.
- Set $\mathcal{V}_t^2(A_i) = V_{t+1}^W(M_{t+1})$ and $\text{RHS}_2(A_i) = u'(c_{t+1}^W(M_{t+1}))$. The must-work branch is not averaged over ϕ.
- If $\sigma_s = 0$, set $\varepsilon_t(A_i) = 1$ when $\mathcal{V}_t^1(A_i) > \mathcal{V}_t^2(A_i)$ and 0 otherwise. Otherwise $\varepsilon_t(A_i) = \max\big\{0,\ 1 - \exp\big(-\beta(\mathcal{V}_t^1(A_i) - \mathcal{V}_t^2(A_i))/\sigma_s\big)\big\}$.
- Compute $\theta(\varepsilon_t(A_i)) = \sigma_s\big(\varepsilon_t(A_i) + (1 - \varepsilon_t(A_i))\log(1 - \varepsilon_t(A_i))\big)$, with $\theta = 0$ if $\varepsilon_t = 0$ and $\theta = \sigma_s$ if $\varepsilon_t = 1$.
- Set $P_1(\varepsilon_t) = \varepsilon_t(A_i)$ and $P_2(\varepsilon_t) = 1 - \varepsilon_t(A_i)$.
- Compute $\text{RHS}(A_i) = R\big(P_1\,\text{RHS}_1(A_i) + P_2\,\text{RHS}_2(A_i)\big)$.
- Compute $\mathbb{E}V_{t+1}(A_i) = P_1\,\mathcal{V}_t^1(A_i) + P_2\,\mathcal{V}_t^2(A_i)$.
- Compute $c_t^W(A_i) = u'^{-1}\big(\beta\,\text{RHS}(A_i)\big)$ and $M_t(A_i) = A_i + c_t^W(A_i)$.
- Compute $V_t^W(M_t(A_i)) = u(c_t^W(A_i), 1, \varepsilon_t(A_i)) + \beta\,\mathbb{E}V_{t+1}(A_i)$, and store $V_t^{T,W} = -1/V_t^W$.

On the endogenous grid, fill the liquidity-constrained segment below the smallest $M_t(A_i)$ by setting consumption equal to resources.

Apply the upper envelope to $(M_t, c_t^W, V_t^{T,W})$ and interpolate the consumption and inverse-value functions on the common M-grid.

so $\mathcal{V}_t^1$, $\mathcal{V}_t^2$, and the worker's Euler equation integrated over η_{t+1} as well as over ϕ_{t+1}. The figures set $\sigma_\eta = 0$.

Exogenous arrival. The probability of receiving the option was $P_1(\varepsilon_t) = (1-\varepsilon_t)p + \varepsilon_t$, and the probability of not receiving it was $P_2(\varepsilon_t) = (1-\varepsilon_t)(1-p)$. The factor $(1-p)$ multiplied the value gap in the effort rule,

$$\varepsilon_t^* = \max\left\{0,\ 1 - \exp\left(-\frac{\beta(1-p)}{\sigma_s}\left(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\right)\right)\right\}. \tag{37}$$

At $p = 1$ the option arrived for sure and optimal effort was zero. At $p = 0$ the arrival probability equals effort, which is the case in the text above. The Extreme Value and account-only figures used $p = 1$. The search and search-plus-account figures used $p = 0$.

Multiplicative taste shock. A mean-one log-normal shock ξ_t, with scale σ_ξ, multiplied the working branch from outside the consumption and effort choice,

$$V_t^W(M_t, \xi_t) = \xi_t \max_{\substack{0 \le c_t \le M_t \\ 0 \le \varepsilon_t \le 1}} \left\{\log(c_t) - \delta - \theta(\varepsilon_t) + \beta P_1(\varepsilon_t)\,\mathcal{V}_t^1(A_t) + \beta P_2(\varepsilon_t)\,\mathcal{V}_t^2(A_t)\right\}. \tag{38}$$

The envelope on that branch was ξ_t / c_t^W. In the terminal period, $V_T^W(M_T, \xi_T) = \xi_T\big(\log(M_T) - \delta\big)$. This is not the Extreme Value Type I shock: ξ_t rescales the value of working, and the expectation over it is a quadrature, not the closed-form log-sum. The comparison in this section shuts ξ_t down.

3 Durable model (Basic model)

The second example is a consumption-savings model with durable and nondurable goods. It is a simplified version of the model in **?** which uses endogeneous search as a numerical tool to solve the model. This section focuses on the solution method and compares the usage of taste shocks, endogeneous search, and stochastic durable prices.

3.1 Model setup

The model consists of ex-ante identical agents who face uninsurable income risk and can consume non-durable goods, receive flow utility from a durable good, and can save in one single risk-free liquit asset.

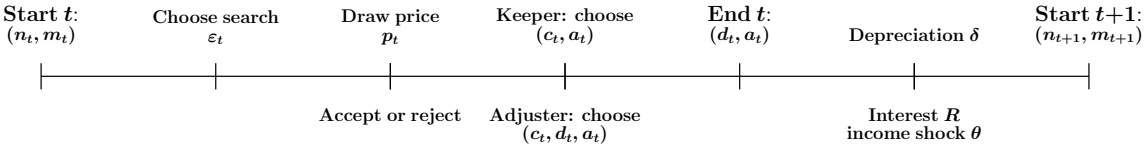


Figure 7: Timeline of the durable model

Figure 7 places search effort before the price draw. The household starts with durable stock n_t and market resources m_t, and chooses effort from the distribution of the purchase price. A match draws that price, and the household accepts the draw or keeps the current stock. The keeper then chooses nondurable consumption and assets. The adjuster chooses nondurable consumption, a new stock, and assets at the realized price. The period ends with (d_t, a_t). Between t and $t+1$, interest is received, the transitory income shock realizes, and the durable depreciates. With search and price

dispersion shut off, this order collapses to a hard maximum over the keeper and the adjuster. That baseline follows. Search and the price draw return together in Section 3.5.

Each household is characterized by two state variables: beginning-of-period market resources m_t and the stock of durables n_t. Preferences exhibit constant relative risk aversion and combine nondurable consumption c_t and the service flow of durables d_t using a Cobb-Douglas aggregator[1]:

$$u(c_t, d_t) = \frac{\left(c_t^{\alpha} d_t^{1-\alpha}\right)^{1-\rho}}{1-\rho}, \qquad \alpha \in (0,1),\ \rho > 0. \tag{39}$$

At the start of the period, the household observes (m_t, n_t) and decides whether to keep the durable stock unchanged or to adjust. If no adjustment occurs, the household sets $d_t = n_t$ and allocates m_t between nondurable consumption and savings. If the household adjusts, it sells the entire beginning-of-period stock at no price risk but incurs a proportional adjustment cost $\tau \in (0,1)$. The resulting liquidity is

$$\ell_t = m_t + (1-\tau) n_t, \tag{40}$$

which can then be allocated across nondurable consumption, a new durable stock d_t, and savings a_t.

The household's problem is the hard discrete choice between keeping and adjusting:

$$v_t(m_t, n_t) = \max\left\{ v_t^{\text{keep}}(m_t, n_t),\ v_t^{\text{adj}}(\ell_t) \right\}, \tag{41}$$

where v_t^{keep} and v_t^{adj} are the value functions conditional on keeping or adjusting.

Labor income is idiosyncratic and purely transitory. Normalized income is

$$y_t = \theta_t, \tag{42}$$

where θ_t is i.i.d. log-normally distributed with $\mathbb{E}[\theta_t] = 1$ and a strictly positive standard deviation. There is no unemployment risk.

Assets earn a single interest rate r, so the gross interest factor is $\mathcal{R} = 1 + r$. The durable stock depreciates at rate $\delta \in (0,1)$. There are no maintenance costs and no operation costs. Therefore, next-period states evolve according to

$$m_{t+1} = \mathcal{R} a_t + \theta_{t+1}, \tag{43}$$

$$n_{t+1} = (1-\delta)\, d_t. \tag{44}$$

Borrowing is prohibited: the loan-to-value ratio is zero and the lower bound on assets is $\underline{a} = 0$, so end-of-period assets must satisfy $a_t \geq 0$.

[1] Using this form is standard (eg see ??) which is in line with micro-evidence ?. An alternative approach is to use a CES aggregator as in ?

Hence, the Bellman equations conditional on keep and adjust are:

$$v_t^{\text{keep}}(m_t, n_t) = \max_{c_t} \; u(c_t, d_t) + \beta \mathbb{E}_t\big[v_{t+1}(m_{t+1}, n_{t+1})\big], \tag{45}$$

$$\text{s.t. } a_t = m_t - c_t \geq 0, \tag{46}$$

$$d_t = n_t, \tag{47}$$

$$m_{t+1} = \mathcal{R} a_t + \theta_{t+1}, \tag{48}$$

$$n_{t+1} = (1-\delta)\, d_t. \tag{49}$$

$$v_t^{\text{adj}}(\ell_t) = \max_{c_t, d_t} \; u(c_t, d_t) + \beta \mathbb{E}_t\big[v_{t+1}(m_{t+1}, n_{t+1})\big], \tag{50}$$

$$\text{s.t. } a_t = \ell_t - c_t - d_t \geq 0, \tag{51}$$

$$m_{t+1} = \mathcal{R} a_t + \theta_{t+1}, \tag{52}$$

$$n_{t+1} = (1-\delta)\, d_t. \tag{53}$$

With the following first-order conditions with respect to nondurable consumption and durable stock:

$$\frac{\partial v_t^{adj}}{\partial c_t} = u_c(c_t, d_t) \geq \beta \mathcal{R}\, \mathbb{E}_t\big[v_{t+1}^m\big], \tag{54}$$

$$\frac{\partial v_t^{adj}}{\partial d_t} = u_d(c_t, d_t) \geq \beta \mathcal{R}\, \mathbb{E}_t\big[v_{t+1}^m\big] - \beta(1-\delta)\, \mathbb{E}_t\big[v_{t+1}^n\big], \tag{55}$$

Equations 54 and 55 are stated as KKT conditions. Equality holds when the borrowing constraint does not bind, and the left-hand side strictly exceeds the right-hand side when it does. Equation 54 equates the marginal utility of nondurable consumption with the discounted shadow value of savings. Equation 55 shows that the marginal utility of durables equals the net user cost: the first term on the right is the discounted opportunity cost $\mathcal{R}$ of diverting a unit of resources into durables rather than liquid assets, and the second term is the discounted benefit of carrying the remaining durable stock $(1-\delta)d_t$ into the next period.

3.2 Algorithm

This section outlines the algorithm for one recursive period of the baseline model: no Extreme Value Type I taste shocks, no directed search, and no durable-price dispersion. The discrete keep-versus-adjust choice is resolved by a hard maximum. The terminal period is closed form—the keeper sets $c_T = m_T$ and $d_T = n_T$, the adjuster sets $c_T = \alpha \ell_T$ and $d_T = (1-\alpha)\ell_T$, and the two branches are combined by a hard max—so Algorithm 2 starts from a given next-period solution $(v_{t+1},\, v_{t+1}^m,\, v_{t+1}^n)$.

Define the end-of-period (gothic) objects on (d_t, a_t) by

$$\mathfrak{v}_c'(d_t, a_t) = \beta\, \mathcal{R}\, \mathbb{E}_t\big[v_{t+1}^m(n_{t+1}, m_{t+1})\big], \tag{56}$$

$$\mathfrak{v}_d'(d_t, a_t) = \beta\, (1-\delta)\, \mathbb{E}_t\big[v_{t+1}^n(n_{t+1}, m_{t+1})\big], \tag{57}$$

$$\mathfrak{v}(d_t, a_t) = \beta\, \mathbb{E}_t\big[v_{t+1}(n_{t+1}, m_{t+1})\big], \tag{58}$$

with $m_{t+1} = \mathcal{R} a_t + \theta_{t+1}$ and $n_{t+1} = (1-\delta)d_t$. The expectation over θ_{t+1} may be evaluated by quadrature (as written below) or by the transition-matrix integration described in Appendix B. Because adjustment costs make the first-order conditions only necessary, the keeper step applies a DC-EGM upper envelope [**?**] and the adjuster step applies FUES [**?**].

Algorithm 2: EGM step for the baseline durable model (hard max)

Let $\vec{\theta} = \{\theta_1, \ldots, \theta_J\}$ be quadrature points for the transitory income shock with weights $\vec{\omega} = \{\omega_1, \ldots, \omega_J\}$.

Form ascending grids $\vec{a}_t = \{a^1, \ldots, a^I\}$, $\vec{d}_t = \{d^1, \ldots, d^L\}$, a fine durable grid $\vec{d}_t^{\,\text{adj}}$ for the adjuster FOC scan, and beginning-of-period grids $\vec{n}_t$, $\vec{m}_t$.

`/* End-of-period (gothic) functions */`

for $i = 1, \ldots, I$ **do**
  for $q = 1, \ldots, L$ **do**
    for $j = 1, \ldots, J$ **do**
      Compute $m_{t+1}^j(a^i) = \mathcal{R} a^i + \theta_j$ and $n_{t+1}(d^q) = (1-\delta) d^q$
      Interpolate v_{t+1}, v_{t+1}^m, and v_{t+1}^n at $(n_{t+1}(d^q),\, m_{t+1}^j(a^i))$
    Compute $\mathfrak{v}'_c(d^q, a^i) = \beta \mathcal{R} \sum_{j=1}^J \omega_j \, v_{t+1}^m(n_{t+1}(d^q), m_{t+1}^j(a^i))$
    Compute $\mathfrak{v}'_d(d^q, a^i) = \beta (1-\delta) \sum_{j=1}^J \omega_j \, v_{t+1}^n(n_{t+1}(d^q), m_{t+1}^j(a^i))$
    Compute $\mathfrak{v}(d^q, a^i) = \beta \sum_{j=1}^J \omega_j \, v_{t+1}(n_{t+1}(d^q), m_{t+1}^j(a^i))$

`/* Keeper: EGM with` $d_t = n_t$`, then DC-EGM upper envelope */`

for *each* $n \in \vec{n}_t$ **do**
  for $i = 1, \ldots, I$ **do**
    Compute $c_t^{\text{keep}}(a^i; n) = u_c^{-1}\big(\mathfrak{v}'_c(n, a^i);\, n\big)$
    Compute endogenous resources $m_t^{\text{EGM}}(a^i; n) = a^i + c_t^{\text{keep}}(a^i; n)$
    Compute $v_t^{\text{EGM}}(a^i; n) = u\big(c_t^{\text{keep}}(a^i; n), n\big) + \mathfrak{v}(n, a^i)$
  Apply a DC-EGM upper envelope to $\{(m_t^{\text{EGM}},\, c_t^{\text{keep}},\, v_t^{\text{EGM}})\}$ and interpolate onto $\vec{m}_t$ to obtain $c_t^{\text{keep}}(n, m)$, $a_t^{\text{keep}} = \max\{m - c_t^{\text{keep}}, 0\}$, and $d_t^{\text{keep}} = n$
  Compute $v_t^{\text{keep}}(n, m) = u(c_t^{\text{keep}}, n) + \mathfrak{v}(n, a_t^{\text{keep}})$, $v_t^{m,\text{keep}} = u_c(c_t^{\text{keep}}, n)$, and $v_t^{n,\text{keep}} = u_d(c_t^{\text{keep}}, n) + \mathfrak{v}'_d(n, a_t^{\text{keep}})$

`/* Adjuster: FOC roots in` a `on a fine` d`-grid, then FUES in liquidity */`

Initialize empty candidate clouds for $(\ell_t, c_t, d_t, a_t, v_t)$

for *each* $d \in \vec{d}_t^{\,\text{adj}}$ **do**
  Find all roots a of $u_d(c, d) - \mathfrak{v}'_c(d, a) + \mathfrak{v}'_d(d, a) = 0$ with $c = u_c^{-1}\big(\mathfrak{v}'_c(d, a);\, d\big)$
  for *each root* a **do**
    Compute $\ell_t = a + c + d$ and $v_t^{\text{adj}} = u(c, d) + \mathfrak{v}(d, a)$
    Append $(\ell_t, c, d, a, v_t^{\text{adj}})$ to the candidate cloud

Sort candidates by ℓ_t and apply FUES to obtain refined points $(\ell_t, c_t^{\text{adj}}, d_t^{\text{adj}}, a_t^{\text{adj}}, v_t^{\text{adj}})$

Construct one-dimensional interpolants $c_t^{\text{adj}}(\ell)$, $d_t^{\text{adj}}(\ell)$, and $a_t^{\text{adj}}(\ell)$

for *each* $(n, m) \in \vec{n}_t \times \vec{m}_t$ **do**
  Compute $\ell_t = m + (1-\tau) n$
  Evaluate $c_t^{\text{adj}}(\ell_t)$, $d_t^{\text{adj}}(\ell_t)$, and $a_t^{\text{adj}} = \max\{\ell_t - c_t^{\text{adj}} - d_t^{\text{adj}}, 0\}$
  Compute $v_t^{\text{adj}}(\ell_t) = u(c_t^{\text{adj}}, d_t^{\text{adj}}) + \mathfrak{v}(d_t^{\text{adj}}, a_t^{\text{adj}})$, $v_t^{m,\text{adj}} = u_c(c_t^{\text{adj}}, d_t^{\text{adj}})$, and $v_t^{n,\text{adj}} = (1-\tau)\, u_c(c_t^{\text{adj}}, d_t^{\text{adj}})$

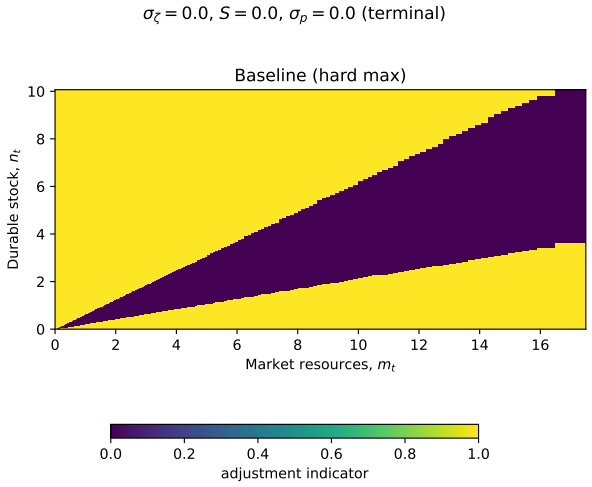
`/* Hard-max discrete choice */`

for *each* $(n, m) \in \vec{n}_t \times \vec{m}_t$ **do**
  Set $\mathcal{I}_t(n, m) = 1$ if $v_t^{\text{adj}}\big(m + (1-\tau) n\big) \geq v_t^{\text{keep}}(n, m)$ and 0 otherwise
  Set $(c_t, d_t, a_t, v_t, v_t^m, v_t^n)$ to the adjuster branch if $\mathcal{I}_t = 1$ and to the keeper branch if $\mathcal{I}_t = 0$

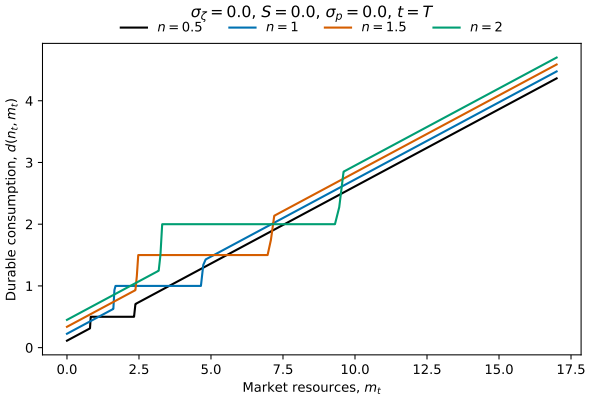
Collect the combined policies and (inverse) value functions on $(\vec{n}_t, \vec{m}_t)$ for use in period $t-1$.

3.3 Illustration of the numerical challenge

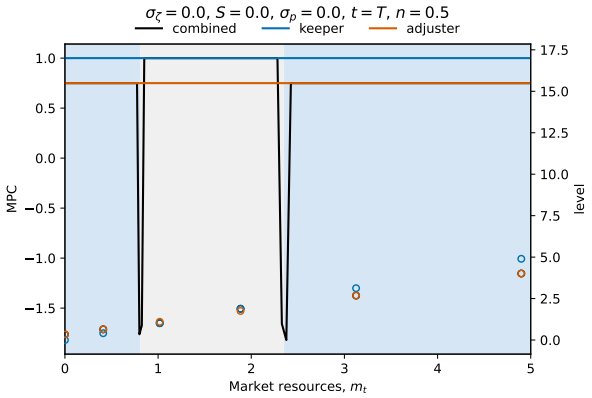
Figure 8 is the last period, when nothing in the continuation can smooth today's choice. Figure 9 is an earlier period, after the jump at that boundary has been passed backward.



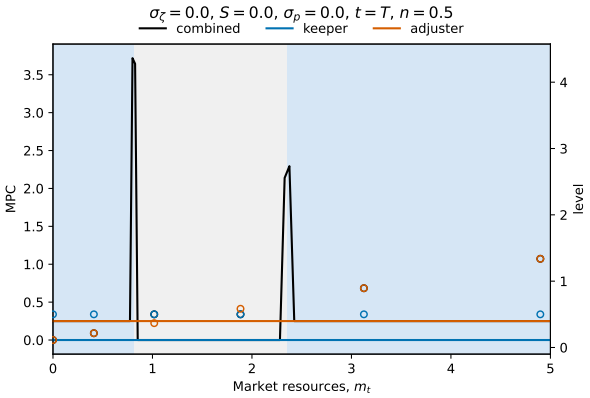
(a) Adjustment region (S,s bands).



(b) Durable policy $d_t(n_t, m_t)$.



(c) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.



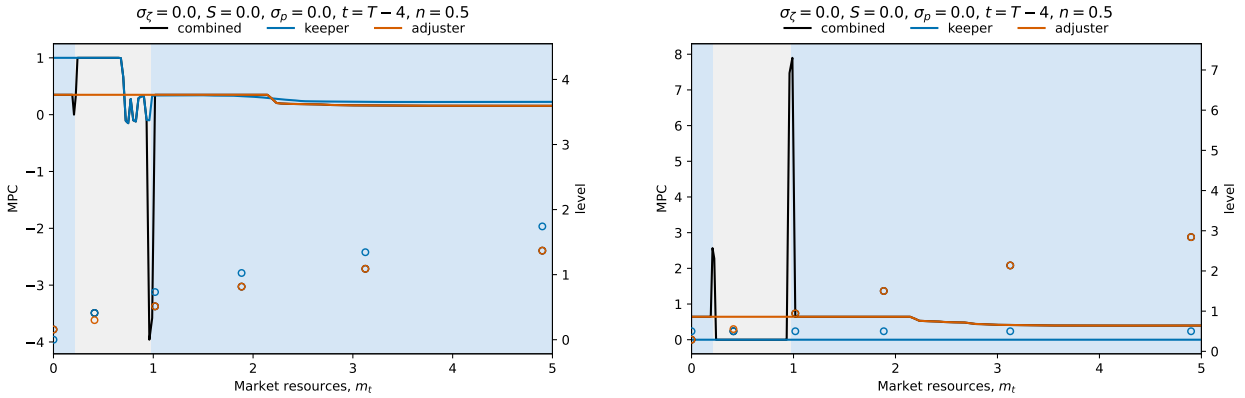
(d) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 8: Terminal-period solution of the baseline durable model (hard max). The upper left panel shows the keep–adjust regions in (m_t, n_t) space; the upper right panel shows the combined durable policy for selected beginning-of-period durable stocks n_t. The lower panels show the MPC decomposition of nondurable and durable consumption given a fixed beginning-of-period durable stock n_t.

The upper left panel is the adjustment region in (m_t, n_t). The household adjusts with probability one in one region and keeps with probability one in the other. The boundary between them is the (S, s) band. Inside it, the beginning-of-period stock is close enough to the stock the household would choose that the proportional cost is not worth paying. Outside it, the household pays the cost and reoptimizes. Because the choice is an indicator, the combined durable policy in the upper right panel jumps between the keeper, who is stuck at $d_T = n_T$, and the adjuster, who sets $d_T = (1 - \alpha)\ell_T$. The lower panels hold the beginning-of-period stock fixed and plot the combined rule together with the two branches it switches between. Where the band is crossed, the marginal propensity inherits

the jump. The spike is the discrete change of branch, rather than a feature of either branch's own Euler equation.

In every earlier period the household anticipates that jump. Next period's value is kinked along the future band, so today's first-order conditions are necessary and not sufficient, and the endogenous grid for the keeper can fold back on itself. The adjuster's scan of the durable first-order condition likewise returns several candidates. Algorithm 2 therefore applies an upper envelope on each branch, DC-EGM for the keeper and FUES for the adjuster. The envelope removes the dominated candidates on a branch. Joining the two refined branches by a hard maximum puts the jump back into the combined policy, and that combined policy is the continuation of the previous period. The nonconvexity is passed backward, as in the retirement problem. Figure 9 shows the resulting rules at a fixed beginning-of-period stock. Where adjusting is optimal, the combined rule and its marginal propensity lie on the adjuster. Once market resources cross the band they lie on the keeper, and both marginal propensities spike at the crossing. The extensions that follow replace this indicator with a probability. The remainder of the section asks whether that replacement removes the spike, and whether the three replacements can be told apart.



(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.

(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 9: MPC and policy functions of non-durable and durable consumption.

3.4 Extension: Extreme Value Type I Taste Shocks

The hard maximum in the discrete keep-versus-adjust decision produces a kink in the value function and a discontinuity in the policy functions at the state combinations (m_t, n_t) where the household is indifferent between the two branches. Following the discrete–continuous choice literature (e.g. **?**), this subsection smooths that discontinuity with i.i.d. Extreme Value Type I (Gumbel) taste shocks on the two discrete alternatives, while leaving search frictions aside (these are introduced separately in Section 3.5). EV Type I shocks enter *additively* on each branch and are integrated out in closed form: they randomize the discrete outcome directly through logit choice probabilities, without introducing an additional continuous (or discretized) state.

At the start of period t, after observing (m_t, n_t), the household draws i.i.d. Extreme Value Type I shocks $(\zeta_t^{\text{keep}}, \zeta_t^{\text{adj}})$ and chooses the branch that maximizes the shock-augmented conditional value,

$$v_t(m_t, n_t) = \max\left\{v_t^{\text{keep}}(m_t, n_t) + \sigma_\zeta\, \zeta_t^{\text{keep}},\ v_t^{\text{adj}}(\ell_t) + \sigma_\zeta\, \zeta_t^{\text{adj}}\right\}, \tag{59}$$

where $\sigma_\zeta > 0$ is the taste-shock scale and $\ell_t = m_t + (1-\tau)n_t$. The conditional value functions $v_t^{\text{keep}}(m_t, n_t)$ and $v_t^{\text{adj}}(\ell_t)$ retain exactly the same Bellman definitions as in the baseline model: the taste shocks do not enter the inner continuous maximizations, so the first-order conditions (54)–(55) and the continuous policies c_t^{keep}, c_t^{adj}, d_t^{adj} are unchanged.

Because $(\zeta_t^{\text{keep}}, \zeta_t^{\text{adj}})$ are i.i.d. Extreme Value Type I, the expectation of (59) has the closed-form log-sum representation

$$v_t(m_t, n_t) = \sigma_\zeta \log\Big(\exp\big(v_t^{\text{keep}}(m_t, n_t)/\sigma_\zeta\big) + \exp\big(v_t^{\text{adj}}(\ell_t)/\sigma_\zeta\big)\Big), \tag{60}$$

with logit choice probabilities

$$P_t^{\text{adj}}(m_t, n_t) = \frac{\exp\big(v_t^{\text{adj}}(\ell_t)/\sigma_\zeta\big)}{\exp\big(v_t^{\text{keep}}(m_t, n_t)/\sigma_\zeta\big) + \exp\big(v_t^{\text{adj}}(\ell_t)/\sigma_\zeta\big)}, \qquad P_t^{\text{keep}}(m_t, n_t) = 1 - P_t^{\text{adj}}(m_t, n_t). \tag{61}$$

The expected (ex-ante) policy functions are the probability-weighted averages of the two branches,

$$c_t(m_t, n_t) = P_t^{\text{keep}}\, c_t^{\text{keep}}(m_t, n_t) + P_t^{\text{adj}}\, c_t^{\text{adj}}(\ell_t), \tag{62}$$

$$d_t(m_t, n_t) = P_t^{\text{keep}}\, n_t + P_t^{\text{adj}}\, d_t^{\text{adj}}(\ell_t). \tag{63}$$

As $\sigma_\zeta \to 0$, the logit probabilities collapse to the hard indicators of the baseline model and (60) collapses to the hard maximum. Larger σ_ζ spreads probability mass onto the inferior branch near indifference, so c_t and d_t (and the marginal-value functions derived from the inclusive value (60)) are smooth in (m_t, n_t). No quadrature over taste nodes is required: the EV Type I assumption delivers the inclusive value and choice probabilities in closed form. The taste shocks do not alter the household's budget set, so the no-borrowing constraint $a_t \geq 0$ is unaffected. I derive the envelope conditions, the Euler equations, and the EGM algorithm with EV Type I taste shocks in appendix B.2.

The inclusive value (60) is differentiable in (m_t, n_t) for any $\sigma_\zeta > 0$, so the terminal keep–adjust boundary is no longer a line. Figure 10 plots the adjustment probability for $\sigma_\zeta \in \{0.1, 0.15, 0.2\}$. Where the two conditional values are close, the logit probability lies strictly between zero and one, and a larger scale widens that band because the household needs a larger value gap before it is willing to put almost all of its probability on the better branch. The combined rules (62)–(63) are the same probabilities applied to the two branches. Around the old threshold the plotted consumption and durable policies are therefore mixtures of the keeper and the adjuster, rather than a jump from one to the other.

One period earlier, the continuation value that enters the endogenous grid is this smooth inclusive value, so the nonconvexity that the hard maximum passes backward is already integrated out. Figure 11 plots the combined consumption and durable rules, and their marginal propensities, at a fixed beginning-of-period stock. The jump at the old threshold and the spike in the marginal propensity are both replaced by a gradual transition between the keeper and the adjuster. A small taste scale is enough for that transition to be continuous. The same scale also moves probability onto the inferior branch at states far from indifference, which is the feature the steady-state comparison below will use to separate taste shocks from search.

3.5 Extension: Search and stochastic prices

This subsection replaces the hard keep-versus-adjust maximum with directed search and a stochastic purchase price, introduced together. Extreme Value Type I taste shocks and directed search are

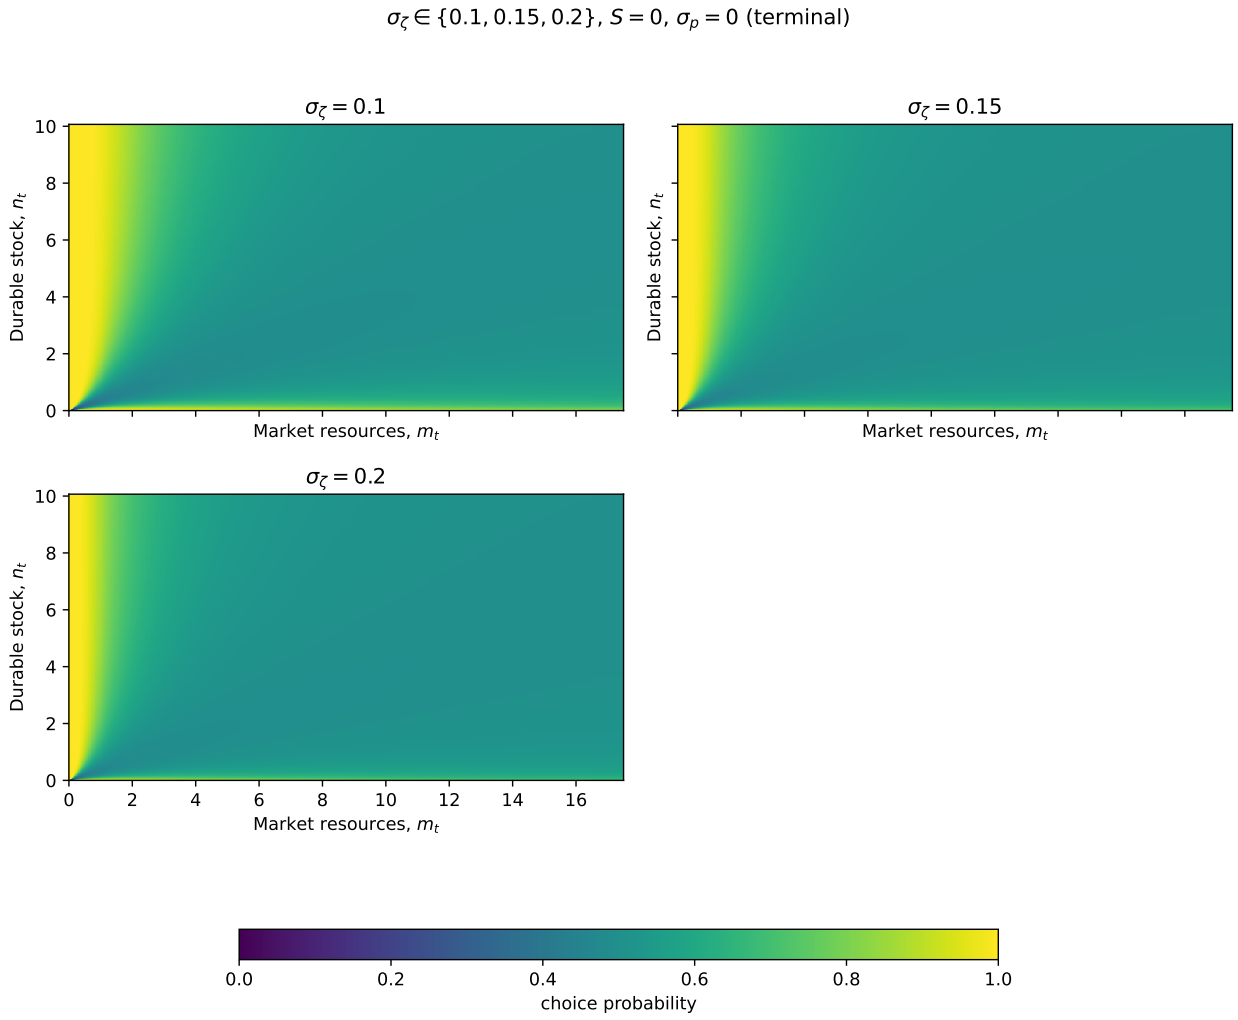
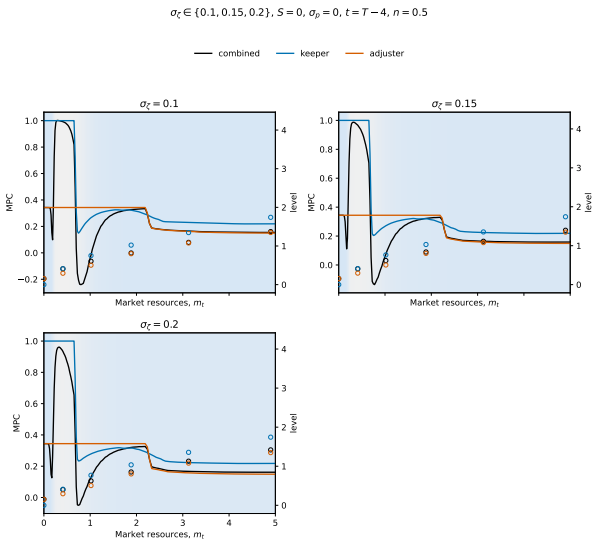
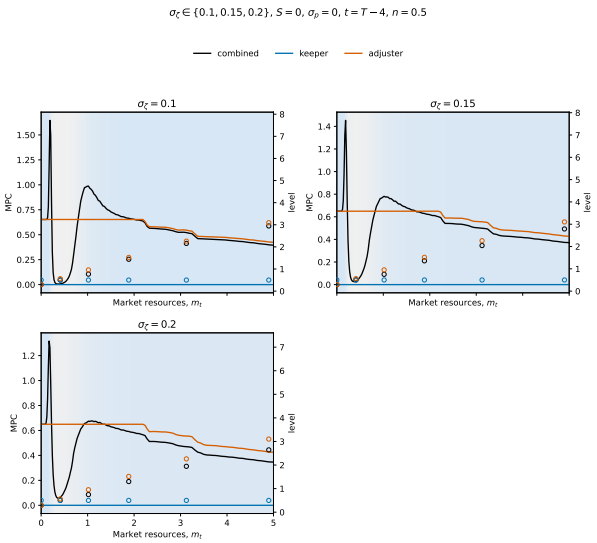


Figure 10: Terminal-period keep–adjust regions in (m_t, n_t) space under EV Type I taste shocks for $\sigma_\zeta \in \{0.1, 0.15, 0.2\}$.



(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.



(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 11: MPC and policy functions of non-durable and durable consumption using EV taste shocks.

alternative smoothers of the discrete margin; nesting the two would double-smooth that choice, so the environment here uses search together with price risk on the adjuster's continuous problem. The household chooses a search intensity $\varepsilon_t \in [0, 1]$ and, when it adjusts, buys the new durable at a realized relative price p_t.

Search effort is costly,

$$\sigma(\varepsilon_t) = S\big[\varepsilon_t + (1 - \varepsilon_t)\log(1 - \varepsilon_t)\big], \tag{64}$$

where $S > 0$ governs the friction in the market for durables, following **?**. The purchase price is i.i.d. across periods and log-normal with mean one,

$$\log p_t \sim \mathcal{N}\Big(-\frac{\sigma_p^2}{2}, \sigma_p^2\Big), \qquad \mathbb{E}[p_t] = 1, \tag{65}$$

$$p_t \perp \theta_t. \tag{66}$$

Only a purchase of new durables is exposed to p_t. The household liquidates the existing stock at the fixed recovery rate $1 - \tau$, so the liquidity available upon adjustment, $\ell_t = m_t + (1 - \tau)n_t$, does not depend on the price. The keeper never buys, so $v_t^{\text{keep}}(m_t, n_t)$ keeps its baseline Bellman definition. The adjuster's budget is

$$a_t = \ell_t - c_t - p_t\, d_t \geq 0, \tag{67}$$

and the price-contingent value is

$$v_t^{\text{adj}}(\ell_t, p_t) = \max_{c_t, d_t}\ u(c_t, d_t) + \beta\mathbb{E}_t\big[\bar{v}_{t+1}(m_{t+1}, n_{t+1})\big], \tag{68}$$

$$\text{s.t. equation (67)}, \tag{69}$$

$$m_{t+1} = \mathcal{R}a_t + \theta_{t+1}, \qquad n_{t+1} = (1 - \delta)\, d_t. \tag{70}$$

The interest rate $\mathcal{R}$, the depreciation rate δ, and the no-borrowing constraint $a_t \geq 0$ are unchanged. Next period's durable stock is measured in physical units, so p_t changes the current cost of acquiring d_t, and the future service flow changes only through depreciation.

The adjuster's Lagrangian, with p_t held fixed, has first-order conditions

$$\lambda_t^{\text{adj}} = u_c(c_t, d_t), \tag{71}$$

$$\lambda_t^{\text{adj}} = \beta \mathcal{R}\, \mathbb{E}_t\big[\bar{v}_{t+1}^m\big], \tag{72}$$

$$u_d(c_t, d_t) = p_t\, \lambda_t^{\text{adj}} - \beta(1-\delta)\, \mathbb{E}_t\big[\bar{v}_{t+1}^n\big]. \tag{73}$$

In KKT form these are

$$u_c(c_t, d_t) \geq \beta \mathcal{R}\, \mathbb{E}_t\big[\bar{v}_{t+1}^m\big], \tag{74}$$

$$u_d(c_t, d_t) \geq p_t\, \beta \mathcal{R}\, \mathbb{E}_t\big[\bar{v}_{t+1}^m\big] - \beta(1-\delta)\, \mathbb{E}_t\big[\bar{v}_{t+1}^n\big], \tag{75}$$

with equality when the borrowing constraint does not bind. Because p_t multiplies only d_t in the budget, the nondurable condition is the baseline condition (54) given d_t. The durable condition prices the marginal unit at p_t rather than at one. The adjuster's continuous policies $c_t^{\text{adj}}(\ell_t, p_t)$ and $d_t^{\text{adj}}(\ell_t, p_t)$ are therefore solved at each price node, and the reported ex-ante rules are price-weighted averages of those state-contingent choices.

After θ_t has materialized and (m_t, n_t) are known, the household chooses search intensity from the distribution of p_t, before the realization. A successful search then draws p_t. The household adjusts when the realized price is worth taking and otherwise keeps the current stock. The search cost is sunk, so a bad draw can be rejected. Let

$$v_t^{\text{match}}(m_t, n_t) \equiv \mathbb{E}_{p_t}\Big[\max\big\{v_t^{\text{adj}}(\ell_t, p_t),\ v_t^{\text{keep}}(m_t, n_t)\big\}\Big] \tag{76}$$

be the expected value of a match. The discrete problem is

$$v_t(m_t, n_t) = \max_{\varepsilon_t \in [0,1]} \Big[(1-\varepsilon_t)\, v_t^{\text{keep}}(m_t, n_t) + \varepsilon_t\, v_t^{\text{match}}(m_t, n_t) - \sigma(\varepsilon_t)\Big]. \tag{77}$$

Write $\Delta(p_t) = v_t^{\text{adj}}(\ell_t, p_t) - v_t^{\text{keep}}(m_t, n_t)$ and $\bar{\Delta}(m_t, n_t) = \mathbb{E}_{p_t}[\max\{\Delta(p_t), 0\}]$. Optimal effort is one number at (m_t, n_t),

$$\varepsilon_t^*(m_t, n_t) = \max\big\{0,\ 1 - \exp\big(-\bar{\Delta}(m_t, n_t)/S\big)\big\}. \tag{78}$$

It does not depend on the realized price. Let $\mathbb{1}(p_t) = 1$ when $\Delta(p_t) \geq 0$, and let $\pi_t = \mathbb{E}_{p_t}[\mathbb{1}(p_t)]$ be the probability that a draw is acceptable. The probability of buying is $\varepsilon_t^* \pi_t$. The ex-ante value that enters the previous period is

$$\bar{v}_t(m_t, n_t) = (1-\varepsilon_t^*)\, v_t^{\text{keep}}(m_t, n_t) + \varepsilon_t^*\, v_t^{\text{match}}(m_t, n_t) - \sigma(\varepsilon_t^*), \tag{79}$$

which replaces $v_t(m_t, n_t)$ inside the keeper and adjuster problems one period earlier. The two mechanisms are complementary. Price dispersion raises v_t^{match} above v_t^{keep} in states where adjusting at the mean price is worse but some draws are better, and that option value raises effort. Search governs whether a draw arrives. The realized price is the content of the draw.

This is the information order of Section 2. There, effort is chosen from the distribution of the account factor, the realization is observed only if the option arrives, and the household then claims or denies the draw. The remaining difference is the date of the benefit. Retirement effort chosen at

the end of period t is the probability that an option arrives at $t+1$, so β enters that effort rule and effort is irrelevant at the terminal date. Here the match and the acceptance decision fall in the same period, the effort rule has no discount factor, and search is defined in the terminal period because the benefit is the current comparison. The probability of retiring in that section is last period's effort times the probability that the hard maximum selects claiming. The probability of adjusting here is this period's effort times π_t.

Two limits isolate each mechanism. As $\sigma_p \to 0$, the price collapses to one. Then $\pi_t \in \{0, 1\}$, and whenever adjusting is better, (77) is the search problem

$$v_t(m_t, n_t) = \max_{\varepsilon_t \in [0,1]} \Big[(1 - \varepsilon_t)\, v_t^{\text{keep}}(m_t, n_t) + \varepsilon_t\, v_t^{\text{adj}}(\ell_t) - \sigma(\varepsilon_t) \Big], \tag{80}$$

with the same continuation one period ahead. As $S \to 0$, search is costless. Whenever $\bar{\Delta} > 0$, effort goes to one, the household draws a price, and the acceptance decision is the hard maximum

$$v_t(m_t, n_t, p_t) = \max \Big\{ v_t^{\text{keep}}(m_t, n_t),\ v_t^{\text{adj}}(\ell_t, p_t) \Big\}, \tag{81}$$

whose expectation is

$$\bar{v}_t(m_t, n_t) = \mathbb{E}_{p_t} \Big[\max \big\{ v_t^{\text{keep}}(m_t, n_t),\ v_t^{\text{adj}}(\ell_t, p_t) \big\} \Big]. \tag{82}$$

As $S \to \infty$, effort goes to zero and the household never adjusts. Both limits together, $\sigma_p \to 0$ and $S \to 0$, recover the hard-maximum baseline of Section 3.1. The optimal effort rule, the envelope conditions, and the EGM steps are derived in Appendices B.3, B.4, and B.5.

Search alone. Shutting the price down ($\sigma_p = 0$) isolates the extensive margin. Effort is then the probability of moving from the keeper to the adjuster, and the closed form (78) sets that probability to zero wherever adjusting does not raise value and raises it continuously once the gap is positive. The sharp boundary of the hard maximum therefore becomes a band of partial adjustment. Figure 12 plots that band in the terminal period for $S \in \{0.15, 0.2, 0.25\}$: a higher search cost makes the same value gap buy less effort, so the region of interior adjustment shrinks toward the states where adjusting is most valuable. One period earlier, the combined consumption and durable rules are the effort-weighted average of the two branches. Figure 13 shows the consequence. The jump at the old threshold is replaced by a gradual transition, and the spikes in the marginal propensities, which in the baseline are the discrete switch from one branch to the other, are damped with it.

Prices alone. Shutting search down ($S \to 0$) isolates the other margin. A draw arrives with probability one, and at each realized price the household still takes the hard maximum (81). A threshold therefore survives inside every price node. What softens is the probability, taken before the draw, that the realized price is worth accepting. Figure 14 plots that ex-ante adjustment region for $\sigma_p \in \{0.05, 0.1, 0.15\}$ at a fixed number of price nodes. Near the old keep–adjust boundary the probability lies strictly between zero and one, and a wider price distribution widens the set of states in which some draws are accepted and others are rejected. Figure 15 shows the corresponding period-t rules. They are smoother than the hard maximum and less smooth than search, because the average over prices does not remove the kink at a given price. Figure 16 holds σ_p fixed and varies the number of nodes used to compute that average. It is a check on the discretization, not a third economic experiment.

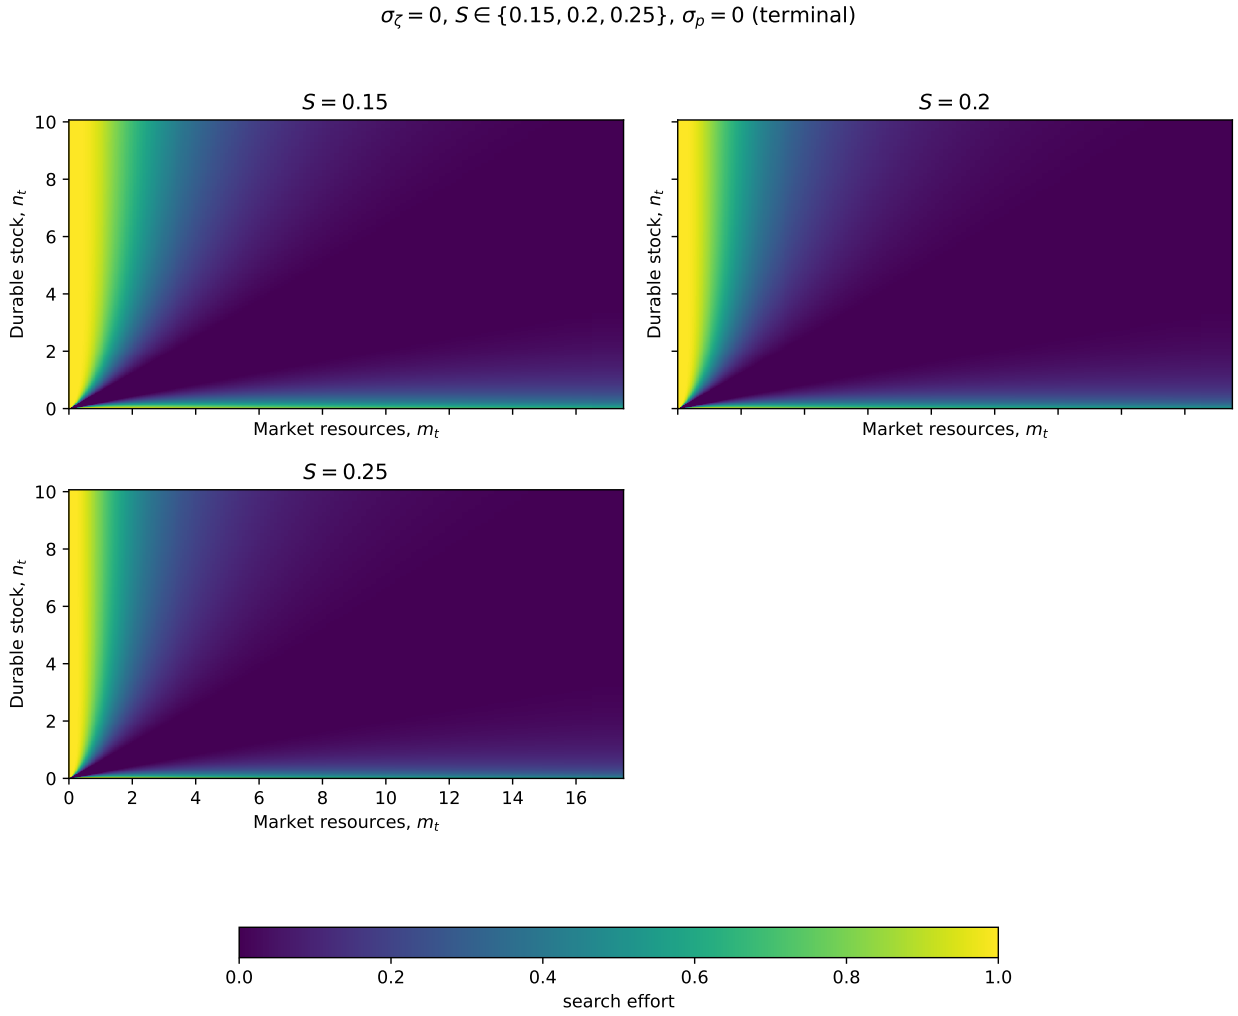
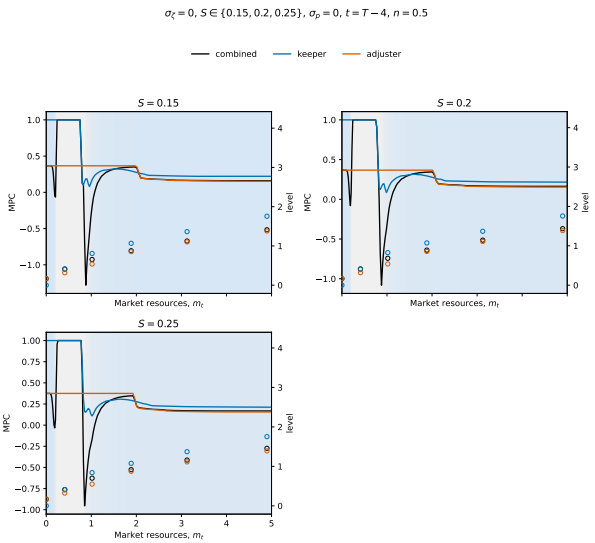
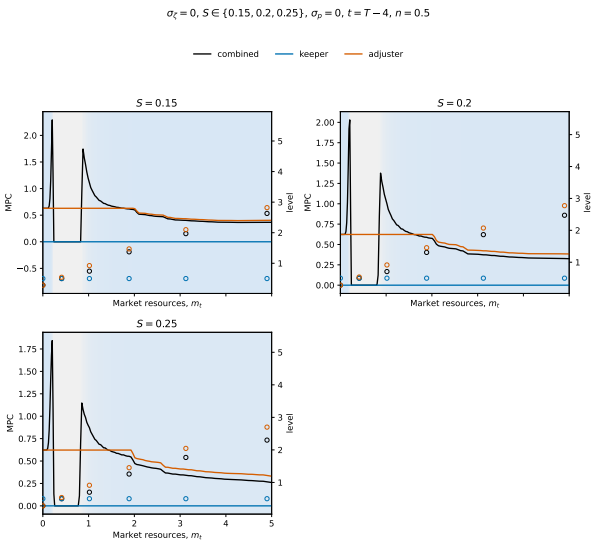


Figure 12: Terminal-period keep–adjust regions in (m_t, n_t) space under endogenous search for search-cost scales $S \in \{0.15, 0.2, 0.25\}$.



(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.



(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 13: MPC and policy functions of non-durable and durable consumption using endogeneous search.

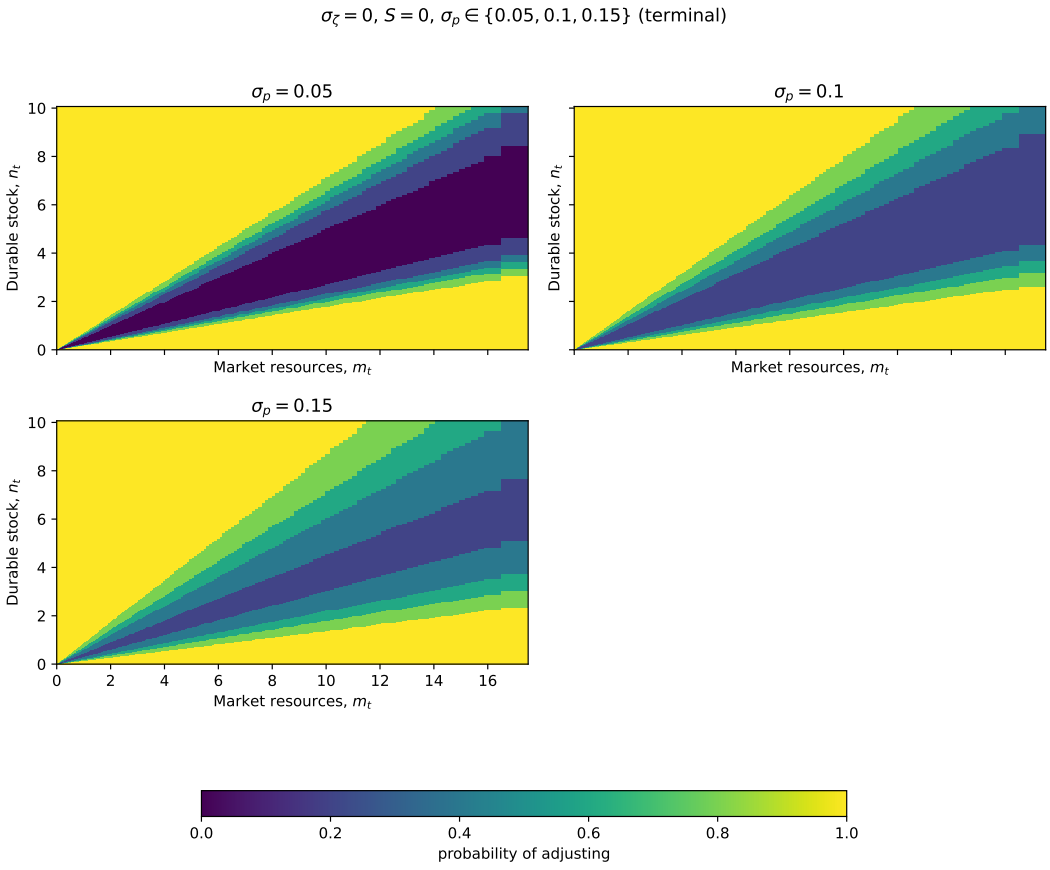
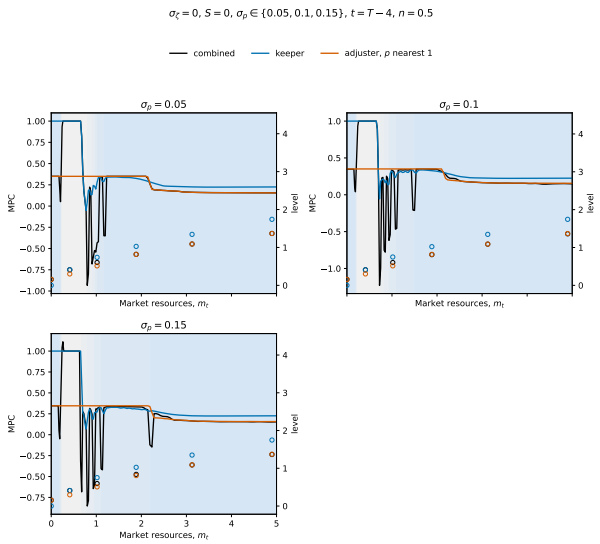
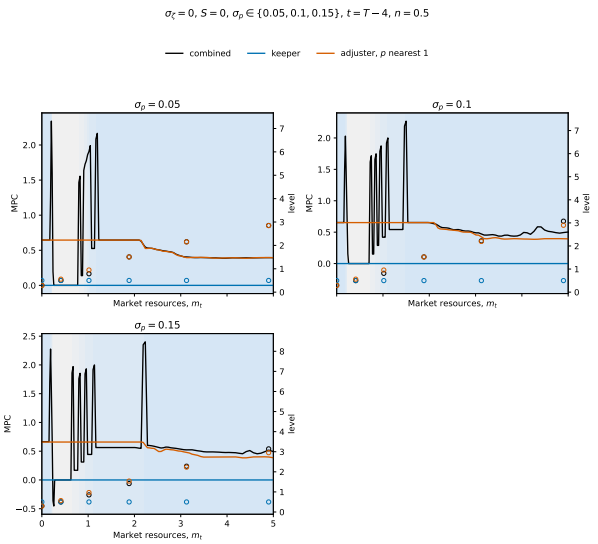


Figure 14: Adjustment regions in period T under stochastic durable prices for $\sigma_p \in \{0.05, 0.1, 0.15\}$ (fixed number of price nodes). Greater price dispersion softens the hard keep–adjust threshold.

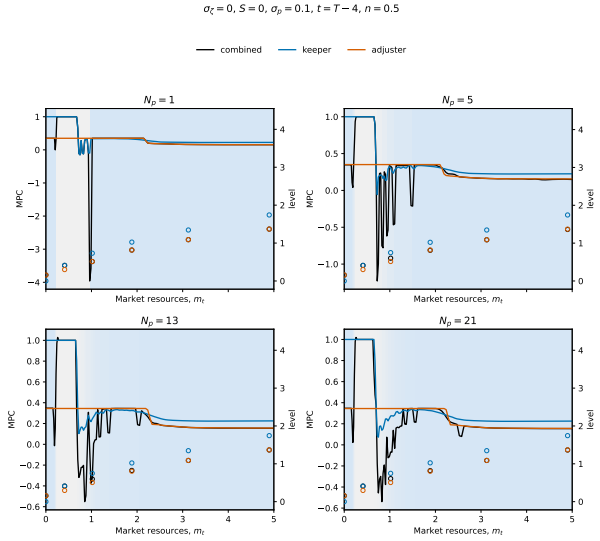


(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.

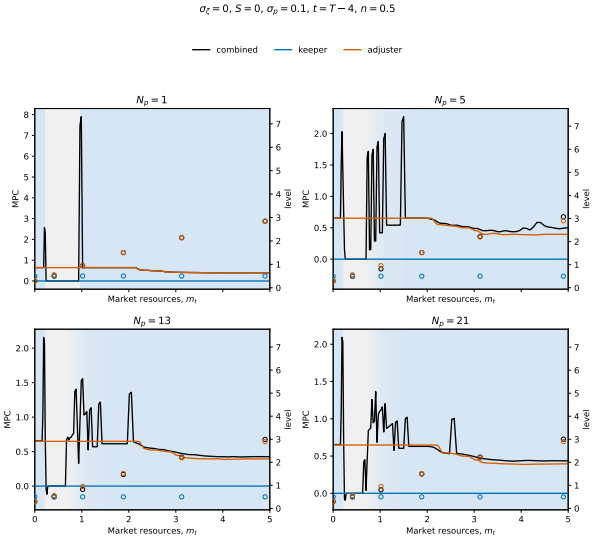


(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 15: MPC and policy functions of non-durable and durable consumption using stochastic prices.



(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.



(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 16: MPC and policy functions of non-durable and durable consumption varying the number of price points.

Search and prices together. The specification compared with Extreme Value taste shocks uses both margins at once. Effort is chosen from the distribution of the purchase price, a successful search draws that price, and a bad draw is rejected. States in which keeping dominates at the mean price can still adjust when a good draw arrives, and states in which adjusting dominates can still keep when the draw is poor. The probability of buying, effort times the acceptance probability, is therefore interior on both sides of the threshold that the hard maximum draws as a line. Figure 17 shows the terminal region under that pair, and Figure 18 shows the period-t consumption and durable rules. The one-at-a-time figures above are the two limits that isolate each piece. The comparison in the next subsection is about the pair.

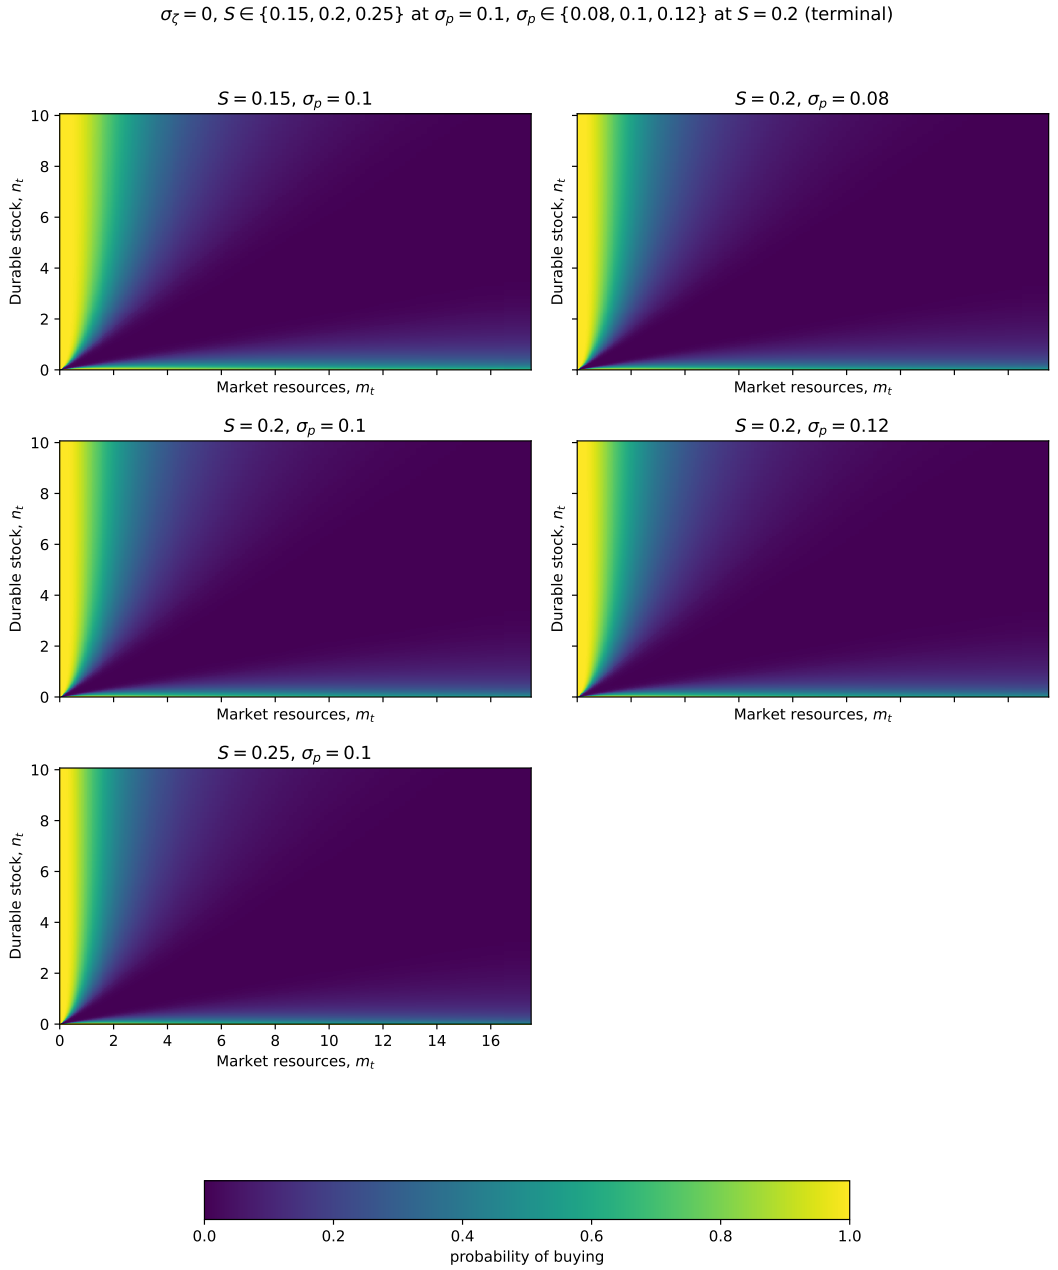
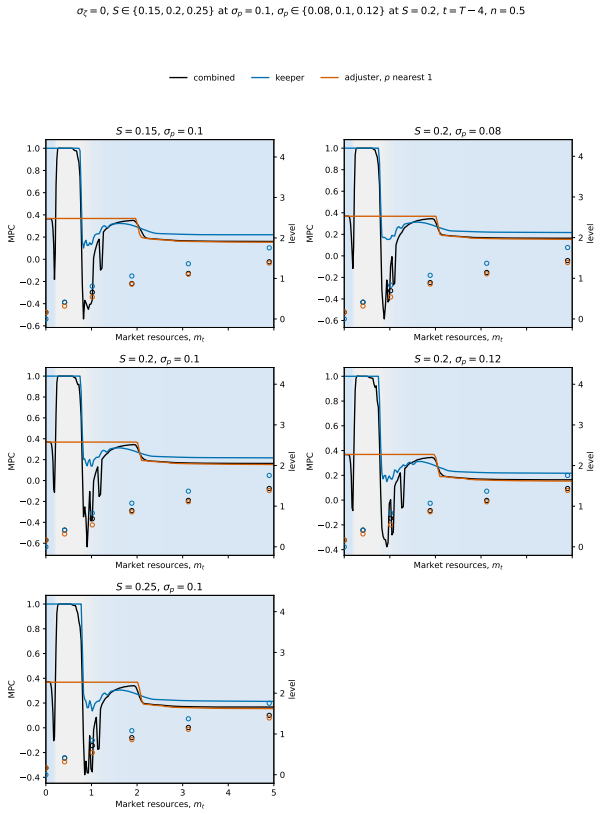
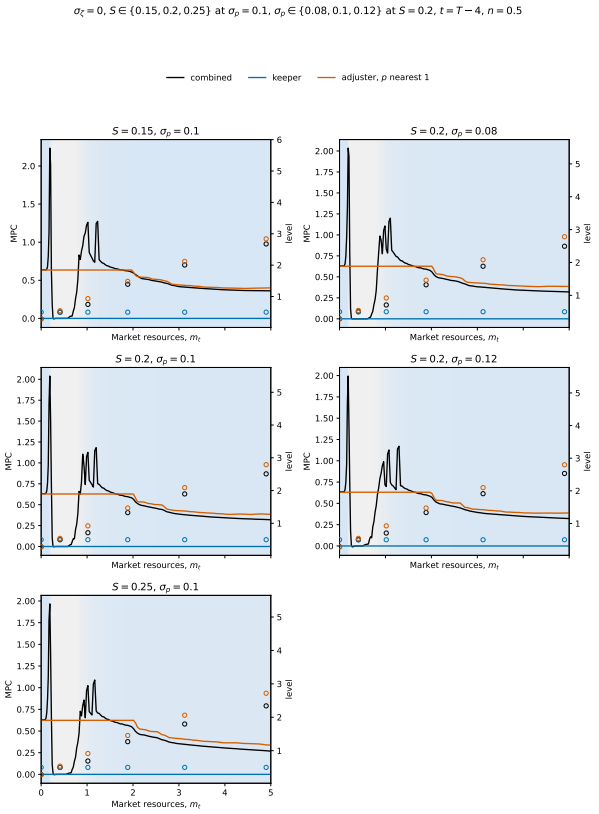


Figure 17: Adjustment regions in period T under stochastic durable prices and endogeneous search.



(a) MPC decomposition of nondurable consumption $c_t(n_t, m_t)$.



(b) MPC decomposition of durable consumption $d_t(n_t, m_t)$.

Figure 18: MPC and policy functions of non-durable and durable consumption using stochastic prices and search.

3.6 Comparing smoothness mechanisms: Lifecycle

The three extensions above each replace the hard keep-versus-adjust maximum with a probability that can lie between zero and one. This subsection asks whether that replacement produces comparable smoothness over a finite horizon, and whether one period's decision rule is enough to tell the devices apart. The answer to the second question is what sends the argument to the stationary distribution in Section 3.7.

Figures 10, 12, and 14 are the same terminal adjustment region once each device is turned on. In the baseline of Figure 8a the household adjusts with probability one outside the (S, s) band and keeps with probability one inside it. Under Extreme Value shocks that indicator becomes the logit probability (61). It rises continuously through the old boundary, and a larger σ_ζ spreads the rise over a wider set of states, because a larger scale makes the household willing to take the inferior branch at a larger value gap. Under search with the price shut off, the indicator becomes effort. Effort is zero where keeping is better and rises with the value gap where adjusting is better, so the band is the set of states in which a match is worth paying for. Under price dispersion with search shut off, each realized price still has its own threshold, and the figure plots the probability that a draw from the price distribution falls on the acceptable side of that threshold. A larger σ_p widens the set of states in which some draws are accepted and others are rejected. All three figures therefore replace that sharp boundary with a region in which the adjustment probability lies strictly between zero and one. The three regions are not the same object, which the kink counts make precise.

Table 1 counts the candidates that the upper envelope still has to delete, in the same earlier period as the policy figures below, on the raw endogenous-grid cloud before it is interpolated. For the keeper, a decreasing segment is a fold in the locus that maps end-of-period assets back to beginning-of-period market resources. Such a fold appears when next period's value is nonconcave, so the first-order condition is necessary and not sufficient, and the upper envelope of **?** discards it. For the adjuster, the scan of the durable first-order condition returns several candidates at nearby levels of liquidity, and FUES keeps only those on the upper envelope of the value cloud. Under the hard maximum the keeper still has five decreasing segments, about 4.5 percent of the locus, and the adjuster still loses four candidates. A taste scale of $\sigma_\zeta = 0.05$ and a search cost of $S = 0.1$ drive both counts to zero. The continuation that enters today's endogenous grid is already smooth enough that every first-order candidate lies on the envelope. Price dispersion at $\sigma_p = 0.1$ does the same for the keeper and leaves the adjuster with seventeen candidates still removed, about 3.5 percent of the cloud summed over price nodes. Each price node is its own hard-maximum problem, so nonconvexity survives inside the price-contingent choice even after the ex-ante probability has been averaged.

Search alone has a further limitation that a deletion count cannot see. Where the expected gain from a match is not strictly positive, effort in (78) is exactly zero, and the household keeps with probability one. Extreme Value shocks randomize that side of the threshold as well, because the logit assigns positive probability to adjustment at every finite value gap. Price dispersion alone randomizes through the draw and retains a threshold at each realized price. The specification whose smoothness is comparable to the taste shock is the pair: effort governs whether a draw arrives, and the household can still reject it, so the probability of buying is interior on both sides of the old boundary. A zero deletion count is also a statement about the raw cloud, not about the interpolant that is stored and plotted. Section 3.7 pairs these counts with off-grid Euler residuals.

The decision rule in a single period does not carry the distinction. Figures 19 and 20 hold the beginning-of-period durable stock fixed and overlay the combined consumption and durable rules, and their slopes with respect to market resources, as one scale is strengthened. Panels (a)–(c) are taste, search, and price dispersion, each with the other two shut off. In every one of those panels

the combined nondurable rule shifts up, the combined durable rule shifts down, and the spikes in the marginal propensities around the band are damped. Near indifference the household puts weight on both branches. The adjuster liquidates the current stock, pays the proportional cost, and reoptimizes, which on these grids means more nondurable consumption and a smaller chosen durable than the keeper, who is stuck with $d_t = n_t$. Strengthening the smoother widens the set of states that mix the two branches, so the combined rule moves toward the adjuster's bundle. Panel (d) raises the standard deviation of the transitory income shock. Income risk changes the continuation value of both branches. It does not replace the discrete choice, and it is drawn next to the other three so that a smoother operating through the income process is visible and is kept out of the comparison.

This common period-t shift is why the finite horizon cannot identify the three devices. In the retirement illustration, taste and search moved the worker's consumption function in opposite directions, so the decision rule itself separated them. Keep-versus-adjust reallocates a stock the household already owns, and it leaves the labor-income process unchanged, so the combined rules move together. The devices separate once households are allowed to accumulate. The probability of adjusting then changes how often the stock is reset and how much liquid wealth is held against that reset, which is a property of the stationary distribution rather than of one period's policy.

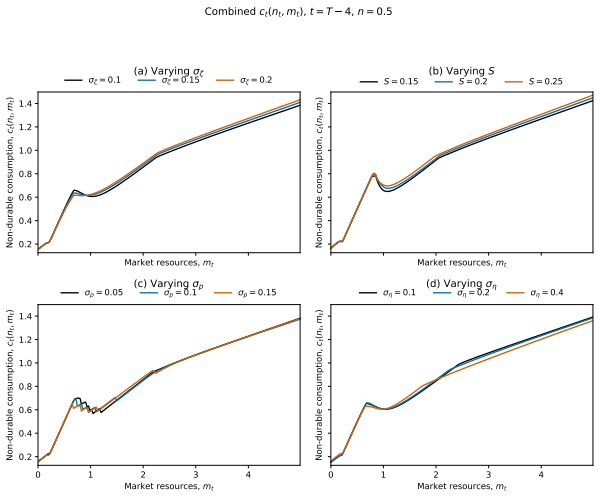
Table 1: Period-t kink diagnostics under the hard-max baseline and one smoothed magnitude of each discrete-choice channel. Keeper: number and share of decreasing EGM segments. Adjuster: number and share of FOC points removed by FUES. The σ_p adjuster count sums over price nodes.

	Keeper		Adjuster	
Specification	Decr. seg.	Share	Kinks rem.	Share
Hard max	5	0.045	4	0.040
$\sigma_\zeta = 0.05$	0	0	0	0
$S = 0.1$	0	0	0	0
$\sigma_p = 0.1$	0	0	17	0.035

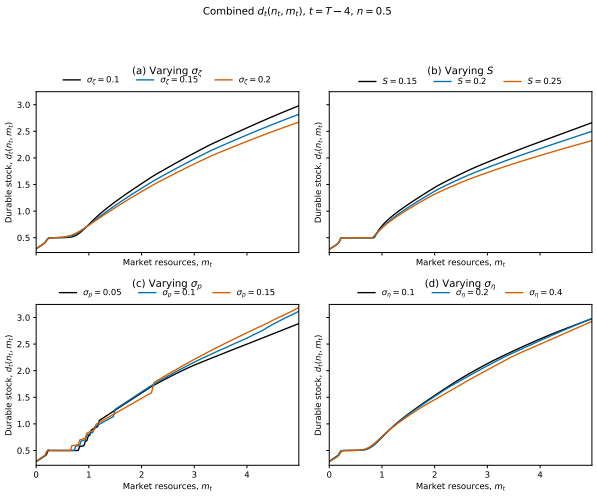
3.7 Comparing smoothness mechanisms: Infinite horizon

The finite-horizon comparison showed that taste shocks, search, and price dispersion each replace the keep–adjust line with a band, and that one period's combined rules move in the same direction when any of the three is strengthened. Households in that exercise do not accumulate. The probability of adjusting changes how often the durable is reset and how much liquid wealth is held against the reset, so the three devices can separate only in a stationary distribution. This subsection solves each economy to a time-invariant policy and reads the comparison off the distribution of households on that policy.

Table 2 reports the stationary means. C is nondurable consumption, D is the durable stock, DEX is expenditure on durables, TEX $= C +$ DEX is total expenditure, A is liquid assets, and ADJ is the fraction of households that buy in a period, in percent. The first block is the Extreme Value economy by itself. Search is shut off, the purchase price is degenerate at one, and the taste scale runs through $\sigma_\zeta \in \{0.1, 0.15, 0.2\}$. The other two blocks are one search-and-price economy, solved twice as a comparative static, rather than two further experiments with the other mechanism shut off. Effort is chosen from the price distribution before the draw, as in (77), and a bad draw is rejected. The second block holds the search cost at $S = 0.2$ and raises price dispersion through $\sigma_p \in \{0.08, 0.1, 0.12\}$. The third holds price dispersion at $\sigma_p = 0.1$ and raises the search cost through $S \in \{0.15, 0.2, 0.25\}$. The middle specification is the same household in both blocks,

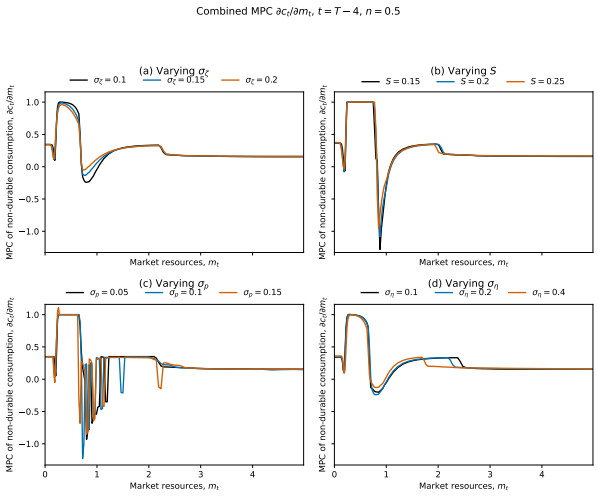


(a) Combined nondurable consumption $c_t(n_0, m_t)$.

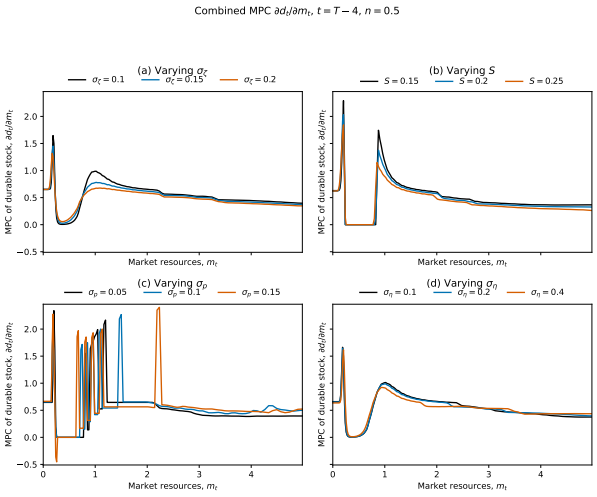


(b) Combined durable stock $d_t(n_0, m_t)$.

Figure 19: Combined policies at a fixed beginning-of-period stock n_0, one period before the terminal date. Panels (a)–(c) strengthen the taste scale, the search cost, and price dispersion one at a time, with the other two shut off. In each of those panels combined nondurable consumption shifts up and the combined durable stock shifts down. Panel (d) raises transitory income risk, which changes the continuation value and is outside the comparison of the three devices.



(a) MPC of nondurable consumption $\partial c_t/\partial m_t$.



(b) MPC of durable stock $\partial d_t/\partial m_t$.

Figure 20: Slopes of the combined rules in Figure 19 with respect to market resources. The spikes are the discrete switch between the keeper and the adjuster. Strengthening taste, search, or price dispersion replaces that switch with a gradual transition, so the spikes in both marginal propensities are damped. The income panel is the same experiment as panel (d) of the levels figure.

which is why that row appears twice. Search with the price shut off would be the one-sided limit of the previous subsection: effort is zero wherever keeping is better, and a higher search cost then need not reduce how often households adjust. The sign of S is therefore read inside the pair, where the option value of a rejectable draw is what effort is buying.

Table 2: Stationary aggregates. The first block is Extreme Value taste shocks alone, with search and price dispersion shut off. The second block raises price dispersion at a fixed search cost. The third raises the search cost at a fixed price dispersion. The middle row of the second block and the middle row of the third block are the same household.

		Steady state					
		C	D	DEX	TEX	A	ADJ (%)
σ_ζ	0.1	0.80	3.56	0.18	0.97	0.94	12.2
	0.15	0.79	3.27	0.16	0.95	0.71	18.6
	0.2	0.79	3.03	0.15	0.94	0.60	24.7
σ_p $(S = 0.2)$	0.08	0.81	4.28	0.21	1.03	1.39	7.7
	0.1	0.81	4.51	0.23	1.04	1.32	8.7
	0.12	0.82	4.87	0.24	1.06	1.27	9.9
S $(\sigma_p = 0.1)$	0.15	0.81	4.59	0.23	1.04	1.27	9.3
	0.2	0.81	4.51	0.23	1.04	1.32	8.7
	0.25	0.81	4.45	0.22	1.04	1.37	8.3

Steady state. A larger taste scale raises the adjustment probability because the logit puts more weight on the inferior branch at every state, including states in which adjusting is only slightly worse than keeping. Households replace the durable for reasons that do not track the value gap. They carry a smaller stock, and they hold fewer liquid assets against a discrete expenditure that now arrives more often. From $\sigma_\zeta = 0.1$ to 0.15 to 0.2, the adjustment probability rises from 12.2% to 18.6% to 24.7%, the durable stock falls from 3.56 to 3.27 to 3.03, and liquid assets fall from 0.93 to 0.71 to 0.60. Nondurable consumption barely moves, from 0.80 to 0.79. The entire comparative static is a change in how often an exogenous shock selects the adjuster.

Price dispersion, inside the pair, raises the adjustment probability for a different reason and moves the stock the other way. At $S = 0.2$, raising σ_p from 0.08 to 0.12 raises the adjustment probability from 7.7% to 9.9%, and liquid assets fall from 1.39 to 1.27, while the durable stock rises from 4.28 to 4.87. The price distribution is normalized to mean one before the household accepts or rejects the draw, so a wider distribution is a wider spread of offers rather than a cheaper posted price. Unusually good draws are accepted and unusually bad draws are rejected, and the purchases that go through install a larger stock. Taste and price dispersion can therefore produce the same rise in the adjustment frequency while moving durables in opposite directions. An adjustment frequency by itself cannot tell those two mechanisms apart.

The search cost moves the adjustment probability in the remaining direction. At $\sigma_p = 0.1$, raising S from 0.15 to 0.25 lowers the adjustment probability from 9.3% to 8.3%. The durable stock falls from 4.59 to 4.45, and liquid assets rise from 1.27 to 1.37. Effort has become more expensive, so fewer matches arrive, replacement is less frequent, and the resources that would have been spent on those purchases stay in the liquid asset. The movements are smaller than in the taste block because effort responds where the option value of a draw is positive, whereas the logit randomizes every state. Across the two blocks of the pair, the parameters push the adjustment probability in opposite directions: wider price dispersion raises it, and a higher search cost lowers it. A target for that frequency therefore leaves S and σ_p separately free, which is the sense in which the pair

can be disciplined by more than one moment. The quantitative step, matching the response of consumption and the response of the adjustment probability together, belongs to the fuller model and is not carried out here.

Labor Jacobians. The same stationary policies yield sequence-space Jacobians. A column of the Jacobian is the path of an aggregate after a one-time unanticipated change in an aggregate shock, so the date of the shock itself is the impact response and the later dates are the propagation. Figures 21–23 plot that object for a rise in mean labor income, for the three blocks of Table 2. The paths describe dynamics around the steady states just reported. They are a different question from the signs in the table. On this grid a larger taste scale raises the impact response of nondurable consumption slightly, from 0.089 at $\sigma_\zeta = 0.1$ to 0.095 at $\sigma_\zeta = 0.2$, so the comparative static in the taste scale is not a dampening of the impact propensity to consume. Holding the search cost at $S = 0.2$ and varying price dispersion leaves that impact response in a narrow range, between 0.076 and 0.079. Holding price dispersion at $\sigma_p = 0.1$ and varying the search cost moves it: the impact response is about zero at $S = 0.15$, and it is 0.077 and 0.060 at $S = 0.2$ and $S = 0.25$. Any later exercise that matches impulse responses has to confront both the impact and the shape of these paths, including the search-cost case in which the impact response changes with S.

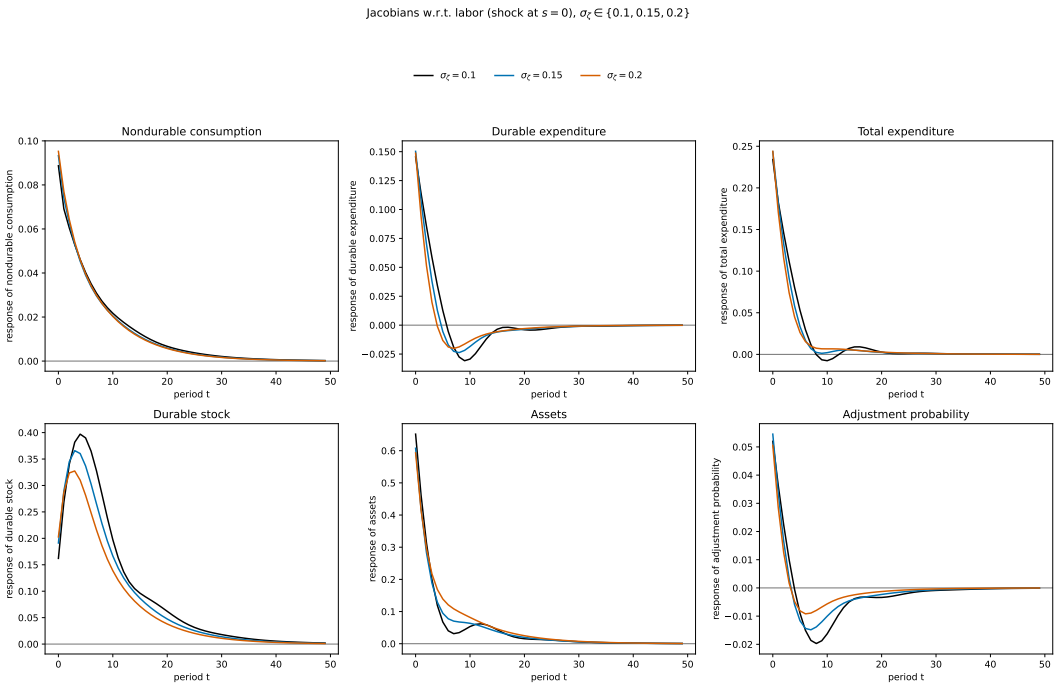


Figure 21: Jacobians with respect to a labor shock for $\sigma_\zeta \in \{0.1, 0.15, 0.2\}$. Search and price dispersion are shut off.

Kinks and Euler residuals Table 1 records cleanliness of the raw EGM cloud. The object that is stored and plotted is an interpolant of the refined cloud, so a zero deletion count does not by itself establish that the computed rule solves the problem that was posed. A policy mixed over dominated Euler roots can look as smooth as one whose envelope has nothing left to delete. I therefore pair the kink counts with off-grid Euler residuals in the sense of **?**. For an unconstrained continuous

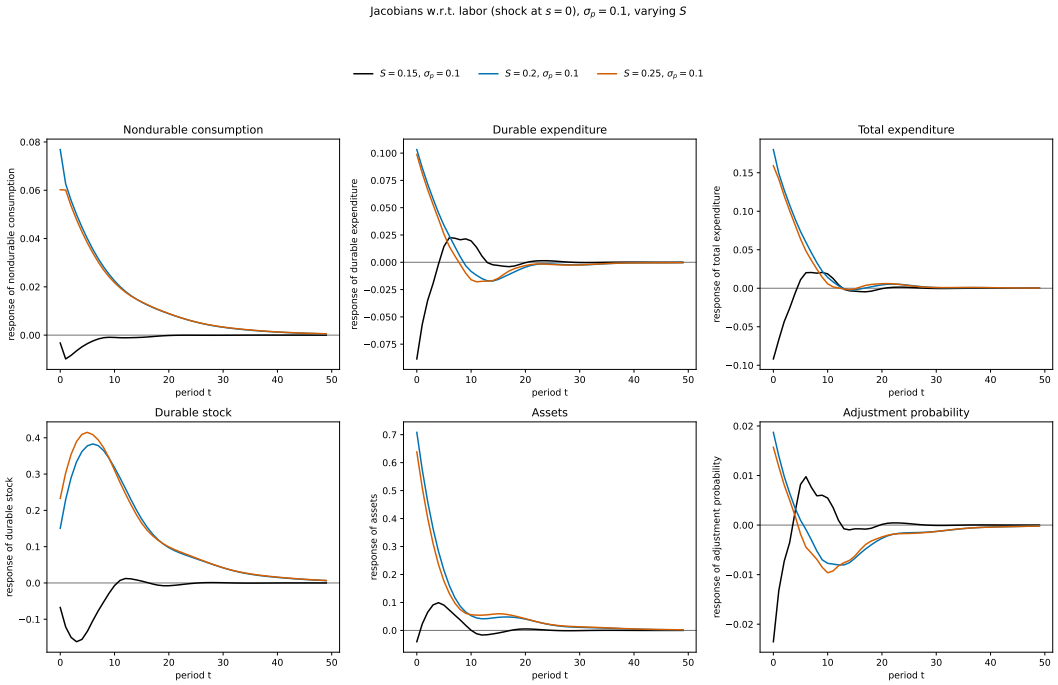


Figure 22: Labor Jacobians of the search-plus-price model at $\sigma_p = 0.1$, for $S \in \{0.15, 0.2, 0.25\}$.

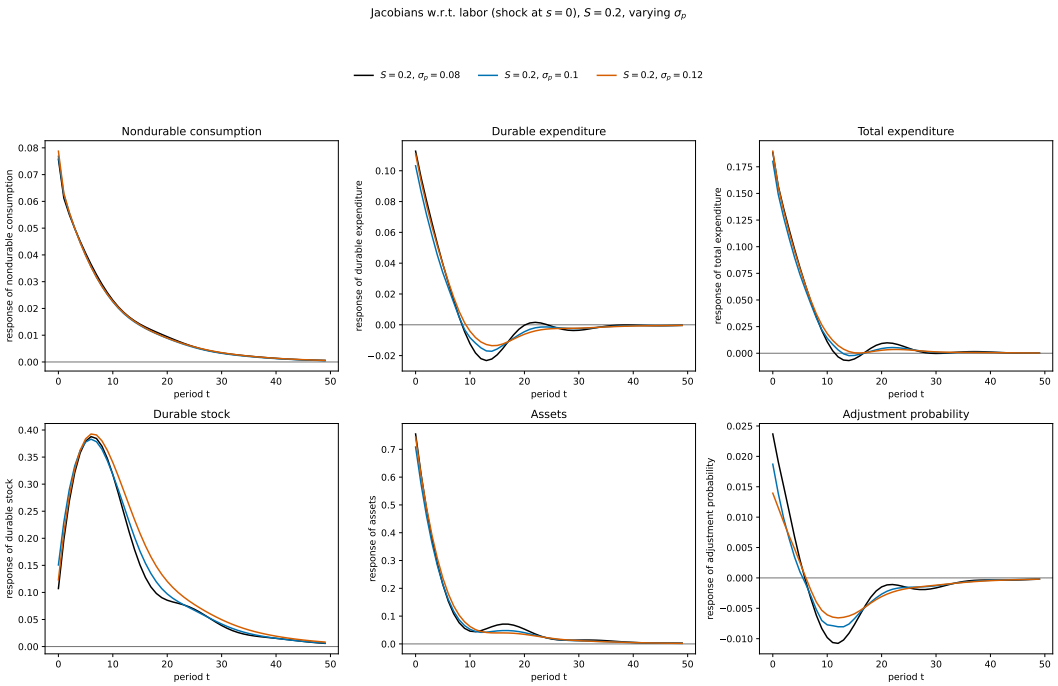


Figure 23: Labor Jacobians of the search-plus-price model at $S = 0.2$, for $\sigma_p \in \{0.08, 0.1, 0.12\}$.

choice c with implied Euler consumption c^{impl},

$$\text{EE}(c) \;=\; \log_{10}\!\left(\left|\frac{c - c^{\text{impl}}}{c}\right| + 10^{-16}\right). \tag{83}$$

More negative is more accurate: -4 is a 0.01% consumption error, -3 a 0.1% error. Residuals are branch-conditional—keepers on c_t^{keep}, adjusters on the durable first-order condition converted to the same consumption units (EE_d^{adj})—and are never evaluated at the mixed policy $\bar{c}_t$, which does not satisfy the Euler equation even when both branches do. The test mesh is off the solution grid; constrained points are dropped. I report the mean and the 95th percentile, because the mean can look acceptable while a thin set of points near the S,s band is not.

Table 3 places both diagnostics in the same period as Figures 19–20, now for every listed magnitude of σ_ζ, S, and σ_p. For the keeper the computational claim holds. Decreasing segments are zero at all three taste and search values, and $\text{EE}_c^{\text{keep}}$ improves with the smoother: the hard-max mean and 95th percentile are -4.19 and -2.50; at $\sigma_\zeta = 0.2$ they are -4.52 and -3.39, and at $S = 0.3$ they are -4.42 and -3.24. The smoother combined consumption rule is not an average over dominated keeper roots. Price dispersion likewise zeros the keeper count at all three σ_p, with only a small deterioration in the mean (from -4.11 at $\sigma_p = 0.05$ to -4.00 at 0.2).

For the adjuster the deletion count is not enough. Taste and search eliminate FUES removals at every listed magnitude, but the 95th percentile of EE_d^{adj} remains positive—0.45 to 0.69 as σ_ζ rises, 0.27 to 0.22 as S rises—while the mean stays around -3. Raising the smoother does not repair the tail. Price dispersion continues to remove points (20, 17, 34 as σ_p increases) and has a similar tail. Cleanliness of the raw FOC cloud is not the same as accuracy of the stored (c, d) interpolant, so FUES is not shown to be redundant by the deletion count alone. Taste, search, and prices are different economies and are not ranked by residual size; the residuals certify internal accuracy of each specification, not proximity to the frictionless target. Table 4 reports the same kink and Euler diagnostics for the infinite-horizon stationary policies of Section 3.7.

Table 3: Period-t kink counts and Euler residuals under the hard-max baseline and the same σ_ζ, S, and σ_p sweep as Figure 19. Residuals are mean and 95th percentile of (83) on an off-grid test mesh; constrained points excluded. The σ_p adjuster kink count sums over price nodes.

		Kinks		$\text{EE}_c^{\text{keep}}$		EE_d^{adj}	
		Decr. seg.	FUES rem.	Mean	P95	Mean	P95
Hard max		5	4	−4.19	−2.50	−3.49	−2.42
σ_ζ	0.05	0	0	−4.22	−3.23	−3.04	0.45
	0.1	0	0	−4.39	−3.32	−3.03	0.48
	0.2	0	0	−4.52	−3.39	−3.25	0.69
S	0.1	0	0	−4.23	−3.13	−3.06	0.27
	0.2	0	0	−4.34	−3.21	−3.22	0.22
	0.3	0	0	−4.42	−3.24	−3.20	0.22
σ_p	0.05	0	20	−4.11	−2.92	−2.96	0.53
	0.1	0	17	−4.04	−2.93	−2.97	0.52
	0.2	0	34	−4.00	−2.94	−2.87	0.65

Table 4: Infinite-horizon kink counts and Euler residuals. The taste, search-only, and price-only blocks vary one device with the other two shut off. The last two blocks vary σ_p at $S = 0.2$ and S at $\sigma_p = 0.1$, as in Table 2. Residuals are the mean and 95th percentile of (83) on an off-grid test mesh at the stationary policy; constrained points are excluded. The σ_p adjuster kink count sums over price nodes.

		Kinks		$\mathrm{EE}_c^{\mathrm{keep}}$		$\mathrm{EE}_d^{\mathrm{adj}}$	
		Decr. seg.	FUES rem.	Mean	P95	Mean	P95
σ_ζ	0.1	0	0	−5.08	−4.09	−1.49	1.35
	0.15	0	0	−5.20	−4.15	−1.92	0.94
	0.2	0	0	−5.29	−4.17	−2.17	0.66
S	0.15	1	0	−3.59	−2.23	−1.24	1.59
	0.2	0	0	−3.89	−2.64	−2.24	0.50
	0.25	0	0	−4.80	−3.60	−2.05	0.53
σ_p	0.05	0	102	−4.76	−3.31	−1.89	1.24
	0.1	0	59	−5.01	−3.83	−2.73	0.10
	0.15	1	22	−4.92	−3.87	−3.20	−1.00
σ_p $(S = 0.2)$	0.08	0	55	−4.88	−3.86	−2.12	1.21
	0.1	0	48	−4.89	−3.82	−2.08	1.31
	0.12	2	35	−4.92	−3.81	−2.25	0.99
S $(\sigma_p = 0.1)$	0.15	0	64	−4.62	−3.27	−2.11	1.08
	0.2	0	48	−4.89	−3.82	−2.08	1.31
	0.25	0	49	−4.89	−3.84	−2.20	1.10

4 Durable model (Full Model)

In this section, we show how the usage of endogeneous search and price dispersion is not only more microfounded than EV taste shocks, but since both mechanisms have opposite effects on the share of adjusters as well as asset and durable stock holdings for the steady state level, as well as for Jacobians, we can use it to match macroeconomic variables better than using just taste shocks.

4.1 Model description

The full model has the same baseline inregients as the basic model. However, to make it quantitatively more accurate, we introduce maintenance costs

5 Conclusion

Directed search and a rejectable price draw can match the smoothness that Extreme Value Type I taste shocks give to a discrete-continuous problem. The agent chooses effort in the unit interval at a convex utility cost. Effort raises the probability that an option to switch arrives. The draw is the content of that option: an account factor in the retirement model, and the purchase price of the new stock in the durable model. A bad draw is rejected. Search alone leaves one side of the threshold unsmoothed, because effort is zero when the switch is not strictly preferred. The price draw fills that side. The option value of the draw raises effort.

The retirement model illustrates this mechanism. The quantitative comparison is the keep-versus-adjust model of durables. Taste shocks, price dispersion, and search all smooth the (S, s) band. In steady state they move the adjustment probability in opposite directions. A larger taste

scale raises it, and durables and assets fall. Wider price dispersion raises it, because better draws are accepted; durables rise and assets fall. A higher search cost lowers it; durables fall and assets rise. Search and price dispersion can therefore be disciplined separately. Impulse responses of consumption and of the adjustment probability, at a common adjustment frequency, are the open exercise.

Search acts near the (S, s) thresholds, so the Euler equation still describes agents far from indifference. Optimal effort maps into an expected search duration. A search cost specified to scale with permanent income does not add a state to the normalized problem. An Extreme Value shock of the same scale for every household does not scale in that way.

A Retirement-model appendix

This appendix derives the objects used in Section 2. Section A.1 records the retiree block and the search problem without account-value risk. Section A.2 shuts search down and keeps the stochastic account realization. Section A.3 is the baseline of Section 2: directed search together with that realization. Throughout, labor income while working is the constant y, and the probability of receiving the retirement option equals search effort.

A.1 Baseline envelopes and Euler equations

This subsection records the blocks that do not depend on the account price ϕ_t: log utility, disutility of work δ, convex search cost $\theta(\varepsilon_t) = \sigma_s\big(\varepsilon_t + (1-\varepsilon_t)\log(1-\varepsilon_t)\big)$, option probabilities $P_1(\varepsilon_t) = \varepsilon_t$ and $P_2(\varepsilon_t) = 1-\varepsilon_t$, and $M_{t+1} = A_t R + y$ on the working branch. When there is no account-price draw, claiming delivers book resources one-for-one, so both branches use $A_t = M_t - c_t$. The baseline of Section 2 evaluates the retiree block below at claimed wealth $\phi_t M_t$.

Retiree. Retirement is self-absorbing. The Bellman equation is

$$V_t^{\text{Ret}}(M_t) = \max_{0\le c_t\le M_t} \left\{\log(c_t) + \beta V_{t+1}^{\text{Ret}}(A_t R)\right\}, \qquad A_t = M_t - c_t. \tag{84}$$

The Lagrangian $\mathscr{L} = \log(c_t) + \beta V_{t+1}^{\text{Ret}}(A_t R) + \lambda_t^{\text{Ret}}(M_t - c_t - A_t)$ yields

$$\lambda_t^{\text{Ret}} = \frac{1}{c_t} = \beta R\, \frac{\partial V_{t+1}^{\text{Ret}}(A_t R)}{\partial M_{t+1}}. \tag{85}$$

The envelope theorem gives $\partial V_t^{\text{Ret}}/\partial M_t = \lambda_t^{\text{Ret}} = 1/c_t^{\text{Ret}}$, which iterated forward is $\partial V_{t+1}^{\text{Ret}}/\partial M_{t+1} = 1/c_{t+1}^{\text{Ret}}$. Substituting produces the baseline retiree Euler equation

$$\frac{1}{c_t} = R\beta\, \frac{1}{c_{t+1}^{\text{Ret}}(A_t R)}, \tag{86}$$

or $c_t^{\text{Ret}} = \big(R\beta (c_{t+1}^{\text{Ret}}(A_t R))^{-1}\big)^{-1}$, as in the main text. The gothic retiree value is $\mathcal{V}_t^{\text{Ret}}(A_t) \equiv V_{t+1}^{\text{Ret}}(A_t R)$. This retiree problem is an input to Sections A.2–A.3: a claimed account $\tilde{M} = \phi M$ is just a different starting point for the same functions V_t^{Ret} and c_t^{Ret}.

Worker who must work. The worker's problem is

$$V_t^{W}(M_t) = \max_{\substack{0\le c_t\le M_t \\ 0\le \varepsilon_t\le 1}} \left\{\log(c_t) - \delta - \theta(\varepsilon_t) + \beta P_1(\varepsilon_t)\, \mathcal{V}_t^1(A_t) + \beta P_2(\varepsilon_t)\, \mathcal{V}_t^2(A_t)\right\}. \tag{87}$$

Splitting consumption and effort into sequential sub-periods, the effort first-order condition is

$$\theta'(\varepsilon_t^*) = \beta\big(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\big), \tag{88}$$

and since $\theta'(\varepsilon) = -\sigma_s \log(1-\varepsilon)$,

$$\varepsilon_t^*(A_t) = \max\left\{0,\ 1 - \exp\left(-\frac{\beta}{\sigma_s}\big(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\big)\right)\right\}. \tag{89}$$

The consumption first-order condition, evaluated at ε_t^*, is

$$\frac{1}{c_t} = \beta \left[P_1(\varepsilon_t^*) \frac{\partial \mathcal{V}_t^1}{\partial A_t} + P_2(\varepsilon_t^*) \frac{\partial \mathcal{V}_t^2}{\partial A_t} \right]. \tag{90}$$

Differentiating the gothic values at next period's certain labor income $M_{t+1} = A_t R + y$,

$$\frac{\partial \mathcal{V}_t^1}{\partial A_t} = R \frac{\partial V_{t+1}^{R,W}(M_{t+1})}{\partial M_{t+1}}, \tag{91}$$

$$\frac{\partial \mathcal{V}_t^2}{\partial A_t} = R \frac{\partial V_{t+1}^{W}(M_{t+1})}{\partial M_{t+1}}. \tag{92}$$

The envelope theorem applied to (87), using that ε_t^* is already optimal so its indirect effect vanishes, gives

$$\frac{\partial V_t^W(M_t)}{\partial M_t} = \frac{1}{c_t^W} = u'(c_t^W, 1, \varepsilon_t^*). \tag{93}$$

For the worker who has the option,

$$V_t^{R,W}(M_t) = \max \left\{ V_t^{\text{Ret}}(M_t),\ V_t^W(M_t) \right\}, \tag{94}$$

so

$$\frac{\partial V_{t+1}^{R,W}}{\partial M_{t+1}} = \mathbb{1}_{t+1}^0 \cdot \frac{1}{c_{t+1}^{\text{Ret}}} + \mathbb{1}_{t+1}^1 \cdot \frac{1}{c_{t+1}^W}, \tag{95}$$

with $\mathbb{1}_{t+1}^0 = 1$ iff $V_{t+1}^{\text{Ret}} \geq V_{t+1}^W$ and $\mathbb{1}_{t+1}^1 = 1 - \mathbb{1}_{t+1}^0$. Substituting (93)–(95) into (90) gives the worker's Euler equation for search without account dispersion. The consumption function with search and the account price is (36).

Terminal period. With $V_{T+1} \equiv 0$ and no borrowing, both branches set $A_T = 0$. The retiree consumes $c_T^{\text{Ret}} = M_T$ and $V_T^{\text{Ret}} = \log(M_T)$. The worker consumes $c_T^W = M_T$ and $V_T^W(M_T) = \log(M_T) - \delta$ (effort is irrelevant because there is no $T+1$). The option-contingent value is the hard maximum of the two.

A.2 Derivations: Stochastic retirement wealth

This subsection derives the envelope conditions, the Euler equations, the terminal-period solution, and the EGM algorithm for account dispersion with search shut down. The account draw arrives with certainty. The retiree's Bellman problem, first-order condition, and envelope of Section A.1 carry over unchanged: once a draw has been accepted, the agent is an ordinary retiree starting from $\tilde{M}_t = \phi_t M_t$. The worker's budget remains $A_t = M_t - c_t$. What changes is the discrete comparison and, through that comparison, the continuation that enters the worker's Euler equation.

Claiming maps into the baseline retiree. Define claimed wealth by (12), $\tilde{M}_t = \phi_t M_t$. The value of accepting the draw is the baseline retiree value evaluated at that point,

$$V_t^{\text{claim}}(M_t, \phi_t) \equiv V_t^{\text{Ret}}(\phi_t M_t), \tag{96}$$

with claiming-period consumption $c_t^{\text{claim}}(M_t, \phi_t) = c_t^{\text{Ret}}(\phi_t M_t)$. The retiree Euler equation remains (86); ϕ_t does not appear in it. After claiming, $A_t = \phi_t M_t - c_t^{\text{Ret}}(\phi_t M_t)$ and $M_{t+1} = A_t R$, with no further account draws.

Envelope conditions. By the chain rule, the derivative of claiming value with respect to *book* resources M_t, holding ϕ_t fixed, is

$$\frac{\partial V_t^{\text{claim}}(M_t, \phi_t)}{\partial M_t} = \phi_t \cdot \frac{\partial V_t^{\text{Ret}}(\tilde{M}_t)}{\partial \tilde{M}_t} = \phi_t \cdot \frac{1}{c_t^{\text{Ret}}(\phi_t M_t)}. \tag{97}$$

An extra unit of book wealth is worth ϕ_t units of retirement wealth, each of which has marginal value $1/c_t^{\text{Ret}}(\tilde{M}_t)$. Relative to the baseline envelope $\partial V_t^{\text{Ret}}/\partial M_t = 1/c_t^{\text{Ret}}$, the right-hand side is scaled by ϕ_t and the consumption function is evaluated at $\phi_t M_t$ rather than at M_t. This is the analog of the durable-price envelope in which p_t changes the adjuster's optimum: here ϕ_t changes the state at which that optimum is evaluated, not the first-order condition itself.

The worker never claims, so

$$\frac{\partial V_t^{W}(M_t)}{\partial M_t} = \frac{1}{c_t^{W}(M_t)}, \tag{98}$$

exactly as in the baseline.

Discrete choice with a stochastic account. Recall the realization-contingent discrete choice (15),

$$V_t(M_t, \phi_t) = \max\Big\{V_t^{\text{Ret}}(\phi_t M_t),\ V_t^{W}(M_t)\Big\}. \tag{99}$$

Define the state- and draw-contingent indicators $\mathbb{1}_t^0(M_t, \phi_t) = 1$ iff $V_t^{\text{Ret}}(\phi_t M_t) \geq V_t^{W}(M_t)$ and $\mathbb{1}_t^1(M_t, \phi_t) = 1 - \mathbb{1}_t^0(M_t, \phi_t)$. Applying the envelope theorem to the outer maximum,

$$\frac{\partial V_t(M_t, \phi_t)}{\partial M_t} = \mathbb{1}_t^0 \cdot \phi_t \cdot \frac{1}{c_t^{\text{Ret}}(\phi_t M_t)} + \mathbb{1}_t^1 \cdot \frac{1}{c_t^{W}(M_t)}. \tag{100}$$

The account price ϕ_t affects (100) both through the indicator (the extensive margin) and through the argument of c_t^{Ret} and the factor ϕ_t (the intensive margin and the mark-to-market of book wealth).

Integrating out the account draw. The objects that enter the time-$(t-1)$ Euler equations are the ϕ-integrated value and marginal-value functions

$$\bar{V}_t(M_t) \equiv \mathbb{E}_{\phi_t}\big[V_t(M_t, \phi_t)\big], \tag{101}$$

$$\bar{V}_{M,t}(M_t) \equiv \mathbb{E}_{\phi_t}\left[\mathbb{1}_t^0 \cdot \phi_t \cdot \frac{1}{c_t^{\text{Ret}}(\phi_t M_t)} + \mathbb{1}_t^1 \cdot \frac{1}{c_t^{W}(M_t)}\right]. \tag{102}$$

Because the draw arrives with certainty, the gothic value of a worker is $\mathcal{V}_{t-1}^1(A_{t-1}) = \bar{V}_t(RA_{t-1}+y)$. The gothic retiree value is unchanged: $\mathcal{V}_{t-1}^{\text{Ret}}(A_{t-1}) = V_t^{\text{Ret}}(RA_{t-1})$, because a household that has already claimed no longer faces ϕ. As $\sigma_\phi \to 0$, $\phi_t \to 1$ a.s., $\mathbb{1}_t^0 \to \mathbb{1}\{V_t^{\text{Ret}}(M_t) \geq V_t^{W}(M_t)\}$, and (102) collapses to the hard-maximum envelope (95).

The model's ex-ante (ϕ-weighted) consumption function follows the same logic. Because the household observes ϕ_t before choosing, the realized policy for a given draw is a hard 0/1 choice,

$$c_t(M_t, \phi_t) = \mathbb{1}_t^0\, c_t^{\text{Ret}}(\phi_t M_t) + \mathbb{1}_t^1\, c_t^{W}(M_t), \tag{103}$$

while the ex-ante policy reported for plotting is

$$\bar{c}_t(M_t) \equiv \mathbb{E}_{\phi_t}\big[c_t(M_t, \phi_t)\big] = \mathbb{E}_{\phi_t}\Big[\mathbb{1}_t^0\, c_t^{\text{Ret}}(\phi_t M_t) + \mathbb{1}_t^1\, c_t^{W}(M_t)\Big]. \tag{104}$$

Euler equations. The retiree's Euler equation is the baseline (86), evaluated at whatever claimed wealth the household accepted. For the worker, because the draw arrives with certainty,

$$\frac{1}{c_t^W(A_t)} = R\beta\,\bar{V}_{M,t+1}(M_{t+1}) = R\beta\int_{\phi_{t+1}}\left[\mathbb{1}^0_{t+1}\cdot\phi_{t+1}\cdot\frac{1}{c^{\text{Ret}}_{t+1}(\phi_{t+1}M_{t+1})} + \mathbb{1}^1_{t+1}\cdot\frac{1}{c^W_{t+1}(M_{t+1})}\right]dF(\phi_{t+1}), \tag{105}$$

where $\mathbb{1}^0_{t+1} = \mathbb{1}^0_{t+1}(M_{t+1},\phi_{t+1})$, $c^W_{t+1} = c^W_{t+1}(M_{t+1})$, and $M_{t+1} = A_tR + y$. Relative to the baseline worker Euler, the retiree-branch contribution to next period's marginal value is $\phi_{t+1}/c^{\text{Ret}}_{t+1}(\phi_{t+1}M_{t+1})$ rather than $1/c^{\text{Ret}}_{t+1}(M_{t+1})$, and the indicator is ϕ_{t+1}-contingent.

Terminal period. With $V_{T+1}\equiv 0$ and borrowing prohibited, both branches set $A_T = 0$. The worker's terminal solution is unaffected by the account draw: $c^W_T = M_T$ and $V^W_T = \log(M_T) - \delta$. For a claimant, $c^{\text{Ret}}_T(\phi_T M_T) = \phi_T M_T$ and

$$V^{\text{Ret}}_T(\phi_T M_T) = \log(\phi_T M_T) = \log(M_T) + \log(\phi_T). \tag{106}$$

The terminal discrete choice is $V_T(M_T,\phi_T) = \max\{V^{\text{Ret}}_T(\phi_T M_T),\, V^W_T(M_T)\}$. Substituting the closed forms, the household claims iff

$$\log(M_T) + \log(\phi_T) \geq \log(M_T) - \delta \quad\Longleftrightarrow\quad \phi_T \geq e^{-\delta}. \tag{107}$$

The terminal extensive margin is independent of wealth: both branches consume all resources, so M_T cancels and the comparison reduces to whether the log account realization exceeds the (negative of) work disutility. High ϕ_T favors claiming; low ϕ_T is denied. (In earlier periods wealth matters, because remaining at work preserves a claim on future labor income *and* on future draws.) The continuation handed to period $T-1$ is $\bar{V}_T(M_T) = \mathbb{E}_{\phi_T}[V_T(M_T,\phi_T)]$.

EGM algorithm with stochastic retirement wealth. Let $\vec{\phi} = \{\phi_1,\ldots,\phi_K\}$ be quadrature nodes for the log-normal account density with weights $\vec{\mu} = \{\mu_1,\ldots,\mu_K\}$. Because ϕ_t rescales the state at which the baseline retiree is evaluated rather than the retiree's first-order condition, the retiree is solved *once*—in contrast to the durable-price algorithm of Appendix B.4, which re-solves the adjuster at every p_k. The algorithm for one recursive period, starting from a given next-period solution $(c^{\text{Ret}}_{t+1},\, c^W_{t+1},\, V^{\text{Ret}}_{t+1},\, V^W_{t+1})$, proceeds as follows.

1. **Solve the retiree once,** by EGM on $A_t = \tilde{M}_t - c_t$ exactly as in Section A.1, yielding $c^{\text{Ret}}_t(\tilde{M})$ and $V^{\text{Ret}}_t(\tilde{M})$ on a grid that extends to at least $\phi_{\max}M_{\max}$ so that every claiming evaluation $\phi_k M$ lies in range.

2. **Solve the worker once.** The worker does not depend on current ϕ_t. For each A_i on a grid $\vec{A}_t$, compute $M_{t+1} = RA_i + y$ and, for each k,

$$V_{t+1}(M_{t+1},\phi_k) = \max\left\{V^{\text{Ret}}_{t+1}(\phi_k M_{t+1}),\; V^W_{t+1}(M_{t+1})\right\}, \tag{108}$$

$$\text{RHS}^{\text{pre}}(M_{t+1},\phi_k) = \mathbb{1}^0_{t+1}\cdot\phi_k\cdot\frac{1}{c^{\text{Ret}}_{t+1}(\phi_k M_{t+1})} + \mathbb{1}^1_{t+1}\cdot\frac{1}{c^W_{t+1}(M_{t+1})}. \tag{109}$$

Then

$$\mathcal{V}_t^1(A_i) = \sum_{k=1}^{K} \mu_k \, V_{t+1}(M_{t+1}, \phi_k), \tag{110}$$

$$\mathrm{RHS}(A_i) = R\beta \sum_{k=1}^{K} \mu_k \, \mathrm{RHS}^{\mathrm{pre}}(M_{t+1}, \phi_k), \tag{111}$$

$$c_t^W(A_i) = \mathrm{RHS}(A_i)^{-1}, \tag{112}$$

with endogenous grid $M_t^W(A_i) = A_i + c_t^W(A_i)$ and $V_t^W(M_t^W(A_i)) = \log(c_t^W(A_i)) - \delta + \beta \, \mathcal{V}_t^1(A_i)$. Interpolate onto a common M-grid.

3. **Combine on the M-grid, integrating over ϕ_k.** For each M and each node ϕ_k, form $\mathbb{1}^0(M, \phi_k)$ from $V_t^{\mathrm{Ret}}(\phi_k M) \geq V_t^W(M)$, then integrate as in (101)–(102) (and, for reporting, (104)). Feed $\bar{V}_t$ and $\bar{V}_{M,t}$ into the period-$(t-1)$ gothic step; feed V_t^{Ret} and c_t^{Ret} forward unchanged.

Computational cost scales with K only in the combine step and in the sum that builds $\mathcal{V}_t^1$; the retiree and worker EGM inversions remain single passes. As $\sigma_\phi \to 0$, $K = 1$ with $\phi_1 = 1$ recovers the hard-maximum baseline.

A.3 Derivations: Search and stochastic retirement wealth

This subsection derives the optimal search-effort rule, the combined envelope conditions, the Euler equations, the terminal-period solution, and the EGM algorithm for the baseline of Section 2. The retiree is the problem of Section A.1, evaluated at claimed wealth $\phi_t M_t$. The search cost is $\theta(\varepsilon_t) = \sigma_s\big(\varepsilon_t + (1-\varepsilon_t)\log(1-\varepsilon_t)\big)$, with derivative $\theta'(\varepsilon_t) = -\sigma_s \log(1-\varepsilon_t)$, and the worker's budget remains $A_t = M_t - c_t$. Search effort is the probability that a rejectable account draw arrives.

Claiming first-order conditions. Holding ϕ_t fixed and accepting the draw, the retiree's problem is the baseline problem at $\tilde{M}_t = \phi_t M_t$. The Euler equation is (86) and the chain-rule envelope is (97). No search takes place after claiming.

Optimal search effort. Search is chosen only on the working branch, after consumption, and current ϕ_t does not enter the worker's problem. Splitting (14) into sequential sub-periods as in the main text,

$$\mathcal{V}_t(A_t) = \max_{0 \leq \varepsilon_t \leq 1} \left\{ \log(c_t) - \delta - \theta(\varepsilon_t) + \beta \, \varepsilon_t \, \mathcal{V}_t^1(A_t) + \beta \, (1-\varepsilon_t) \, \mathcal{V}_t^2(A_t) \right\}, \tag{113}$$

with $\mathcal{V}_t^1$ and $\mathcal{V}_t^2$ given by (17)–(18). Because θ is strictly convex, the objective is strictly concave in ε_t. The interior first-order condition is $\theta'(\varepsilon_t^*) = \beta\big(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\big)$. With $\theta'(\varepsilon) = -\sigma_s \log(1-\varepsilon)$,

$$-\sigma_s \log(1-\varepsilon_t^*) = \beta\big(\mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t)\big), \tag{114}$$

which yields the closed form (24). The formula is the baseline (89); what changes with account dispersion is the content of $\mathcal{V}_t^1$. A mean-preserving spread in ϕ_{t+1} raises $\mathcal{V}_t^1$ because $V_{t+1}(M, \{R, W\}, \phi) = \max\{V_{t+1}^{\mathrm{Ret}}(\phi M), V_{t+1}^W(M)\}$ is convex in the claimed value and the household can reject bad draws. The excess value $\mathcal{V}_t^1 - \mathcal{V}_t^2$ therefore rises with σ_ϕ, and so does ε_t^*. As in the baseline, $\varepsilon_t^* = 0$ whenever $\mathcal{V}_t^1 \leq \mathcal{V}_t^2$, and $\varepsilon_t \leq 1$ never binds for finite σ_s because $\theta'(\varepsilon) \to \infty$ as $\varepsilon \to 1^-$. With $\phi_t \equiv 1$, (24) is search at book resources. Shutting search down, so that the draw arrives with certainty, is the account-only problem of Section A.2.

Envelope conditions and combined marginal values. Because ε_t^* is chosen optimally, the envelope theorem eliminates indirect effects through search effort when differentiating the worker's value. On the working branch,

$$\frac{\partial V_t^W(M_t)}{\partial M_t} = \frac{1}{c_t^W(M_t)}, \tag{115}$$

as in (98): current ϕ_t does not appear. For a worker who has the option, the envelope of the outer maximum (16) is

$$\frac{\partial V_t(M_t, \{R, W\}, \phi_t)}{\partial M_t} = \mathbb{1}_t^0(M_t, \phi_t) \cdot \phi_t \cdot \frac{1}{c_t^{\text{Ret}}(\phi_t M_t)} + \mathbb{1}_t^1(M_t, \phi_t) \cdot \frac{1}{c_t^W(M_t)}, \tag{116}$$

which is (100) (the discrete comparison is still a hard maximum; search operates one period ahead). Differentiating the gothic values (17)–(18) at $M_{t+1} = A_t R + y$,

$$\frac{\partial \mathcal{V}_t^1}{\partial A_t} = R \int_{\phi_{t+1}} \frac{\partial V_{t+1}(M_{t+1}, \{R, W\}, \phi_{t+1})}{\partial M_{t+1}} \, dF(\phi_{t+1}), \tag{117}$$

$$\frac{\partial \mathcal{V}_t^2}{\partial A_t} = R \frac{\partial V_{t+1}^W(M_{t+1})}{\partial M_{t+1}}, \tag{118}$$

with the integrands given by (116) and (115) one period ahead.

Integrating out account draws. The beginning-of-period objects that enter time-$(t-1)$ gothic values are

$$\bar{V}_t^{R,W}(M_t) \equiv \mathbb{E}_{\phi_t}\Big[\max\left\{V_t^{\text{Ret}}(\phi_t M_t),\ V_t^W(M_t)\right\}\Big], \tag{119}$$

$$\bar{V}_{M,t}^{R,W}(M_t) \equiv \mathbb{E}_{\phi_t}\left[\mathbb{1}_t^0 \cdot \phi_t \cdot \frac{1}{c_t^{\text{Ret}}(\phi_t M_t)} + \mathbb{1}_t^1 \cdot \frac{1}{c_t^W}\right]. \tag{120}$$

Equation (119) restates (19). The worker who must work is not integrated over current ϕ_t. The ex-ante (ϕ-weighted) consumption of an agent who has the option is

$$\bar{c}_t^{R,W}(M_t) \equiv \mathbb{E}_{\phi_t}\Big[\mathbb{1}_t^0 \, c_t^{\text{Ret}}(\phi_t M_t) + \mathbb{1}_t^1 \, c_t^W(M_t)\Big], \tag{121}$$

while the working-branch consumption $c_t^W(M_t)$ and effort $\varepsilon_t^*(A_t)$ are not functions of current ϕ_t. Because the household observes ϕ_t before the discrete choice, the *realized* policy for a given draw is the hard choice at that realization; (121) is the object reported for plotting.

Euler equations. The retiree's Euler equation remains (86). For the worker, the consumption first-order condition of (14) at ε_t^* is

$$\frac{1}{c_t^W} = \beta\left[\varepsilon_t^* \frac{\partial \mathcal{V}_t^1}{\partial A_t} + (1 - \varepsilon_t^*) \frac{\partial \mathcal{V}_t^2}{\partial A_t}\right]. \tag{122}$$

Substituting (117)–(118) and the envelopes (115)–(116) one period ahead,

$$\begin{aligned} \frac{1}{c_t^W} = R\beta\Bigg[\varepsilon_t^* \int_{\phi_{t+1}} \left(\mathbb{1}_{t+1}^0 \cdot \phi_{t+1} \cdot \frac{1}{c_{t+1}^{\text{Ret}}(\phi_{t+1} M_{t+1})} + \mathbb{1}_{t+1}^1 \cdot \frac{1}{c_{t+1}^W}\right) dF(\phi_{t+1}) \\ + (1 - \varepsilon_t^*) \frac{1}{c_{t+1}^W}\Bigg], \end{aligned} \tag{123}$$

where $M_{t+1} = A_t R + y$. Equation (123) is the worker Euler of Section 2. Relative to the account-only Euler (105), the continuation mixes over next period's arrival of the option through ε_t^*.

Terminal period. The within-branch terminal solutions are those of the account-only extension: $c_T^W = M_T$, $V_T^W = \log(M_T) - \delta$, and $V_T^{\text{Ret}}(\phi_T M_T) = \log(\phi_T M_T)$ from (106). Search effort is irrelevant at T because there is no $T+1$. The terminal discrete problem of a worker who has the option is the hard maximum (107), and $\bar{V}_T^{R,W}(M_T) = \mathbb{E}_{\phi_T}[V_T(M_T, \phi_T)]$ is the continuation handed to period $T-1$ through $\mathcal{V}_{T-1}^1$. A worker who does not have the option at T must work, so $\mathcal{V}_{T-1}^2 = V_T^W(RA_{T-1} + y)$.

EGM algorithm with search and stochastic retirement wealth. Let $\vec{\phi} = \{\phi_1, \ldots, \phi_K\}$ be quadrature nodes with weights $\vec{\mu}$. The algorithm nests the account step of Section A.2 and the search step in the main text.

1. **Solve the retiree once,** by EGM exactly as in Section A.1, on a $\tilde{M}$-grid that covers $\phi_{\max} M_{\max}$. The retiree does not search and does not depend on ϕ except through the starting state.

2. **Form gothic values on $\vec{A}_t$.** For each A_i, set $M_{t+1} = RA_i + y$. Interpolate the period-$(t+1)$ worker solution to obtain $V_{t+1}^W(M_{t+1})$ and $c_{t+1}^W(M_{t+1})$, and interpolate the retiree solution at each claimed wealth $\phi_k M_{t+1}$. Then

$$V_{t+1}(M_{t+1}, \phi_k) = \max\left\{V_{t+1}^{\text{Ret}}(\phi_k M_{t+1}),\ V_{t+1}^W(M_{t+1})\right\}, \tag{124}$$

$$\text{RHS}_1^{\text{pre}}(M_{t+1}, \phi_k) = \mathbb{1}_{t+1}^0 \cdot \phi_k \cdot \frac{1}{c_{t+1}^{\text{Ret}}(\phi_k M_{t+1})} + \mathbb{1}_{t+1}^1 \cdot \frac{1}{c_{t+1}^W(M_{t+1})}, \tag{125}$$

$$\text{RHS}_2(A_i) = \frac{1}{c_{t+1}^W(M_{t+1})}. \tag{126}$$

   Integrate

$$\mathcal{V}_t^1(A_i) = \sum_{k=1}^K \mu_k\, V_{t+1}(M_{t+1}, \phi_k), \tag{127}$$

$$\mathcal{V}_t^2(A_i) = V_{t+1}^W(M_{t+1}), \tag{128}$$

$$\text{RHS}_1(A_i) = \sum_{k=1}^K \mu_k\, \text{RHS}_1^{\text{pre}}(M_{t+1}, \phi_k). \tag{129}$$

3. **Effort and worker EGM.** Compute

$$\varepsilon_t(A_i) = \max\left\{0,\ 1 - \exp\big(-\beta(\mathcal{V}_t^1(A_i) - \mathcal{V}_t^2(A_i))/\sigma_s\big)\right\}, \tag{130}$$

$$\theta(\varepsilon_t(A_i)) = \sigma_s\big(\varepsilon_t(A_i) + (1 - \varepsilon_t(A_i))\log(1 - \varepsilon_t(A_i))\big), \tag{131}$$

$$\text{RHS}(A_i) = R\beta\big(\varepsilon_t(A_i)\,\text{RHS}_1(A_i) + (1 - \varepsilon_t(A_i))\,\text{RHS}_2(A_i)\big), \tag{132}$$

$$c_t^W(A_i) = \text{RHS}(A_i)^{-1}, \tag{133}$$

$$\mathbb{E}V_{t+1}(A_i) = \varepsilon_t(A_i)\,\mathcal{V}_t^1(A_i) + (1 - \varepsilon_t(A_i))\,\mathcal{V}_t^2(A_i), \tag{134}$$

$$V_t^W(M_t(A_i)) = \log(c_t^W(A_i)) - \delta - \theta(\varepsilon_t(A_i)) + \beta\,\mathbb{E}V_{t+1}(A_i), \tag{135}$$

$$M_t(A_i) = c_t^W(A_i) + A_i. \tag{136}$$

   Interpolate c_t^W and V_t^W onto the common M-grid. If there is not enough smoothness, apply an upper-envelope refinement as in the algorithm of Section 2.

4. **Combine on the M-grid, integrating over ϕ_k.** For each M and each ϕ_k, form $\mathbb{1}^0(M, \phi_k)$ from $V_t^{\mathrm{Ret}}(\phi_k M) \geq V_t^W(M)$ and integrate as in (119)–(120) (and, for reporting, (121)). Feed $\bar{V}_t^{R,W}$, $\bar{V}_{M,t}^{R,W}$, V_t^W, $\partial V_t^W / \partial M_t$, V_t^{Ret}, and c_t^{Ret} into the period-$(t-1)$ gothic step.

As in the earlier algorithms, inverse representations of the (marginal) value functions may be stored so that infinite conditional values do not propagate into interpolation error. Computational cost scales with K only through evaluations of the already-solved retiree at $\phi_k M$ and through the sum for $\mathcal{V}_t^1$; the retiree EGM itself is not repeated.

B Model Appendix

B.1 Derive Beginning-of-Period Marginal Value Functions

Problem in Bellman Form: The value function is the hard maximum over keeping $d_t = n_t$ and adjusting $d_t \neq n_t$:

$$v_t(m_t, n_t) = \max\left\{v_t^{keep}(m_t, n_t),\ v_t^{adj}(\ell_t)\right\} \tag{137}$$

where:

$$\ell_t = m_t + (1-\tau)n_t \tag{138}$$
$$a_t = \ell_t - c_t - d_t \geq 0, \tag{139}$$
$$m_{t+1} = \mathcal{R}a_t + \theta_{t+1}, \tag{140}$$
$$n_{t+1} = (1-\delta)\, d_t \tag{141}$$
$$u(c_t, d_t) = \frac{(c_t^\alpha d_t^{1-\alpha})^{1-\rho}}{1-\rho} \tag{142}$$

Where we have the Bellman for the keeper:

$$v_t^{keep}(m_t, n_t) = \max_{c_t} u(c_t, n_t) + \beta\mathbb{E}_t\Big(v_{t+1}(m_{t+1}, n_{t+1})\Big) \tag{143}$$

s.t.

$$a_t = m_t - c_t \geq 0, \tag{144}$$
$$m_{t+1} = \mathcal{R}a_t + \theta_{t+1}, \tag{145}$$
$$n_{t+1} = (1-\delta)\, d_t \tag{146}$$
$$d_t = n_t \tag{147}$$
$$u(c_t, d_t) = \frac{(c_t^\alpha d_t^{1-\alpha})^{1-\rho}}{1-\rho} \tag{148}$$

And for the adjuster:

$$v_t^{adj}(\ell_t) = \max_{c_t, d_t} u(c_t, d_t) + \beta\mathbb{E}_t\Big(v_{t+1}(m_{t+1}, n_{t+1})\Big) \tag{149}$$

s.t.

$$\ell_t = m_t + (1-\tau)n_t \tag{150}$$
$$a_t = \ell_t - c_t - d_t \geq 0, \tag{151}$$
$$m_{t+1} = \mathcal{R}a_t + \theta_{t+1}, \tag{152}$$
$$n_{t+1} = (1-\delta)\, d_t \tag{153}$$
$$u(c_t, d_t) = \frac{(c_t^\alpha d_t^{1-\alpha})^{1-\rho}}{1-\rho} \tag{154}$$

Step 1 (chain rule / envelope): The adjuster value is a function of ℓ_t,

$$v_t^{adj} = v_t^{adj}(\ell_t), \qquad \ell_t = m_t + (1-\tau)n_t.$$

By the chain rule

$$\frac{\partial v_t^{adj}(\ell_t)}{\partial m_t} = \frac{\partial v_t^{adj}(\ell_t)}{\partial \ell_t} \cdot \frac{\partial \ell_t}{\partial m_t}, \tag{155}$$
$$\frac{\partial v_t^{adj}(\ell_t)}{\partial n_t} = \frac{\partial v_t^{adj}(\ell_t)}{\partial \ell_t} \cdot \frac{\partial \ell_t}{\partial n_t}. \tag{156}$$

Step 2 (derivatives of ℓ_t):

$$\frac{\partial \ell_t}{\partial m_t} = 1, \qquad \frac{\partial \ell_t}{\partial n_t} = 1 - \tau.$$

Step 3 (envelope for $\partial v^{adj}/\partial \ell_t$): The adjuster problem is

$$\max_{c_t, d_t} \; u(c_t, d_t) + \beta \mathbb{E}_t\big[v_{t+1}(m_{t+1}, n_{t+1})\big] \quad \text{s.t.} \quad a_t = \ell_t - c_t - d_t,$$

with the stated laws of motion. Because ℓ_t enters only through the linear resource constraint $a_t = \ell_t - c_t - d_t$, the envelope theorem implies that

$$\frac{\partial v_t^{adj}(\ell_t)}{\partial \ell_t} = \lambda_t^{adj},$$

where λ_t^{adj} is the Lagrange multiplier on the resource constraint (i.e. the marginal value of an extra unit of ℓ_t).

Step 4 (relation between λ_t^{adj} and marginal utility): The FOC w.r.t. c_t (ignoring binding additional constraints) sets

$$\lambda_t^{adj} = u_c(c_t^*, d_t^*).$$

With the specified utility,

$$u_c(c, d) = (c^\alpha d^{1-\alpha})^{-\rho}\, \alpha c^{\alpha-1} d^{1-\alpha} = \alpha\, c^{\alpha(1-\rho)-1}\, d^{(1-\alpha)(1-\rho)}.$$

If the borrowing constraint binds and has multiplier μ_t, then λ_t^{adj} equals u_c plus the relevant shadow price (i.e. $\lambda_t^{adj} = u_c +$ extra multipliers).

Step 5 (combine chain rule and envelope): Using Steps 1–4,

$$\frac{\partial v_t^{adj}}{\partial m_t} = \lambda_t^{adj} = u_c(c_t^*, d_t^*) \tag{157}$$

and

$$\frac{\partial v_t^{adj}}{\partial n_t} = (1 - \tau)\, \lambda_t^{adj} = (1 - \tau)\, u_c(c_t^*, d_t^*) \tag{158}$$

evaluated at the adjuster optimum (c_t^*, d_t^*) (and with λ_t^{adj} augmented by constraint multipliers if applicable).

Keeper: stepwise derivation (same style) In the keeper regime we have $d_t = n_t$ (fixed) and the household chooses only c_t. The state is (m_t, n_t) and

$$a_t = m_t - c_t, \qquad m_{t+1} = \mathcal{R} a_t + \theta_{t+1}, \qquad n_{t+1} = (1 - \delta)\, n_t.$$

Keeper Step 1 (envelope for m_t): Since m_t enters linearly via $a_t = m_t - c_t$, the envelope theorem gives

$$\frac{\partial v_t^{keep}(m_t, n_t)}{\partial m_t} = \lambda_t^{keep} = u_c(c_t^*, n_t) \tag{159}$$

where again λ_t^{keep} is the Lagrange multiplier on the budget and u_c is as above (add constraint multipliers if the borrowing constraint binds).

Keeper Step 2 (envelope for n_t): The durable stock n_t affects the keeper value through two channels:

1. Directly in current utility because $d_t = n_t$: marginal contribution $u_d(c_t^*, n_t)$.

2. Indirectly via $n_{t+1} = (1-\delta)\, n_t$, which affects $v_{n,t+1}$.

Applying the envelope theorem and the chain rule,

$$\frac{\partial v_t^{keep}}{\partial n_t} = u_d(c_t^*, n_t) + \beta\, \mathbb{E}_t \left[v_{n,t+1} \cdot \frac{\partial n_{t+1}}{\partial n_t} \right] \tag{160}$$

$$= u_d(c_t^*, n_t) + \beta\, \mathbb{E}_t \Big[v_{n,t+1}\, (1-\delta) \Big]. \tag{161}$$

Thus

$$\frac{\partial v_t^{keep}(m_t, n_t)}{\partial n_t} = u_d(c_t^*, n_t) + \beta(1-\delta)\, \mathbb{E}_t \big[v_{n,t+1} \big] \tag{162}$$

with

$$u_d(c, d) = (1-\alpha)\, c^{\alpha(1-\rho)}\, d^{(1-\alpha)(1-\rho)-1}.$$

Combined expression with adjust / keep indicator Let $\mathcal{I}_t$ be the indicator for *adjusting* at time t ($\mathcal{I}_t = 1$ if adjust, 0 if keep). Then the beginning-of-period partials of the value function can be summarized as

$$\frac{\partial v_t(m_t, n_t)}{\partial m_t} = \mathcal{I}_t \frac{\partial v_t^{adj}}{\partial m_t} + (1 - \mathcal{I}_t) \frac{\partial v_t^{keep}}{\partial m_t} \tag{163}$$

$$= \mathcal{I}_t \lambda_t^{adj} + (1 - \mathcal{I}_t) \lambda_t^{keep} = \frac{\partial u(c_t, d_t)}{\partial c_t} \tag{164}$$

(the final equality holds when λ's equal marginal utility of consumption; add constraint multipliers if necessary).

Similarly,

$$\frac{\partial v_t(m_t, n_t)}{\partial n_t} = \mathcal{I}_t \frac{\partial v_t^{adj}}{\partial n_t} + (1 - \mathcal{I}_t) \frac{\partial v_t^{keep}}{\partial n_t} \tag{165}$$

$$= \mathcal{I}_t (1-\tau) \lambda_t^{adj} + (1 - \mathcal{I}_t) \left\{ u_d(c_t, n_t) + \beta(1-\delta)\, \mathbb{E}_t \big[v_{n,t+1} \big] \right\}. \tag{166}$$

If $\lambda_t^{adj} = u_c(c_t^*, d_t^*)$ and $\lambda_t^{keep} = u_c(c_t^*, n_t)$, you may substitute the explicit formulas for u_c and u_d given above to obtain fully expanded expressions.

Last period This terminal-period solution serves as the initial condition for the backward induction used to find the infinite-horizon fixed point; it is not a true end-of-life period but rather a starting guess from which the backward recursion is iterated until convergence. We set the continuation value to zero, $v_{t+1} \equiv 0$. Since there is no future utility cost of debt and borrowing is prohibited ($LTV = 0$, $\underline{a} = 0$), both the keeper and the adjuster optimally set end-of-period assets to $a = 0$. For the keeper ($d_t = n_t$), this gives $c^{keep} = m_t$.

For the adjuster problem, the binding borrowing constraint is simply:

$$a = 0 \tag{167}$$

Hence, the budget constraint:

$$\ell = c + d + a = c + d \tag{168}$$

Meaning we solve:

$$\mathscr{L} = u(c,d) + \lambda\big(\ell - c - d\big) \tag{169}$$

$$\frac{\partial \mathscr{L}}{\partial c} = u_c(c,d) = \lambda \tag{170}$$

$$\frac{\partial \mathscr{L}}{\partial d} = u_d(c,d) = \lambda \tag{171}$$

Setting both first-order conditions equal and using CRRA utility $u(c,d) = \dfrac{(c^\alpha d^{1-\alpha})^{1-\rho}}{1-\rho}$:

$$u_c(c,d) = u_d(c,d) \tag{172}$$

$$\frac{\alpha(c^\alpha d^{1-\alpha})^{-\rho}}{c} = \frac{(1-\alpha)(c^\alpha d^{1-\alpha})^{-\rho}}{d} \tag{173}$$

$$\frac{c}{d} = \frac{\alpha}{1-\alpha} \tag{174}$$

$$c = d\,\frac{\alpha}{1-\alpha} \tag{175}$$

Plugging this into the budget constraint:

$$\ell = c + d = d\,\frac{\alpha}{1-\alpha} + d = d\,\frac{1}{1-\alpha} \tag{176}$$

Therefore, we have the solutions:

$$d^* = (1-\alpha)\,\ell \tag{177}$$

$$c^* = \alpha\,\ell \tag{178}$$

Note that $c^* = \alpha\ell$ is a consequence of the Cobb-Douglas aggregator, which fixes budget-share α for nondurables.

Calculate End-Of-Period Functions The standard approach is to discretize the stochastic variables with a limited amount of discrete points such as in Tauchen and others. The arising problem with a discrete choice problem is that policy functions become discontinuous and the value function is non-smooth. As a result, evaluating the next period's (marginal) value function on a few points can lead to *phantom* secondary non-concavities; that is large point masses are distributed at a few points instead of using all points in the surrounding area. The most straightforward solution is to increase the number of points with which a distribution is discretized.

Instead, I use a method of *transition matrix* or *lotteries* often used to receive model solutions[2]. That is, splitting up the transitory shock θ such that the difference between beginning of period

$$b_t = \mathcal{R}\,a_{t-1} \tag{179}$$

$$m_t - b_t = \theta_t \tag{180}$$

The algorithm is as follows:

[2]This method was developed by Matthew White for this particular purpose

1. Pre-solution step
   (a) Specify a fairly dense grid of beginning-of-period market resources and bank balances m and b
   (b) Calculate the pdf of $\theta : f_{\theta t}(m - b)$ and create a matrix with probabilities with b values by row and m values by column.
   (c) Create a Markov transition matrix by taking the row-wise sum and divide the row by it.
2. Solution step
   (a) Evaluate $v_{t+1}(m_{t+1}, n_{t+1})$ on the cross product of the dense m-grid and some grid of n.
   (b) Pre-multiply the Markov matrix by the created matrix. The resulting matrix is b row-wise and n column wise.
   (c) Construct a linear interpolation function.
   (d) Create the end-of-period (marginal) functions over the end-of-period variables a (assets) and d (durable stock) by evaluating the interim function denoted by $\mathfrak{v}(d_t, a_t)$, $\mathfrak{v}_c^{'}(d_t, a_t)$, and $\mathfrak{v}_d^{'}(d_t, a_t)$

Solve Keeper The keeper solution specifies that the beginning-of-period durable stock n_t is equal to the end-of-period durable stock d_t. Therefore, we only have to calculate optimal level of c_t and a_t which we can do by EGM and an upper-envelope algorithm to deal with non-concavities.

1. Loop over each point of the end-of-period durable stock d_t
2. Invert the end-of-period marginal value function to obtain c_t. From $u_c(c_t, n_t) = \mathfrak{v}_c^{'}(n_t, a_t)$ with $d_t = n_t$ fixed:

$$c_t = \left[\frac{\mathfrak{v}_c^{'}(n_t, a_t)}{\alpha \, n_t^{(1-\alpha)(1-\rho)}} \right]^{\frac{1}{\alpha(1-\rho)-1}}$$

3. Calculate the endogenous beginning-of-period market resources using the budget constraint

$$m_t = c_t + a_t$$

4. Use the upper envelope algorithm
5. Construct a function over n and m

Solve Adjuster Solving the adjuster problem follows the paper by **?**:

1. Loop over each end of period durable stock d_t
2. Calculate c_t over a dense grid of a_t using the EGM step. Setting the first-order condition $u_c(c_t, d_t) = \mathfrak{v}_c^{'}(d_t, a_t)$ and inverting with respect to c_t (holding d_t fixed) yields

$$c_t = \left[\frac{\mathfrak{v}_c^{'}(d_t, a_t)}{\alpha \, d_t^{(1-\alpha)(1-\rho)}} \right]^{\frac{1}{\alpha(1-\rho)-1}} \tag{181}$$

3. Find all roots of a_t which obey the first order conditions

$$0 = u_d^{'}(c,d) - \mathfrak{v}_c^{'}(d_t,a_t) + \mathfrak{v}_d^{'}(d_t,a_t) \tag{182}$$

4. Calculate ℓ_t with the budget constraint

$$\ell_t = c_t + d_t + a_t \tag{183}$$

5. Calculate the value:

$$v_t = u(c_t,d_t) + \mathfrak{v}(d_t,a_t) \tag{184}$$

6. Apply FUES

7. Create functions over (ℓ)

Discrete Choice To solve the discrete choice between adjusting and keeping, take the hard maximum of the two conditional value functions:

$$v_t(m_t,n_t) = \max\Big\{v_t^{keep}(m_t,n_t),\ v_t^{adj}\big(m_t + (1-\tau)n_t\big)\Big\}. \tag{185}$$

Define the indicator $\mathcal{I}_t = 1$ if $v_t^{adj}(\ell_t) \geq v_t^{keep}(m_t,n_t)$ and $\mathcal{I}_t = 0$ otherwise. The combined policy functions are then

$$c_t = (1-\mathcal{I}_t)\, c_t^{keep} + \mathcal{I}_t\, c_t^{adj}, \qquad d_t = (1-\mathcal{I}_t)\, n_t + \mathcal{I}_t\, d_t^{adj}, \tag{186}$$

and the combined (marginal) value functions follow equations (164)–(166). When computing (marginal) value functions, note that a positive probability of an infinite conditional value implies that the combined function also takes the value infinity. To facilitate the evaluation of functions involving infinities, I store and assess their inverse representations, and subsequently invert them again when the arrays are used.

B.2 Derivations: Extreme Value Type I Taste Shocks (Section 3.4)

This subsection derives the envelope conditions, the Euler equations, and the EGM algorithm for the EV Type I taste-shock extension of Section 3.4. Throughout, $v_t^{keep}(m_t,n_t)$ and $v_t^{adj}(\ell_t)$ are exactly as defined in Section B, so the within-branch first-order conditions (54)–(55), the Step 1–5 envelope derivation for the adjuster (equations (157)–(158)), and the keeper envelope derivation (equations (159)–(162)) carry over unchanged. What changes is how the two branches are combined and, consequently, what continuation value enters the Euler equations one period back.

Log-sum inclusive value and logit probabilities Recall the EV Type I discrete choice from equations (59)–(61). The inclusive (ex-ante) value is the closed-form log-sum

$$v_t(m_t,n_t) = \sigma_\zeta \log\Big(\exp\big(v_t^{keep}(m_t,n_t)/\sigma_\zeta\big) + \exp\big(v_t^{adj}(\ell_t)/\sigma_\zeta\big)\Big), \tag{187}$$

with logit choice probabilities $P_t^{adj}(m_t,n_t)$ and $P_t^{keep} = 1 - P_t^{adj}$ as in (61). The taste shocks do not enter the inner maximizations defining v_t^{keep} and $v_t^{adj}(\ell_t)$, so those maximizations—and hence $c_t^{keep,*}$, $c_t^{adj,*}$, and $d_t^{adj,*}$—are independent of $(\zeta_t^{keep}, \zeta_t^{adj})$. Differentiating the log-sum (187) and

using the envelope results $\partial v_t^{keep}/\partial m_t = u_c(c_t^{keep,*}, n_t)$ and $\partial v_t^{adj}/\partial m_t = u_c(c_t^{adj,*}, d_t^{adj,*})$ from (159) and (157) yields the probability-weighted combined marginal values

$$\frac{\partial v_t(m_t, n_t)}{\partial m_t} = P_t^{keep}\, u_c(c_t^{keep,*}, n_t) + P_t^{adj}\, u_c(c_t^{adj,*}, d_t^{adj,*}), \tag{188}$$

$$\frac{\partial v_t(m_t, n_t)}{\partial n_t} = P_t^{keep}\Big[u_d(c_t^{keep,*}, n_t) + \beta(1-\delta)\,\mathbb{E}_t\big[v_{n,t+1}\big]\Big] + P_t^{adj}\,(1-\tau)\,u_c(c_t^{adj,*}, d_t^{adj,*}), \tag{189}$$

where the n-derivative of the adjuster branch uses $\partial \ell_t/\partial n_t = (1-\tau)$ as in (158). Unlike the multiplicative mean-one log-normal taste shock used in the retirement application, there is no additional scaling factor on the adjuster's marginal contribution: the EV shocks enter only through the logit weights P_t^{keep} and P_t^{adj}. As $\sigma_\zeta \to 0$, $P_t^{adj} \to \mathcal{I}_t$ and (188)–(189) collapse to the hard-maximum combined envelopes (164)–(166).

Expected policies The ex-ante (probability-weighted) policy functions are

$$c_t(m_t, n_t) = P_t^{keep}\, c_t^{keep,*}(m_t, n_t) + P_t^{adj}\, c_t^{adj,*}(\ell_t), \tag{190}$$

$$d_t(m_t, n_t) = P_t^{keep}\, n_t + P_t^{adj}\, d_t^{adj,*}(\ell_t), \tag{191}$$

which restates (62)–(63). Because the inclusive value (187) already integrates out the Gumbel shocks in closed form, no quadrature over taste nodes is required: v_t, $\partial v_t/\partial m_t$, $\partial v_t/\partial n_t$, c_t, and d_t are evaluated directly from the two conditional branches and the logit probabilities.

Euler equations The within-branch first-order conditions are unchanged in form, but the continuation value is now the EV-smoothed inclusive value v_{t+1} from (187) rather than the hard maximum:

$$u_c(c_t^{adj,*}, d_t^{adj,*}) \geq \beta\mathcal{R}\,\mathbb{E}_t\big[v_{m,t+1}\big], \tag{192}$$

$$u_d(c_t^{adj,*}, d_t^{adj,*}) \geq \beta\mathcal{R}\,\mathbb{E}_t\big[v_{m,t+1}\big] - \beta(1-\delta)\,\mathbb{E}_t\big[v_{n,t+1}\big], \tag{193}$$

and analogously for the keeper with $d_t = n_t$ fixed, where $v_{m,t+1} = \partial v_{t+1}/\partial m_{t+1}$ and $v_{n,t+1} = \partial v_{t+1}/\partial n_{t+1}$ are given by (188)–(189) one period ahead. Writing $\mathbb{E}_t[\cdot]$ as the expectation over θ_{t+1} and substituting (188) makes explicit that the right-hand side averages next period's income *and* the logit-weighted keep/adjust margins, e.g.

$$u_c(c_t^{adj,*}, d_t^{adj,*}) = \beta\mathcal{R}\,\mathbb{E}_{\theta_{t+1}}\Big[P_{t+1}^{keep}\, u_c(c_{t+1}^{keep,*}, n_{t+1}) + P_{t+1}^{adj}\, u_c(c_{t+1}^{adj,*}, d_{t+1}^{adj,*})\Big] \tag{194}$$

when the no-borrowing constraint does not bind.

Terminal period The within-branch terminal solutions are unaffected by the taste shocks, since $(\zeta_T^{keep}, \zeta_T^{adj})$ do not enter the inner maximizations: $c_T^{keep} = m_T$ for the keeper, and $c_T^{adj*} = \alpha\ell_T$, $d_T^{adj*} = (1-\alpha)\ell_T$ for the adjuster, exactly as derived in the "Last period" paragraph of Section B. Only the discrete resolution changes: the terminal inclusive value is the log-sum (187) at T, and the terminal expected policies follow (190)–(191).

EGM algorithm with EV Type I taste shocks No taste-shock quadrature grid is needed. In addition to the income-shock integration already used to build the end-of-period functions $\mathfrak{v}(d_t, a_t)$, $\mathfrak{v}'_c(d_t, a_t)$, and $\mathfrak{v}'_d(d_t, a_t)$ (Section B, "Calculate End-Of-Period Functions"), the algorithm proceeds as follows:

1. **Solve the within-branch problems unchanged.** Solve the keeper by EGM and upper envelope, and the adjuster by root-finding and FUES, exactly as in the "Solve Keeper" and "Solve Adjuster" paragraphs above. This yields $c_t^{keep}(m_t, n_t)$, $v_t^{keep}(m_t, n_t)$, $c_t^{adj}(\ell_t)$, $d_t^{adj}(\ell_t)$, and $v_t^{adj}(\ell_t)$.

2. **Combine with the log-sum / logit formulas on a grid of (m_t, n_t).** For each grid point, evaluate $v_t^{keep}(m_t, n_t)$ and $v_t^{adj}(\ell_t)$, form P_t^{adj} and P_t^{keep} from (61), and set

$$v_t(m_t, n_t) = \sigma_\zeta \log\Big(\exp\big(v_t^{keep}/\sigma_\zeta\big) + \exp\big(v_t^{adj}/\sigma_\zeta\big)\Big), \tag{195}$$

$$v_{m,t}(m_t, n_t) = P_t^{keep}\, u_c(c_t^{keep,*}, n_t) + P_t^{adj}\, u_c(c_t^{adj,*}, d_t^{adj,*}), \tag{196}$$

$$v_{n,t}(m_t, n_t) = P_t^{keep}\, v_{n,t}^{keep}(m_t, n_t) + P_t^{adj}\,(1-\tau)\, u_c(c_t^{adj,*}, d_t^{adj,*}), \tag{197}$$

   together with the expected policies (190)–(191). Numerically, evaluate the log-sum in the numerically stable form $\max\{v^{keep}, v^{adj}\} + \sigma_\zeta \log\big(\sum_j \exp((v^j - \max)/\sigma_\zeta)\big)$ (as in `calcLogSumChoiceProbs` with `smoothtype='type1'`).

3. **Feed v_t, $v_{m,t}$, $v_{n,t}$ into the period-$(t-1)$ end-of-period step,** in place of the hard-maximum objects used in the baseline model, so that the transition-matrix income integration described above operates on the EV-smoothed functions.

As $\sigma_\zeta \to 0$, step 2 collapses to the hard-maximum combine of the baseline algorithm. For $\sigma_\zeta > 0$, the inclusive value and probability-weighted policies are smooth in (m_t, n_t), which is precisely what reduces the propagation of kinks through the backward recursion.

B.3 Derivations: Search Frictions (Section 3.5)

This subsection derives the optimal search-effort rule and the smooth combined (marginal) value functions for the search extension of Section 3.5, and states the resulting EGM algorithm. As in the taste-shock extension, the within-branch Bellman problems for $v_t^{keep}(m_t, n_t)$ and $v_t^{adj}(\ell_t)$, and hence the first-order conditions (54)–(55) and the envelope results (157)–(162), are unchanged; only the way the two branches are combined, and the continuation value that results, is different.

Optimal search effort Recall the search problem from equation (80),

$$v_t(m_t, n_t) = \max_{\varepsilon_t \in [0,1]} \underbrace{\Big[(1-\varepsilon_t)\, v_t^{keep}(m_t, n_t) + \varepsilon_t\, v_t^{adj}(\ell_t) - \sigma(\varepsilon_t)\Big]}_{\equiv\, g(\varepsilon_t;\, m_t, n_t)}. \tag{198}$$

The cost function's derivative is

$$\begin{aligned}
\sigma(\varepsilon_t) &= S\big[\varepsilon_t + (1-\varepsilon_t)\log(1-\varepsilon_t)\big] \\
\sigma'(\varepsilon_t) &= S\Big[1 + \big(-\log(1-\varepsilon_t)\big) + (1-\varepsilon_t)\cdot\Big(-\frac{1}{1-\varepsilon_t}\Big)\cdot(-1)\Big] = S\big[1 - \log(1-\varepsilon_t) - 1\big] = -S\log(1-\varepsilon_t),
\end{aligned} \tag{199}$$

and $\sigma''(\varepsilon_t) = S/(1-\varepsilon_t) > 0$ for $\varepsilon_t \in [0,1)$, so σ is strictly convex and $g(\cdot; m_t, n_t)$ is strictly concave in ε_t (linear minus convex). The first-order condition for an interior optimum is

$$\frac{\partial g}{\partial \varepsilon_t} = v_t^{adj}(\ell_t) - v_t^{keep}(m_t, n_t) - \sigma'(\varepsilon_t^*) = 0 \quad\Longleftrightarrow\quad v_t^{adj}(\ell_t) - v_t^{keep}(m_t, n_t) = -S\log(1-\varepsilon_t^*), \tag{200}$$

which, solving for ε_t^*, gives the closed-form optimal search-effort rule

$$\varepsilon_t^*(m_t, n_t) = \max\left\{0,\ 1 - \exp\left(-\frac{v_t^{adj}(\ell_t) - v_t^{keep}(m_t, n_t)}{S}\right)\right\}. \tag{201}$$

The $\max\{0, \cdot\}$ reflects the constraint $\varepsilon_t \geq 0$: when $v_t^{adj}(\ell_t) \leq v_t^{keep}(m_t, n_t)$, the unconstrained formula would return a non-positive number, but exerting effort to increase the weight on an inferior branch is never optimal, so the constraint binds at $\varepsilon_t^* = 0$. Because $\sigma'(\varepsilon_t) \to \infty$ as $\varepsilon_t \to 1^-$, the upper constraint $\varepsilon_t \leq 1$ never binds for finite $v_t^{adj} - v_t^{keep}$, consistent with search never delivering the adjustment opportunity with certainty. As $S \to 0$, $\varepsilon_t^* \to \mathbb{1}\{v_t^{adj}(\ell_t) > v_t^{keep}(m_t, n_t)\}$, the hard indicator $\mathcal{I}_t$ of the baseline model.

Envelope theorem and combined marginal value functions Because $\varepsilon_t^*(m_t, n_t)$ is chosen to satisfy the first-order (or Karush-Kuhn-Tucker) condition above, the envelope theorem implies that the indirect effect of m_t or n_t operating through ε_t^* vanishes when differentiating $v_t(m_t, n_t) = g(\varepsilon_t^*(m_t, n_t); m_t, n_t)$, leaving only the direct effects through v_t^{keep} and $v_t^{adj}(\ell_t)$:

$$\frac{\partial v_t(m_t, n_t)}{\partial m_t} = (1 - \varepsilon_t^*)\,\frac{\partial v_t^{keep}}{\partial m_t} + \varepsilon_t^*\,\frac{\partial v_t^{adj}}{\partial \ell_t} \cdot \frac{\partial \ell_t}{\partial m_t} = (1 - \varepsilon_t^*)\,\lambda_t^{keep} + \varepsilon_t^*\,\lambda_t^{adj}, \tag{202}$$

$$\frac{\partial v_t(m_t, n_t)}{\partial n_t} = (1 - \varepsilon_t^*)\,\frac{\partial v_t^{keep}}{\partial n_t} + \varepsilon_t^*\,(1 - \tau)\,\lambda_t^{adj} = (1 - \varepsilon_t^*)\Big[u_d(c_t^{keep,*}, n_t) + \beta(1 - \delta)\,\mathbb{E}_t[v_{n,t+1}]\Big] + \varepsilon_t^*\,(1 - \tau)\,\lambda_t^{adj}, \tag{203}$$

using $\lambda_t^{adj} = u_c(c_t^{adj,*}, d_t^{adj,*})$, $\lambda_t^{keep} = u_c(c_t^{keep,*}, n_t)$, and the keeper's $\partial v_t^{keep}/\partial n_t$ from (162). Equations (202)–(203) are the smooth, search-weighted analogues of the hard-indicator combined value functions (164)–(166), with ε_t^* in place of $\mathcal{I}_t$; as $S \to 0$ they coincide exactly, since $\varepsilon_t^* \to \mathcal{I}_t$.

Terminal period The within-branch terminal solutions are unaffected by search: $c_T^{keep} = m_T$ for the keeper and $c_T^{adj*} = \alpha\ell_T$, $d_T^{adj*} = (1 - \alpha)\ell_T$ for the adjuster, exactly as in the "Last period" paragraph of Section B. The terminal discrete choice instead resolves through the optimal terminal search effort $\varepsilon_T^*(m_T, n_T) = \max\{0,\ 1 - \exp(-(v_T^{adj}(\ell_T) - v_T^{keep}(m_T, n_T))/S)\}$, and the smoothed terminal value and marginal-value functions follow (80) and (202)–(203) evaluated at T; these are the objects handed to period $T - 1$.

EGM algorithm with search The algorithm follows the same outline as in the "Discrete Choice" paragraph above, with the hard maximum replaced by the smooth combination:

1. Solve the keeper (EGM plus upper envelope) and the adjuster (root-finding plus FUES) exactly as in Section B, yielding $v_t^{keep}(m_t, n_t)$, $\lambda_t^{keep}(m_t, n_t)$, $v_t^{adj}(\ell_t)$, and $\lambda_t^{adj}(\ell_t)$ on a grid of (m_t, n_t).

2. Compute the optimal search effort $\varepsilon_t^*(m_t, n_t)$ from the closed form (201) at every grid point. As in the full model, this closed form can be evaluated directly, or, for robustness when either conditional value is infinite (e.g. an infeasible branch), effort can instead be selected by a fine grid search over $[0, 1]$ that maximizes $g(\varepsilon_t; m_t, n_t)$; the two approaches coincide whenever both conditional values are finite.

3. Form the combined value function $v_t(m_t, n_t)$ from (80) and the combined marginal-value functions $v_{m,t}(m_t, n_t)$, $v_{n,t}(m_t, n_t)$ from (202)–(203).

4. Feed v_t, $v_{m,t}$, $v_{n,t}$ into the period-$(t-1)$ end-of-period step in place of the hard-maximum objects used in the baseline model.

When a conditional value is infinite (for example because a branch is infeasible), the combined value in step 3 is also infinite; as in the baseline model's discrete-choice paragraph, I store and evaluate inverse representations of the (marginal) value functions and invert them back only when the arrays are used, so that these infinities do not propagate into numerical errors during interpolation.

B.4 Derivations: Stochastic Durable Prices (Section 3.5)

This subsection derives the envelope conditions, the terminal-period solution, and the EGM algorithm for the stochastic-price extension of Section 3.5. The keeper's Bellman problem, first-order condition, and envelope results, equations (159)–(162), carry over unchanged, since the keeper never transacts in the durable market. What changes is the adjuster's budget constraint, its first-order condition with respect to d_t (equations (71)–(73) in the main text), and, consequently, how the two branches are combined.

Adjuster first-order conditions with a stochastic price The price-modified first-order conditions (71)–(73) are stated in the main text (Section 3.5). Equation (71) is identical in form to the baseline first-order condition (54): because p_t multiplies only d_t in the budget constraint, it does not appear in the c_t first-order condition, so $c_t^{adj,*}$ solves the same equation as in the baseline model *given* d_t. Equation (73) is the price-modified analogue of the baseline condition (55): the marginal cost of an extra unit of durables is now $p_t\lambda_t^{adj}$ rather than λ_t^{adj}, since each unit of d_t now costs p_t units of ℓ_t instead of one. As KKT conditions, equality in (71)–(73) holds whenever the borrowing constraint does not bind, and the left-hand side exceeds the right-hand side otherwise.

Envelope conditions Because ℓ_t still enters the adjuster's problem only through the linear resource constraint (67), and p_t is a scalar parameter rather than a function of ℓ_t, the envelope argument of Steps 1–5 above goes through unchanged with p_t held fixed:

$$\frac{\partial v_t^{adj}(\ell_t, p_t)}{\partial m_t} = \frac{\partial v_t^{adj}}{\partial \ell_t} = \lambda_t^{adj} = u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big), \tag{204}$$

$$\frac{\partial v_t^{adj}(\ell_t, p_t)}{\partial n_t} = (1-\tau)\,\frac{\partial v_t^{adj}}{\partial \ell_t} = (1-\tau)\,u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big), \tag{205}$$

which are formally identical to the baseline expressions (157)–(158): the price shock changes *which* $(c_t^{adj,*}, d_t^{adj,*})$ pair is optimal, but not the mapping from the marginal unit of ℓ_t to marginal utility once evaluated at that optimum. Note that p_t does not enter equation (205) directly, because the recovery value $(1-\tau)n_t$ embedded in ℓ_t is, by assumption, insulated from the current-period purchase price.

Discrete choice with a stochastic price Recall the price-contingent discrete choice from equation (81),

$$v_t(m_t, n_t, p_t) = \max\Big\{v_t^{keep}(m_t, n_t),\ v_t^{adj}(\ell_t, p_t)\Big\}. \tag{206}$$

Define the state- and price-contingent indicators $\mathbb{1}_t^{keep}(m_t, n_t, p_t)$ and $\mathbb{1}_t^{adj}(m_t, n_t, p_t) = 1 - \mathbb{1}_t^{keep}(m_t, n_t, p_t)$ analogously to the hard-maximum baseline, with $\mathbb{1}_t^{adj} = 1$ iff $v_t^{adj}(\ell_t, p_t) \geq v_t^{keep}(m_t, n_t)$. Applying

the envelope theorem to the outer maximum exactly as in equations (164)–(166), but now with the price-contingent adjuster policies $c_t^{adj,*}(p_t)$, $d_t^{adj,*}(p_t)$ from (71)–(73),

$$\frac{\partial v_t(m_t, n_t, p_t)}{\partial m_t} = \mathbb{1}_t^{keep}\, u_c(c_t^{keep,*}, n_t) + \mathbb{1}_t^{adj}\, u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big), \tag{207}$$

$$\frac{\partial v_t(m_t, n_t, p_t)}{\partial n_t} = \mathbb{1}_t^{keep}\Big[u_d(c_t^{keep,*}, n_t) + \beta(1-\delta)\,\mathbb{E}_t\big[\bar{v}_{n,t+1}\big]\Big] + \mathbb{1}_t^{adj}\,(1-\tau)\,u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big). \tag{208}$$

Unlike the EV Type I taste shocks of Section B.2, which reweight the two branches through logit probabilities without changing the adjuster's continuous choices, the price shock affects these marginal values only indirectly, through its effect on the optimal $(c_t^{adj,*}, d_t^{adj,*})$ pair and on the indicator functions themselves.

Integrating out the price shock As in the taste-shock case, the object that enters the time-$(t-1)$ Euler equations is the price-integrated value and marginal-value functions,

$$\bar{v}_t(m_t, n_t) \equiv \mathbb{E}_{p_t}\big[v_t(m_t, n_t, p_t)\big], \tag{209}$$

$$\bar{v}_{m,t}(m_t, n_t) \equiv \mathbb{E}_{p_t}\Big[\mathbb{1}_t^{keep}\, u_c(c_t^{keep,*}, n_t) + \mathbb{1}_t^{adj}\, u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big)\Big], \tag{210}$$

$$\bar{v}_{n,t}(m_t, n_t) \equiv \mathbb{E}_{p_t}\Big[\mathbb{1}_t^{keep}\Big(u_d(c_t^{keep,*}, n_t) + \beta(1-\delta)\,\mathbb{E}_t\big[\bar{v}_{n,t+1}\big]\Big) + \mathbb{1}_t^{adj}\,(1-\tau)\,u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big)\Big], \tag{211}$$

where $\bar{v}_{m,t+1}$ and $\bar{v}_{n,t+1}$ on the right-hand side of (72)–(73) one period ahead are themselves defined by the analogous expressions one period later; equation (209) restates equation (82) from the main text. As $\sigma_p \to 0$, $p_t \to 1$ a.s., $\mathbb{1}_t^{adj} \to \mathcal{I}_t$, and equations (210)–(211) collapse to the hard-maximum combined marginal-value functions (164)–(166).

The model's ex-ante (price-weighted) policy functions follow the same logic. Because the household observes p_t before choosing, the realized policy for a given draw is a hard 0/1 choice, $c_t = \mathbb{1}_t^{keep}\, c_t^{keep} + \mathbb{1}_t^{adj}\, c_t^{adj}(p_t)$ and $d_t = \mathbb{1}_t^{keep}\, n_t + \mathbb{1}_t^{adj}\, d_t^{adj}(p_t)$, while the ex-ante policy functions reported for simulation and plotting are

$$\bar{c}_t(m_t, n_t) \equiv \mathbb{E}_{p_t}\Big[\mathbb{1}_t^{keep}\, c_t^{keep} + \mathbb{1}_t^{adj}\, c_t^{adj}(p_t)\Big], \tag{212}$$

$$\bar{d}_t(m_t, n_t) \equiv \mathbb{E}_{p_t}\Big[\mathbb{1}_t^{keep}\, n_t + \mathbb{1}_t^{adj}\, d_t^{adj}(p_t)\Big]. \tag{213}$$

Euler equations The within-branch first-order conditions (71)–(73) are unchanged in form, but, since p_{t+1} is drawn afresh next period, the continuation value entering them is the price-integrated $\bar{v}_{t+1}$ rather than v_{t+1}:

$$u_c\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big) \geq \beta\mathcal{R}\,\mathbb{E}_t\big[\bar{v}_{m,t+1}\big], \tag{214}$$

$$u_d\big(c_t^{adj,*}(p_t), d_t^{adj,*}(p_t)\big) \geq p_t\,\beta\mathcal{R}\,\mathbb{E}_t\big[\bar{v}_{m,t+1}\big] - \beta(1-\delta)\,\mathbb{E}_t\big[\bar{v}_{n,t+1}\big], \tag{215}$$

and analogously for the keeper with $d_t = n_t$ fixed and no price term, since the keeper does not purchase durables. Relative to the baseline Euler equation (55), the only change is the multiplication of the discounted marginal value of assets by the current-period price p_t on the right-hand side of the durable first-order condition.

Terminal period The keeper's terminal solution is unaffected by the price shock: $c_T^{keep} = m_T$, exactly as in the "Last period" paragraph above. For the adjuster, since $v_{T+1} \equiv 0$ and borrowing is prohibited, the terminal problem is

$$\max_{c,d} \; u(c,d) \qquad \text{s.t.} \qquad \ell_T = c + p_T\, d. \tag{216}$$

The Lagrangian $\mathscr{L} = u(c,d) + \lambda(\ell_T - c - p_T d)$ gives $u_c(c,d) = \lambda$ and $u_d(c,d) = p_T \lambda$, so

$$\frac{u_d(c,d)}{u_c(c,d)} = p_T \quad \Longleftrightarrow \quad \frac{1-\alpha}{\alpha} \cdot \frac{c}{d} = p_T \quad \Longleftrightarrow \quad c = \frac{\alpha}{1-\alpha}\, p_T\, d. \tag{217}$$

Substituting into the budget constraint,

$$\ell_T = c + p_T d = p_T d\Big(\frac{\alpha}{1-\alpha} + 1\Big) = \frac{p_T\, d}{1-\alpha} \quad \Longrightarrow \quad d_T^{adj,*} = \frac{(1-\alpha)\,\ell_T}{p_T}, \qquad c_T^{adj,*} = \alpha\, \ell_T. \tag{218}$$

As in the baseline model, $c_T^{adj,*} = \alpha \ell_T$ retains the fixed Cobb-Douglas expenditure share on non-durables regardless of the price, since a proportional increase in p_T raises the cost per unit of d_T without altering the optimal *expenditure* share $(1-\alpha)\ell_T$ devoted to durables; the corresponding physical quantity purchased, $d_T^{adj,*} = (1-\alpha)\ell_T / p_T$, falls one-for-one with the price. The terminal discrete choice is $v_T(m_T, n_T, p_T) = \max\{v_T^{keep}(m_T, n_T),\, v_T^{adj}(\ell_T, p_T)\}$, and $\bar{v}_T(m_T, n_T) = \mathbb{E}_{p_T}[v_T(m_T, n_T, p_T)]$ is the continuation value handed to period $T-1$.

EGM algorithm with stochastic durable prices Let $\vec{p} = \{p_1, \ldots, p_K\}$ be quadrature nodes for the log-normal price density with weights $\vec{\mu} = \{\mu_1, \ldots, \mu_K\}$, in addition to the income-shock integration used to build the end-of-period functions $\mathfrak{v}(d_t, a_t)$, $\mathfrak{v}'_c(d_t, a_t)$, and $\mathfrak{v}'_d(d_t, a_t)$ (Section B, "Calculate End-Of-Period Functions"). Crucially, because the price shock alters the adjuster's own first-order condition (73) rather than merely reweighting an already-computed value, the adjuster problem must be re-solved at every price node, in contrast to the EV Type I taste-shock extension of Section B.2, where a single adjuster solve and a closed-form log-sum combine suffice. The algorithm proceeds as follows:

1. **Solve the keeper once,** by EGM and upper envelope exactly as in the "Solve Keeper" paragraph above; the keeper's solution does not depend on p_t.
2. **Solve the adjuster once per price node.** For each p_k, loop over the end-of-period durable grid d_t and, for each d_t, over the grid of a_t:
   (a) Invert $u_c(c_t, d_t) = \mathfrak{v}'_c(d_t, a_t)$ for c_t exactly as in the baseline model; equation (71) confirms this step does not involve p_k.
   (b) Find the roots of the price-modified first-order condition implied by (73),
   $$0 = u_d(c_t, d_t) - p_k\, \mathfrak{v}'_c(d_t, a_t) + \mathfrak{v}'_d(d_t, a_t), \tag{219}$$
   in place of $0 = u_d(c_t, d_t) - \mathfrak{v}'_c(d_t, a_t) + \mathfrak{v}'_d(d_t, a_t)$ used in the baseline model.
   (c) Compute the price-contingent liquidity $\ell_t^{(k)} = c_t + p_k\, d_t + a_t$ and the value $v_t = u(c_t, d_t) + \mathfrak{v}(d_t, a_t)$ (unchanged in form).
   (d) Apply FUES over $\ell_t^{(k)}$ to obtain $c_t^{adj}(\ell_t; p_k)$, $d_t^{adj}(\ell_t; p_k)$, and $v_t^{adj}(\ell_t; p_k)$.

3. **Combine on a grid of (m_t, n_t).** For each grid point, compute $\ell_t = m_t + (1-\tau)n_t$ (independent of p_k) and $v_t^{keep}(m_t, n_t)$, then for each node p_k evaluate $v_t^{adj}(\ell_t; p_k)$, form the indicators $\mathbb{1}^{keep}(m_t, n_t, p_k)$, $\mathbb{1}^{adj}(m_t, n_t, p_k)$, and integrate as in equations (209)–(211) (and, for reporting, equations (212)–(213)).

4. **Feed $\bar{v}_t$, $\bar{v}_{m,t}$, $\bar{v}_{n,t}$ into the period-$(t-1)$ end-of-period step,** in place of the hard-maximum v_t, $v_{m,t}$, $v_{n,t}$ used in the baseline model.

As in the baseline and taste-shock algorithms, I store and evaluate inverse representations of the (marginal) value functions so that any infinite conditional values (e.g. from an infeasible branch) do not propagate into numerical errors during interpolation.

B.5 Derivations: Search and Stochastic Prices (Section 3.5)

This subsection derives the optimal search-effort rule, the combined envelope conditions, the Euler equations, the terminal-period solution, and the EGM algorithm for the nested search-and-price environment of Section 3.5. The within-branch building blocks are taken from the standalone extensions: the keeper is exactly as in Section B; the adjuster inherits the price-contingent first-order conditions (71)–(73) and envelope results (204)–(205); and the search cost $\sigma(\varepsilon_t)$ and its derivative $\sigma'(\varepsilon_t) = -S\log(1-\varepsilon_t)$ are those of equations (64) and (199). What is new is the joint combination of search with price risk (without an additional EV Type I or multiplicative taste shock).

Adjuster first-order conditions Holding p_t fixed, the adjuster's Lagrangian is identical to the stochastic-price case,

$$\mathscr{L} = u(c_t, d_t) + \beta\mathbb{E}_t\big[\bar{v}_{t+1}(m_{t+1}, n_{t+1})\big] + \lambda_t^{adj}\big(\ell_t - c_t - p_t\, d_t - a_t\big), \tag{220}$$

where the continuation value is now the price-integrated $\bar{v}_{t+1}$ from equation (79) one period ahead (rather than the price-only integral of Section B.4, which used a hard maximum rather than search). The first-order conditions are therefore

$$\lambda_t^{adj} = u_c(c_t, d_t) = \beta\mathcal{R}\,\mathbb{E}_t\big[\bar{v}_{m,t+1}\big], \tag{221}$$

$$u_d(c_t, d_t) = p_t\,\lambda_t^{adj} - \beta(1-\delta)\,\mathbb{E}_t\big[\bar{v}_{n,t+1}\big], \tag{222}$$

with equality when the borrowing constraint does not bind. The price p_t enters only through the durable first-order condition (222), exactly as in (73). The keeper's first-order condition is likewise unchanged relative to the baseline, with $d_t = n_t$ fixed and no price term.

Optimal search effort Recall the nested discrete problem (77). Effort is chosen before p_t is realized, so the gap in the first-order condition is the expected positive value of a match, $\bar{\Delta}(m_t, n_t) = \mathbb{E}_{p_t}[\max\{v_t^{adj}(\ell_t, p_t) - v_t^{keep}(m_t, n_t),\, 0\}]$, not the gap at a realized price. Because σ is strictly convex, the objective is strictly concave in ε_t. The interior first-order condition is

$$\bar{\Delta}(m_t, n_t) - \sigma'(\varepsilon_t^*) = 0 \quad\Longleftrightarrow\quad \bar{\Delta}(m_t, n_t) = -S\log(1-\varepsilon_t^*), \tag{223}$$

which yields (78). Let $\mathbb{1}(p_t) = 1$ when $v_t^{adj}(\ell_t, p_t) \geq v_t^{keep}(m_t, n_t)$ and write $\pi_t = \mathbb{E}_{p_t}[\mathbb{1}(p_t)]$. The probability of buying is $\varepsilon_t^*\pi_t$. As in the search-only case, $\varepsilon_t^* = 0$ whenever $\bar{\Delta} \leq 0$, and the upper constraint $\varepsilon_t \leq 1$ never binds for finite S. With $p_t \equiv 1$, (78) reduces to (201). With $S \to 0$ and $\bar{\Delta} > 0$, effort goes to one and the acceptance decision is the hard indicator of (81).

Envelope conditions and combined marginal values Because ε_t^* is chosen optimally, the envelope theorem eliminates indirect effects through search effort. The accept indicator is a hard maximum, so it contributes no derivative either. The ex-ante marginal values put weight $1-\varepsilon_t^*\pi_t$ on the keeper and weight ε_t^* on the accepted adjuster nodes:

$$\frac{\partial v_t(m_t,n_t)}{\partial m_t} = (1-\varepsilon_t^*\pi_t)\,u_c(c_t^{keep,*},n_t) + \varepsilon_t^*\,\mathbb{E}_{p_t}\Big[\mathbb{1}\!\!\!\!\mathbb{1}(p_t)\,u_c\big(c_t^{adj,*}(p_t),d_t^{adj,*}(p_t)\big)\Big], \tag{224}$$

$$\begin{aligned}\frac{\partial v_t(m_t,n_t)}{\partial n_t} &= (1-\varepsilon_t^*\pi_t)\Big[u_d(c_t^{keep,*},n_t) + \beta(1-\delta)\,\mathbb{E}_t\big[\bar v_{n,t+1}\big]\Big] \\ &\quad + \varepsilon_t^*\,(1-\tau)\,\mathbb{E}_{p_t}\Big[\mathbb{1}\!\!\!\!\mathbb{1}(p_t)\,u_c\big(c_t^{adj,*}(p_t),d_t^{adj,*}(p_t)\big)\Big].\end{aligned} \tag{225}$$

With a degenerate price and $\pi_t \in \{0,1\}$, these collapse to the search-only envelopes (202)–(203). With $S \to 0$ and $\bar\Delta > 0$, they collapse to the price-only indicators (207)–(208).

Objects passed to the previous period The value passed to period $t-1$ is (79). The marginal values are (224) and (225),

$$\bar v_{m,t}(m_t,n_t) \equiv \frac{\partial v_t(m_t,n_t)}{\partial m_t}, \tag{226}$$

$$\bar v_{n,t}(m_t,n_t) \equiv \frac{\partial v_t(m_t,n_t)}{\partial n_t}. \tag{227}$$

The ex-ante policies use the same lottery. A rejected draw, and a failed search, both leave the household on the keeper:

$$\bar c_t(m_t,n_t) \equiv (1-\varepsilon_t^*\pi_t)\,c_t^{keep} + \varepsilon_t^*\,\mathbb{E}_{p_t}\big[\mathbb{1}\!\!\!\!\mathbb{1}(p_t)\,c_t^{adj}(p_t)\big], \tag{228}$$

$$\bar d_t(m_t,n_t) \equiv (1-\varepsilon_t^*\pi_t)\,n_t + \varepsilon_t^*\,\mathbb{E}_{p_t}\big[\mathbb{1}\!\!\!\!\mathbb{1}(p_t)\,d_t^{adj}(p_t)\big], \tag{229}$$

$$\bar\varepsilon_t(m_t,n_t) \equiv \varepsilon_t^*(m_t,n_t)\,\pi_t. \tag{230}$$

Equation (230) is the probability of buying. Effort itself is $\varepsilon_t^*(m_t,n_t)$ from (78). Conditional on a match at p_t, the realized policy is the pure adjuster when $\mathbb{1}\!\!\!\!\mathbb{1}(p_t)=1$ and the pure keeper otherwise.

Euler equations Combining (221)–(222) with the integrated continuation values (226)–(227) yields the within-branch Euler equations

$$u_c\big(c_t^{adj,*}(p_t),d_t^{adj,*}(p_t)\big) \geq \beta\mathcal{R}\,\mathbb{E}_t\big[\bar v_{m,t+1}\big], \tag{231}$$

$$u_d\big(c_t^{adj,*}(p_t),d_t^{adj,*}(p_t)\big) \geq p_t\,\beta\mathcal{R}\,\mathbb{E}_t\big[\bar v_{m,t+1}\big] - \beta(1-\delta)\,\mathbb{E}_t\big[\bar v_{n,t+1}\big], \tag{232}$$

and analogously for the keeper with $d_t = n_t$ fixed and no price term. Relative to the baseline, the only structural change inside the adjuster's durable Euler equation is the factor p_t on the opportunity-cost term; relative to the price-only extension, the continuation values $\bar v_{m,t+1}$ and $\bar v_{n,t+1}$ additionally average over next period's search effort.

Terminal period The within-branch terminal solutions are those of the price extension: $c_T^{keep} = m_T$ for the keeper, and

$$c_T^{adj,*} = \alpha\,\ell_T, \qquad d_T^{adj,*} = \frac{(1-\alpha)\,\ell_T}{p_T}, \tag{233}$$

from equation (218). The terminal discrete problem then resolves through (78) evaluated at T. Search is defined in the terminal period because the benefit is the current comparison. The continuation value handed to period $T-1$ is $\bar v_T$ from (79).

EGM algorithm with search and stochastic prices Let $\vec{p} = \{p_1, \ldots, p_{K_p}\}$ be quadrature nodes for the log-normal price density with weights $\vec{\mu}^p$, in addition to the income-shock integration used to build $\mathfrak{v}(d_t, a_t)$, $\mathfrak{v}'_c(d_t, a_t)$, and $\mathfrak{v}'_d(d_t, a_t)$. The algorithm nests the standalone price and search steps:

1. **Solve the keeper once,** by EGM and upper envelope exactly as in the "Solve Keeper" paragraph of Section B; the keeper does not depend on p_t.

2. **Solve the adjuster once per price node.** For each p_k, proceed exactly as in the price-only algorithm of Section B.4: invert $u_c = \mathfrak{v}'_c$ for c_t, find roots of

$$0 = u_d(c_t, d_t) - p_k\, \mathfrak{v}'_c(d_t, a_t) + \mathfrak{v}'_d(d_t, a_t), \tag{234}$$

   form $\ell_t^{(k)} = c_t + p_k\, d_t + a_t$, and apply FUES to obtain $c_t^{adj}(\ell_t; p_k)$, $d_t^{adj}(\ell_t; p_k)$, and $v_t^{adj}(\ell_t; p_k)$.

3. **Combine on a grid of (m_t, n_t).** For each grid point, compute $\ell_t = m_t + (1-\tau)n_t$ and $v_t^{keep}(m_t, n_t)$. Evaluate $v_t^{adj}(\ell_t; p_k)$ at every price node, form $\bar{\Delta}$, π_t, and one effort $\varepsilon^*(m_t, n_t)$ from (78), and accept at node k when $v_t^{adj}(\ell_t; p_k) \geq v_t^{keep}$. The value and marginal values are (79) and (224)–(225). The reported policies are (228)–(230).

4. **Feed $\bar{v}_t$, $\bar{v}_{m,t}$, $\bar{v}_{n,t}$ into the period-$(t-1)$ end-of-period step,** in place of the hard-maximum objects used in the baseline model.

As in the earlier algorithms, I store and evaluate inverse representations of the (marginal) value functions so that infinite conditional values do not propagate into numerical errors during interpolation. Computational cost scales with K_p for both the adjuster solve and the combine step.

C A Theory of Smoothness in Discrete–Continuous Models

C.1 Purpose and roadmap

The main text establishes by numerical illustration that taste shocks and directed search reduce the discontinuities generated by the EGM step in discrete–continuous (DC) models. This appendix replaces those illustrations with propositions. It does three things: it shows *why* the hard-max model fails (Section C.5), *how* taste shocks repair it (Section C.6), and *how* directed search repairs it (Section C.7), in each case with the exact smoothness class, an explicit bound on curvature, and an explicit bound on the model bias incurred.

Two preliminary remarks shape everything that follows.

First, "smoothness" is not the operative property. What determines whether the EGM step produces suboptimal points is *strict monotonicity of the endogenous grid map* $a \mapsto m(a)$, and by Lemma C.5 that is a condition on *bounded second derivatives* of the continuation value, not on differentiability as such. Differentiability is a separate matter: it governs whether marginal propensities to consume jump. Separating these two targets is what allows the claim in the abstract (suboptimal points can be eliminated by search alone) and the claim in the conclusion (uniform smoothness additionally requires taste shocks) to be stated as two distinct and simultaneously correct propositions. Under directed search with the log-type cost function (64), the first holds and the second fails; see Proposition C.15.

Second, all three mechanisms are the same mathematical object evaluated at different parameters. Each is the convex conjugate of a cost of randomizing over the discrete alternatives (Lemma C.6). The smoothness class of the solved model is governed entirely by the local behaviour of that cost at the boundary of the choice simplex. This yields the sharpest result in the appendix, Proposition C.17: the apparent conflict between concentrating adjustment near the (S, s) thresholds and obtaining smooth decision rules is an artifact of the particular cost function in (64), not a theorem, and both properties can be obtained simultaneously.

Remark C.1 (Notation)**.** The Extreme Value Type I draws are ζ_t^{keep} and ζ_t^{adj}, with scale σ_ζ, in both the retirement model and Section 3.4. Search intensity remains ε_t.

C.2 A common reduced form

Let $x \in X \subseteq \mathbb{R}^n$ denote the beginning-of-period continuous state: $x = M_t$ in the retirement model of Section 2 and $x = (m_t, n_t)$ in the durable model of Section 3.1. Let $j \in \{0, 1\}$ index the two discrete branches, with *conditional value functions* $v_t^0(x)$ and $v_t^1(x)$, each the value of a continuous optimisation problem solved conditional on j. In the retirement model $(v^0, v^1) = (V_t^{\text{Ret}}, V_t^W)$; in the durable model $(v^0, v^1) = (v_t^{\text{keep}}, v_t^{\text{adj}} \circ \ell)$ with $\ell(m, n) = m + (1 - \tau)n$. Define the *value gap* and the *indifference set*

$$\Delta_t(x) \equiv v_t^1(x) - v_t^0(x), \qquad \mathcal{K}_t \equiv \{x \in X : \Delta_t(x) = 0\}. \tag{235}$$

All three mechanisms resolve the discrete choice through an *aggregator* $\Lambda : \mathbb{R}^2 \to \mathbb{R}$ applied pointwise, $v_t(x) = \Lambda(v_t^0(x), v_t^1(x))$. Because each is translation equivariant, $\Lambda(v^0 + k, v^1 + k) = \Lambda(v^0, v^1) + k$ for all $k \in \mathbb{R}$, every aggregator admits the *reduced form*

$$v_t(x) = v_t^0(x) + \Phi\big(\Delta_t(x)\big) \tag{236}$$

for a single scalar function $\Phi : \mathbb{R} \to \mathbb{R}$ with $\Phi(0)$ finite. This is the device that reduces every regularity question below to a statement about Φ near the origin, composed with a transversality

condition on Δ_t. The three cases of interest are

$$\Phi^{\max}(\delta) = \max\{0, \delta\}, \qquad \Phi^{\sigma}(\delta) = \sigma_\zeta \log\big(1 + e^{\delta/\sigma_\zeta}\big), \qquad \Phi^{S}(\delta) = \max_{\varepsilon \in [0,1]} \big\{\varepsilon\delta - \sigma(\varepsilon)\big\}, \tag{237}$$

corresponding respectively to the hard maximum of Section 3.1, the Extreme Value Type I log-sum (60), and the directed-search problem (80) with cost (64).

Because the hard-max case produces functions that are not globally differentiable, the C^k scale alone is insufficient bookkeeping. For a closed set $\mathcal{N} \subset X$ of measure zero, let $\mathcal{PC}^k(\mathcal{N})$ denote the functions that are C^k on $X \setminus \mathcal{N}$ with finite one-sided limits of all derivatives up to order k at each point of $\mathcal{N}$. Write $\mathcal{N}_t$ for the *singular set* of v_t, i.e. the set at which v_t fails to be C^1. Tracking the cardinality (or the surface measure) of $\mathcal{N}_t$ across backward iterations is what formalises the statement in Section 2 that kinks are "propagated back in time, thus, exacerbating the problem."

Finally, let

$$\mathcal{E}_t(a) \equiv \mathbb{E}_t\big[v_{m,t+1}(\mathcal{R}a + \theta_{t+1})\big], \qquad m_t(a) \equiv a + c_t(a), \tag{238}$$

so that $\mathcal{E}_t$ is the object the EGM step inverts and m_t is the endogenous grid map. The following definition makes the paper's central claim precise.

Definition C.2. The EGM step at t is *clean* if $a \mapsto m_t(a)$ is continuous and strictly increasing on the unconstrained region. If the EGM step is clean, the candidate points returned by the Euler equation require no upper-envelope refinement.

C.3 Assumptions

Assumption C.1 (Felicity)**.** u is C^∞ on the interior of its domain, strictly increasing and strictly concave in each argument, with Inada conditions in each argument. In the durable model $u(c,d) = (c^\alpha d^{1-\alpha})^{1-\rho}/(1-\rho)$ with $\alpha \in (0,1)$, $\rho > 0$; this u is concave, being the composition of the increasing concave map $z \mapsto z^{1-\rho}/(1-\rho)$ with the concave, degree-one homogeneous aggregator $c^\alpha d^{1-\alpha}$.

Assumption C.2 (Interiority and regular constraint set)**.** For each t and each x in the region of interest the optimal continuous choice is interior, and the set on which $a_t \geq 0$ binds is relatively open with a C^1 boundary. Moreover the conditional marginal propensity to consume satisfies $0 < \partial c_t^j/\partial m < 1$ on the unconstrained region.

Assumption C.2 is not cosmetic. The borrowing constraint generates a kink of its own at the constrained/unconstrained boundary, and any statement that kinks in the reported policy functions are *discrete-choice* kinks requires that this second source be separated out. The bound $\partial c/\partial m < 1$ is used only in Proposition C.18.

Assumption C.3 (Terminal regularity)**.** At T both conditional value functions are C^∞ on the interior of the state space.

Assumption C.3 holds in both applications and supplies the base case of every induction below: the retirement terminal problem is the closed-form logarithmic problem, and the durable terminal problem has $c_T^{\text{keep}} = m_T$, $c_T^{\text{adj}} = \alpha\ell_T$ and $d_T^{\text{adj}} = (1-\alpha)\ell_T/p_T$.

Assumption C.4 (Transversality)**.** For each t, $\mathcal{K}_t$ is a nonempty C^1 embedded submanifold of X of codimension one, and $\nabla\Delta_t(x) \neq 0$ for every $x \in \mathcal{K}_t$. Define the *jump magnitude*

$$\kappa_t \equiv \inf_{x \in \mathcal{K}_t} \big|\nabla\Delta_t(x)\big| > 0, \qquad \Xi_t \equiv \sup_{x \in \mathcal{K}_t} \big|\nabla\Delta_t(x)\big| < \infty. \tag{239}$$

In the retirement model this reduces to $\lambda_t^W(M^*) \neq \lambda_t^{\text{Ret}}(M^*)$ at each crossing point M^*; in the durable model $\nabla\Delta_t = \big(\lambda_t^{\text{adj}} - \lambda_t^{\text{keep}},\ (1-\tau)\lambda_t^{\text{adj}} - v_{n,t}^{\text{keep}}\big)$.

Assumption C.4 does double duty. It is what makes the negative result *true* — a degenerate tangential crossing would produce no kink at all — and it is what guarantees that the smoothing region in Propositions C.12 and C.15 has positive width of order σ_ζ and S respectively.

Assumption C.5 (Single crossing)**.** For each fixed n, $\Delta_t(\cdot, n)$ has finitely many zeros, and in the durable model $\partial_m \Delta_t > 0$ on $\mathcal{K}_t$.

Assumption C.5 is what delivers an (S, s) *band* rather than an arbitrary adjustment set. It is the least innocent assumption in this list: proving it from primitives in a model whose continuation value is itself an aggregate of two branches is not attempted here, and it should be verified numerically on the solved model (Section C.12).

Assumption C.6 (Shock regularity)**.** The transitory income shock θ_t has a density $f_\theta \in C^\infty$ with all derivatives integrable against the kernels appearing below; the mean-one log-normal specification of Sections 2 and 3.1 qualifies. Where a taste shock ξ_t is present its density f_ξ satisfies the same conditions.

Assumption C.6 refers to the *model*. The *computed* approximation replaces these densities by J- and K-node quadratures. Lemma C.4 shows that these two objects have entirely different regularity, which is the rigorous content of the phantom-kink discussion in Section **??**.

Assumption C.7 (Impatience)**.** $\beta\mathcal{R}^2 < 1$.

Assumption C.7 is required only for the uniform-in-t estimate of Proposition C.18 and is satisfied by both calibrations used in the paper: $\mathcal{R} = 1, \beta = 0.98$ gives $\beta\mathcal{R}^2 = 0.98$, and $\mathcal{R} = 1.04^{0.25}, \beta = 0.975$ gives $\beta\mathcal{R}^2 \approx 0.994$.

Assumption C.8 (Level normalisation; retirement model only)**.** The felicity function is normalised so that the multiplicative taste shock $\xi_t V_t^W$ of Section 2 is well defined, i.e. the additive constant in u is fixed.

Assumption C.8 deserves emphasis because it is substantive rather than technical. A value function is determined only up to an affine transformation of felicity, so a shock entering multiplicatively on V_t^W is *not* invariant to that transformation: adding a constant to u changes the model's predictions. The accompanying code fixes the normalisation with an additive constant chosen for interpolation accuracy. Any proposition about the multiplicative log-normal taste shock is therefore conditional on Assumption C.8. The Extreme Value Type I specification of Section 3.4, being additive, requires no such assumption, which is an argument in its favour that the main text does not currently make.

C.4 Four lemmas

Lemma C.3 (Regularity transfer and the envelope bootstrap)**.** *Under Assumptions C.1–C.2, for any $k \geq 0$:*

1. $\mathcal{E}_t \in C^k \implies c_t = (u')^{-1}(\beta\mathcal{R}\mathcal{E}_t) \in C^k$;
2. $c_t^j \in C^k \implies v_t^j \in C^{k+1}$;
3. *consequently* $v_{t+1} \in C^{k+1} \implies v_t^j \in C^{k+1}$ *for* $j = 0, 1$.

Proof. (i) u' is C^∞ with $u'' < 0$ on the interior, so by the inverse function theorem $(u')^{-1}$ is C^∞ on its range; a composition of a C^∞ map with a C^k map is C^k. (ii) The envelope condition gives $v^j_{m,t} = u'(c^j_t)$, which is C^k by the same argument; a function whose gradient is C^k is C^{k+1}. (iii) $v_{t+1} \in C^{k+1}$ implies $v_{m,t+1} \in C^k$, hence $\mathcal{E}_t \in C^k$ (a finite sum or an integral of C^k functions), hence $c_t \in C^k$ by (i), hence $v^j_t \in C^{k+1}$ by (ii). □

Lemma C.3(iii) is the workhorse of the appendix: the *envelope condition buys back the derivative lost in the Euler equation*, so conditional value functions preserve their smoothness class across a backward step. It follows that *only the aggregator can destroy smoothness*, and the entire induction reduces to the question of what class Φ admits.

Lemma C.4 (Mollification: exact integration versus quadrature)**.** *Let g be locally bounded and piecewise C^∞ with finitely many jump discontinuities in g', of magnitudes $[g'](z_i)$ at points z_i. Fix $y > 0$.*

1. (Exact integration smooths infinitely.) *Under Assumption C.6, $G(b) \equiv \int g(b + y\eta) f(\eta)\, d\eta \in C^\infty$, regardless of the singularities of g.*
2. (Quadrature does not smooth at all.) *For $G_J(b) \equiv \sum_{j=1}^J \omega_j\, g(b+y\eta_j)$ with $\omega_j > 0$, $\sum_j \omega_j = 1$, the derivative G'_J has a jump of magnitude $\omega_j\, [g'](z_i)$ at each b with $b + y\eta_j = z_i$. Hence $G_J \notin C^1$ for every J.*
3. (Conservation of total jump.) *If the preimages are distinct, then for every J,*

$$\sum_i \sum_j \big|[G'_J]\big| \;=\; \sum_i \big|[g'](z_i)\big|, \tag{240}$$

so the largest *jump falls at rate $O(\max_j \omega_j) = O(1/J)$ while the* total *jump is invariant in J.*

Proof. (i) Substituting $z = b + y\eta$ gives $G(b) = y^{-1} \int g(z) f\big((z-b)/y\big) dz$. Differentiation in b falls on f, not on g: $G^{(r)}(b) = (-1)^r y^{-(r+1)} \int g(z) f^{(r)}\big((z-b)/y\big) dz$, which is finite for every r under Assumption C.6 and continuous in b by dominated convergence. (ii) Immediate from $G'_J(b) = \sum_j \omega_j g'(b + y\eta_j)$ wherever defined, evaluating one-sided limits. (iii) Sum the jumps in (ii) over j and use $\sum_j \omega_j = 1$. □

Lemma C.4 is the precise version of the phantom-kink claim of Section **??**. Part (iii) says that refining the income grid makes the solution *look* smooth without ever making it smooth, and it predicts exactly the pattern in Table **??**: the mean absolute deviation collapses in J while the maximum absolute deviation falls slowly and non-monotonically. It also supplies a better diagnostic than a maximum absolute deviation, namely the total variation of c'_t, which is invariant in J under quadrature and finite under exact integration.

Lemma C.5 (EGM monotonicity criterion)**.** *Wherever c_t is differentiable,*

$$m'_t(a) \;=\; 1 + c'_t(a) \;=\; 1 - \frac{\beta \mathcal{R}^2\, \mathbb{E}_t\big[v_{mm,t+1}\big]}{\big|u''(c_t)\big|}. \tag{241}$$

Consequently:

1. *if $\mathbb{E}_t[v_{mm,t+1}] \le 0$ then $m'_t(a) \ge 1$, and the EGM step is clean;*

2. *in general the EGM step is clean if and only if*

$$\beta\mathcal{R}^2\,\mathbb{E}_t\big[v_{mm,t+1}\big] \;<\; \big|u''(c_t)\big| \qquad \textit{pointwise;} \tag{242}$$

3. *if $\mathcal{E}_t$ has a jump discontinuity then m_t has a jump and the EGM step is not clean.*

Proof. Differentiate $u'(c_t(a)) = \beta\mathcal{R}\mathcal{E}_t(a)$ in a, using $\mathcal{E}_t'(a) = \mathcal{R}\,\mathbb{E}_t[v_{mm,t+1}]$, to obtain $c_t'(a) = \beta\mathcal{R}^2\mathbb{E}_t[v_{mm,t+1}]/u''(c_t)$; substitute $u'' = -|u''|$ and add one. Parts (i)–(iii) follow by inspection of (241). □

Condition (242) is the operative criterion of this appendix. It converts "the solution is smooth enough" into a checkable inequality, and combined with Proposition C.18 it yields an explicit threshold on the smoothing parameter. Two caveats on how to state it usefully. First, under CRRA felicity $\inf|u''|$ vanishes at high wealth, so (242) is vacuous if stated globally; it should be stated on a compact wealth region, or, preferably and in keeping with the implementation, in inverse-marginal-utility (consumption-equivalent) units so that both sides are scale free. Second, Lemma C.5 presumes differentiability of c_t; where c_t merely has a kink, m_t' is to be read as the pair of one-sided derivatives and the criterion applied to each.

Lemma C.6 (Aggregator regularity via conjugate duality)**.** *Let $\sigma : [0,1] \to \mathbb{R}$ be strictly convex, C^2 on $(0,1)$, with $\sigma(0) = 0$, and let $\Phi = \sigma^*$ be its convex conjugate as in (237). Then:*

1. *Φ is convex and nondecreasing with $\Phi'(\delta) = \varepsilon^*(\delta) \in [0,1]$, and $\varepsilon^* = (\sigma')^{-1}(\delta)$ wherever interior;*

2. *$\varepsilon^*(\delta) = 0$ for all $\delta \le 0$ if and only if $\sigma'(0) = 0$ (*one-sidedness*);*

3. *given $\sigma'(0) = 0$, $\Phi''(0^+) = 1/\sigma''(0^+)$ and $\Phi''(0^-) = 0$, so*

$$\Phi \in C^2 \quad\Longleftrightarrow\quad \sigma''(0^+) = +\infty; \tag{243}$$

4. *if $\sigma'(\varepsilon) = A\varepsilon^{1/\gamma}$ near $\varepsilon = 0$ for an integer $\gamma \ge 1$ and $A > 0$, then $\varepsilon^*(\delta) = (\delta^+/A)^\gamma$ near $\delta = 0$ and $\Phi \in C^\gamma \setminus C^{\gamma+1}$.*

Proof. (i) A conjugate is a supremum of affine functions of δ, hence convex; monotonicity follows from $\varepsilon \ge 0$; $\Phi' = \varepsilon^*$ is Danskin's theorem, and the interior first-order condition is $\delta = \sigma'(\varepsilon^*)$. (ii) The constraint $\varepsilon \ge 0$ binds exactly when $\delta \le \sigma'(0)$. (iii) Differentiating $\delta = \sigma'(\varepsilon^*(\delta))$ gives $\Phi''(\delta) = 1/\sigma''(\varepsilon^*(\delta))$, and $\varepsilon^*(0^+) = 0$; on $\delta < 0$, $\Phi \equiv 0$. (iv) Invert σ' and note that $\delta \mapsto (\delta^+)^\gamma$ is $C^{\gamma-1}$ but not C^γ at the origin, so $\Phi' \in C^{\gamma-1} \setminus C^\gamma$. □

Lemma C.6 places both smoothing devices in a single family. The Extreme Value Type I log-sum is the conjugate of the Shannon entropy on the two-point simplex, and directed search is the conjugate of the search cost. This is the additive-perturbed-utility representation of stochastic discrete choice studied by Mattsson and Weibull [**?**], Hofbauer and Sandholm [**?**] and **?**; see **?** for the duality facts used above. Two implications matter for the paper. First, "rationalised" versus "behavioural" is not a structural distinction between the two mechanisms, since both are conjugates of a control cost and the associated random-utility and perturbed-utility representations are observationally equivalent; the durable arguments for search should therefore rest on one-sidedness (Proposition C.16) and on testability, which do survive. Second, and constructively, the modeller *chooses* σ, so by Lemma C.6(iii)–(iv) the modeller chooses the smoothness class; see Proposition C.17.

C.5 Why the hard-max model is not smooth

Proposition C.7 (Convex kink and marginal-value jump). *Let $\Phi = \Phi^{\max}$ and suppose Assumptions C.1–C.4 hold with $v_t^0, v_t^1 \in C^1$. Then v_t is Lipschitz, $v_t \notin C^1$ at every $x \in \mathcal{K}_t$, and along any C^1 curve γ crossing $\mathcal{K}_t$ transversally at $\gamma(0) = x^*$ in the direction of increasing Δ_t,*

$$\frac{d}{ds} v_t(\gamma(s))\Big|_{s=0^+} - \frac{d}{ds} v_t(\gamma(s))\Big|_{s=0^-} = \frac{d}{ds} \Delta_t(\gamma(s))\Big|_{s=0} > 0. \tag{244}$$

The jump is therefore strictly upward*: v_t has a convex kink and, in the distributional sense, $v_{mm,t}$ carries a strictly positive Dirac mass on $\mathcal{K}_t$.*

Proof. Write $h(s) = v_t(\gamma(s)) = v_t^0(\gamma(s)) + \max\{0, \Delta_t(\gamma(s))\}$. For $s < 0$, $\Delta_t(\gamma(s)) < 0$ and $h'(s) = \frac{d}{ds} v_t^0$; for $s > 0$, $h'(s) = \frac{d}{ds} v_t^0 + \frac{d}{ds} \Delta_t = \frac{d}{ds} v_t^1$. Taking one-sided limits at $s = 0$ and differencing gives (244), which is strictly positive by transversality (Assumption C.4) and the choice of direction. Lipschitz continuity is inherited from the branches. □

Proposition C.7 identifies the jump magnitude with $|\nabla \Delta_t|$, i.e. with the difference in the branches' marginal values, $\kappa_t = |\lambda^1 - \lambda^0|$ in the retirement model. The *sign* matters: because the maximum always selects the steeper branch after a crossing, the kink is convex and $\mathbb{E}_t[v_{mm,t+1}]$ acquires a positive singular part, which is precisely the direction that violates (242).

Proposition C.8 (Quadrature guarantees failure). *Let $\Phi = \Phi^{\max}$ and let the income expectation be computed by J-node quadrature. Then $\mathcal{E}_{t-1}$ has a jump of magnitude $\omega_j \kappa_t$ at each a with $\mathcal{R}a + y\eta_j \in \mathcal{K}_t$. Consequently c_{t-1} has at least $J \cdot |\mathcal{K}_t|$ jump discontinuities, m_{t-1} is discontinuous, and by Lemma C.5(iii) the EGM step is not clean — for every J and every parameterisation satisfying Assumption C.4.*

Proof. Apply Lemma C.4(ii) with $g = v_{m,t}$, whose jumps are given by Proposition C.7, then Lemma C.3(i) and Lemma C.5(iii). □

Proposition C.9 (Accumulation). *Under the hypotheses of Proposition C.8, the singular set satisfies*

$$\mathcal{N}_{t-1} \supseteq \mathcal{K}_{t-1} \cup \big\{x : \mathcal{R}a(x) + y\eta_j \in \mathcal{K}_t \text{ for some } j \le J\big\}, \tag{245}$$

so $|\mathcal{N}_{t-1}| \ge |\mathcal{K}_{t-1}| + J\,|\mathcal{K}_t|$ and the number of points at which the conditional value functions fail to be C^1 is nondecreasing, and generically strictly increasing, as t recedes from T.

Proof. By Proposition C.8 the conditional policy c_{t-1}^j has jumps at the stated points, so by the envelope condition $v_{m,t-1}^j = u'(c_{t-1}^j)$ has jumps there and $v_{t-1}^j \notin C^1$; the new indifference set $\mathcal{K}_{t-1}$ contributes additionally by Proposition C.7. Note that Lemma C.3(iii) does not apply, since it requires $k \ge 0$ and here $v_{m,t} \notin C^0$. □

Proposition C.9 is the formal content of the informal claim that kinks propagate backwards and worsen the problem. It also shows precisely where Lemma C.3 breaks: the envelope bootstrap preserves smoothness only from C^0 upwards, and the hard maximum pushes the solution below that floor.

Proposition C.10 (Exact income integration is necessary but not sufficient). *Let $\Phi = \Phi^{\max}$ and let the income expectation be exact under Assumption C.6. Then $\mathcal{E}_{t-1} \in C^\infty$ and c_{t-1} is continuous, but*

$$\mathbb{E}_{t-1}\big[v_{mm,t}\big] \ni \frac{\kappa_t}{y} f_\theta(\eta^*), \qquad \eta^* = \frac{M^* - \mathcal{R}a}{y}, \tag{246}$$

as a spike at each preimage of $\mathcal{K}_t$, so by (242) the EGM step fails to be clean whenever

$$\beta \mathcal{R}^2 \frac{\kappa_t}{y} f_\theta(\eta^*) \;\geq\; \left|u''(c_{t-1})\right|. \tag{247}$$

Proof. Continuity and smoothness of $\mathcal{E}_{t-1}$ follow from Lemma C.4(i). Decompose $v_{m,t} = \tilde{v}_{m,t} + \kappa_t \mathbb{1}\{z \geq M^*\}$ with $\tilde{v}_{m,t}$ continuous. The indicator contributes $\kappa_t \Pr[\mathcal{R}a + y\theta \geq M^*]$ to $\mathcal{E}_{t-1}$, whose derivative in a is $\kappa_t \mathcal{R} f_\theta(\eta^*)/y$; dividing by $\mathcal{R}$ gives the stated contribution to $\mathbb{E}[v_{mm,t}]$. Apply Lemma C.5(ii). □

Proposition C.10 rationalises the finding that income uncertainty adds smoothness but does not remove the discontinuities: exact integration upgrades the policy from discontinuous to continuous, but leaves a curvature spike whose height is inversely proportional to the dispersion of income. The discrete choice must be smoothed as well.

Remark C.11 (Two distinct discontinuities)**.** Two different objects are non-smooth under a hard maximum and the main text does not consistently distinguish them.

1. The *combined* policy $c_t = (1 - \mathbb{1}_t) c_t^{\text{keep}} + \mathbb{1}_t c_t^{\text{adj}}$ jumps on $\mathcal{K}_t$ because the argmax switches branch. This is genuine economics; no solution method removes it, and a smoothing device removes it only by reporting an ex-ante expected policy in its place.
2. The *conditional* policy c_t^j jumps because $v_{m,t+1}$ jumps (Proposition C.8). This is what breaks the EGM step.

The figures in Sections 2 and 4 plot objects of type (ii), and the text should say so explicitly; otherwise the paper is open to the objection that it smooths away the economics rather than the numerics.

C.6 Taste shocks

C.6.1 Extreme Value Type I

Proposition C.12 (EV Type I delivers C^∞)**.** *Let $\Phi = \Phi^\sigma$ with scale $\sigma_\zeta > 0$. Then Φ^σ is real analytic on $\mathbb{R}$, with*

$$(\Phi^\sigma)'(\delta) = P^{\text{adj}} = \frac{1}{1 + e^{-\delta/\sigma_\zeta}} \in (0,1), \qquad (\Phi^\sigma)''(\delta) = \frac{P^{\text{adj}}(1 - P^{\text{adj}})}{\sigma_\zeta} \leq \frac{1}{4\sigma_\zeta}, \tag{248}$$

and $0 < \Phi^\sigma(\delta) - \Phi^{\max}(\delta) = \sigma_\zeta \log\left(1 + e^{-|\delta|}\right) \leq \sigma_\zeta \log 2$. Under Assumptions C.1–C.6 and exact income integration, $v_t \in C^\infty$ for every $t \leq T$, $c_t \in C^\infty$, marginal propensities to consume are continuous, and

$$\left\|v_{mm,t}\right\|_\infty \;\leq\; \max_j \left\|v_{mm,t}^j\right\|_\infty + \frac{\Xi_t^2}{4\sigma_\zeta}. \tag{249}$$

Proof. The derivative formulas are direct, and $P(1-P) \leq \frac{1}{4}$. The bias identity follows from $\sigma_\zeta \log(1 + e^{\delta/\sigma_\zeta}) - \max\{0, \delta\} = \sigma_\zeta \log(1 + e^{-|\delta|})$. For the class: by Assumption C.3 the base case holds at T; if $v_{t+1} \in C^\infty$ then $v_t^j \in C^\infty$ by Lemma C.3(iii), and $v_t = v_t^0 + \Phi^\sigma(\Delta_t)$ is a composition of real analytic with C^∞ maps. Differentiating (236) twice gives $v_{mm,t} = (1-P) v_{mm,t}^0 + P v_{mm,t}^1 + (\Phi^\sigma)''(\Delta_t)(\partial_m \Delta_t)^2$, and bounding each term yields (249). □

Proposition C.12 combined with Lemma C.5 yields the explicit *bias–smoothness trade-off* that the main text asserts only qualitatively: curvature is $O(1/\sigma_\zeta)$ while the per-period model bias is $O(\sigma_\zeta)$. Solving (242) for σ_ζ delivers a threshold σ_ζ^* above which the EGM step is clean, and $\sigma_\zeta^* \log 2$ is then the model distortion required to purchase an envelope-free solve. That number is the honest headline statistic for the exercise and should be reported.

C.6.2 Multiplicative log-normal

Proposition C.13 (Discretised log-normal taste shocks do not smooth)**.** *Consider the retirement specification of Section 2 under Assumption C.8.*

1. *With* exact *integration over* ξ, $\mathbb{E}_\xi\big[\max\{V_{t+1}^{\mathrm{Ret}}, \xi V_{t+1}^{W}\}\big] \in C^\infty$ *by Lemma C.4(i).*
2. *With* K*-node discretisation, the object actually computed,*

$$\mathcal{V}_t^1(A) = \sum_{j=1}^{J}\sum_{k=1}^{K} \omega_j \mu_k \max\left\{V_{t+1}^{\mathrm{Ret}}(M^j),\, \xi_k V_{t+1}^{W}(M^j)\right\}, \tag{250}$$

is Lipschitz but not C^1*: it has* $J \cdot K \cdot |\mathcal{K}_{t+1}|$ *kinks, of individual magnitude* $\omega_j \mu_k \kappa$*, and by Lemma C.4(iii) the total jump is invariant in* K*.*

Proposition C.13 is a negative result for the paper as currently written. Section 2 presents the ability of the log-normal shock to "take up multiple distinct values in the discrete approximation" as an advantage over Extreme Value Type I shocks. For the purpose of smoothing it is a *disadvantage*: the EV Type I assumption integrates out in closed form and yields an exactly C^∞ aggregator with no quadrature at all, whereas a K-point log-normal yields an approximation that is never C^1 and whose total roughness does not decline with K. Only the largest individual kink declines, at rate $O(1/K)$, which is why very large taste-shock node counts are required in practice. Together with the non-invariance flagged in Assumption C.8, this argues for replacing the multiplicative log-normal specification by an additive one.

C.7 Directed search

Lemma C.14 (Closed form of the search aggregator)**.** *For* $\sigma(\varepsilon) = S[\varepsilon + (1-\varepsilon)\log(1-\varepsilon)]$ *as in* (64)*, the conjugate is*

$$\Phi^S(\delta) = \begin{cases} \delta - S\big(1 - e^{-\delta/S}\big), & \delta > 0, \\ 0, & \delta \le 0, \end{cases} \qquad \varepsilon^*(\delta) = \big(\Phi^S\big)'(\delta) = \max\left\{0,\, 1 - e^{-\delta/S}\right\}, \tag{251}$$

and

$$\Phi^S(0) = 0, \quad \big(\Phi^S\big)'(0^{\pm}) = 0, \quad \big(\Phi^S\big)''(0^-) = 0, \quad \big(\Phi^S\big)''(0^+) = \frac{1}{S}, \quad -S < \Phi^S - \Phi^{\max} \le 0. \tag{252}$$

Proof. The interior first-order condition $\delta = \sigma'(\varepsilon) = -S\log(1-\varepsilon)$ gives $1-\varepsilon^* = e^{-\delta/S}$. Substituting, $\sigma(\varepsilon^*) = S\varepsilon^* - \delta e^{-\delta/S}$, hence $\Phi^S = \varepsilon^*\delta - \sigma(\varepsilon^*) = \delta(\varepsilon^* + e^{-\delta/S}) - S\varepsilon^* = \delta - S(1 - e^{-\delta/S})$. The derivative values follow from $(\Phi^S)''(\delta) = S^{-1}e^{-\delta/S}$ on $\delta > 0$; the bias bound from $\Phi^S(\delta) - \delta = -S(1 - e^{-\delta/S}) \in (-S, 0)$ for $\delta > 0$ and equality to zero for $\delta \le 0$. □

Note the contrast in the direction of the bias. The log-sum lies *above* the maximum, by at most $\sigma_\zeta \log 2$, reflecting the option value of taste variation; directed search lies *below* it, by at most S, reflecting the resources absorbed by search. Both biases vanish linearly in the smoothing parameter.

Proposition C.15 (Directed search delivers exactly C^1, invariantly in t). *Let $\Phi = \Phi^S$ with $S > 0$ and suppose Assumptions C.1–C.6 hold with exact income integration. Then for every $t \leq T$:*

1. *$v_t \in C^1$ and, on $\mathcal{K}_t$, $v_t \notin C^2$: the class C^1 is* invariant *under backward iteration;*
2. *$v_{m,t}$ is continuous, hence $\mathcal{E}_{t-1}$ is continuous and m_{t-1} is continuous: the failure mechanism of Proposition C.8 is eliminated, and the EGM step is clean whenever (242) holds;*
3. *$c_{t-1} \in C^0 \setminus C^1$: the conditional policy is continuous but marginal propensities to consume jump, by $\left(\partial_m \Delta_t\right)^2/S$ across $\mathcal{K}_t$;*
4. *$\|(\Phi^S)''\|_\infty = 1/S$, so $\|v_{mm,t}\|_\infty \leq \max_j \|v^j_{mm,t}\|_\infty + \Xi^2_t/S$.*

Proof. (i) Induction. The base case at T is Assumption C.3. Suppose $v_{t+1} \in C^1$. Then $v^j_t \in C^1$ by Lemma C.3(iii) with $k = 0$, and in fact $v^j_t \in C^\infty$ on the complement of the inherited singular set. By Lemma C.14, $\Phi^S \in C^1(\mathbb{R})$, so $v_t = v^0_t + \Phi^S(\Delta_t) \in C^1$; and $\min\{k, 1\} = 1$, so the class does not degrade further. Failure of C^2 on $\mathcal{K}_t$ follows from $(\Phi^S)''(0^+) - (\Phi^S)''(0^-) = 1/S \neq 0$ together with Assumption C.4. (ii) Immediate from (i) and Lemma C.3(i). (iii) $c_{t-1} = (u')^{-1}(\beta \mathcal{R} \mathcal{E}_{t-1})$ inherits the class of $\mathcal{E}_{t-1}$, which is that of $v_{m,t}$, namely $C^0 \setminus C^1$; the size of the MPC jump follows from differentiating $v_{mm,t}$ twice through (236). (iv) From Lemma C.14 and the second derivative of (236). □

Proposition C.15 reconciles the abstract with the conclusion. Part (ii) is the abstract's claim: search alone suffices to eliminate the suboptimal points, because what matters for the EGM step is continuity of the marginal value function, and search delivers it. Part (iii) is the conclusion's claim: search alone does *not* deliver uniform smoothness, because differentiability of the marginal value function is a strictly stronger requirement, and search does not deliver that. These are two distinct properties and the paper should say so explicitly; as currently drafted, the abstract and the conclusion appear to contradict each other.

Proposition C.16 (One-sidedness). *$(\Phi^S)'' \equiv 0$ on $\{\delta < 0\}$ and $\varepsilon^* \equiv 0$ on $\{\delta \leq 0\}$. Hence under directed search: (i) the measure of agents who adjust while the adjust branch is strictly dominated is exactly zero; and (ii) there is exactly zero smoothing on the keep side of $\mathcal{K}_t$, the second derivative of v_t jumping from 0 to $\left(\partial_m \Delta_t\right)^2/S$ across the threshold.*

Proposition C.16 is simultaneously the paper's best economic argument and the formal explanation of its main numerical difficulty. Part (i) is the property that distinguishes search from logit: under logit, $P^{\text{adj}}(0) = \frac{1}{2}$ and probability mass is placed on the dominated branch arbitrarily far from indifference, which is the mechanism behind the finding that a large share of adjusters adjust for taste rather than fundamental reasons. Part (ii) is the exact reason for the asymmetric smoothing visible in Fig. 3: the asymmetry is a property of the functional form, and *no* choice of S removes it. The paper attributes the asymmetry to an economic mechanism (agents with low resources do not search); that account is correct but incomplete, since the deeper cause is the truncation at $\varepsilon^* = 0$.

Proposition C.17 (The trade-off is not a theorem). *Within the family of Lemma C.6, one-sidedness requires only $\sigma'(0) = 0$ while $\Phi \in C^2$ requires $\sigma''(0^+) = +\infty$; the two conditions are compatible. In particular, for the power cost*

$$\sigma_\gamma(\varepsilon) \;=\; A \frac{\gamma}{\gamma + 1} \varepsilon^{(\gamma+1)/\gamma}, \qquad \gamma \in \mathbb{N},\ \gamma \geq 1,\ A > 0, \tag{253}$$

one has $\sigma'_\gamma(\varepsilon) = A\varepsilon^{1/\gamma}$ *with* $\sigma'_\gamma(0) = 0$ *for every* γ*, and* $\sigma''_\gamma(0^+) = +\infty$ *for every* $\gamma > 1$*. By Lemma C.6(iv),* $\Phi \in C^\gamma \setminus C^{\gamma+1}$*, so* $v_t \in C^\gamma$*,* $c_{t-1} \in C^{\gamma-1}$*, and marginal propensities to consume are continuous for every* $\gamma \geq 2$*. The log-type cost (64) satisfies* $\sigma'(\varepsilon) \sim S\varepsilon$ *and is therefore the case* $\gamma = 1$*: the* least smooth *member of the one-sided family.*

Proof. Immediate from Lemma C.6(ii)–(iv) and $-S\log(1-\varepsilon) = S\varepsilon + O(\varepsilon^2)$. □

Proposition C.17 is, I think, the most useful result in this appendix. It shows that the apparent conflict between the economic advantage of search (one-sidedness, hence no dominated-branch adjusters, hence an intact Euler-equation interpretation away from the thresholds) and its numerical disadvantage (only C^1, hence jumping marginal propensities, hence the need for taste shocks to obtain uniform smoothness) is not a genuine trade-off. It is an artifact of a cost function whose marginal cost vanishes linearly at zero. Adopting (253) with $\gamma \geq 2$ delivers both properties at once, and removes the need to combine search with taste shocks — which in turn removes the tension between Section 2, where the two are combined, and Section 3.5, where combining them is rejected as double-smoothing. The economic content of γ is the elasticity of the search hazard with respect to the value gap near the threshold, which is in principle estimable from the same duration data used to discipline S.

C.8 Uniformity in the horizon

The results above are one-step statements. Because the paper's claim concerns the accumulation of non-smoothness across backward iterations, bounds uniform in t are required. The following a priori estimate supplies them.

Proposition C.18 (Uniform curvature bound and explicit threshold)**.** *Let* Φ *satisfy* $\|\Phi''\|_\infty \leq C_\Phi < \infty$*, so* $C_\Phi = 1/S$ *under directed search (Lemma C.14) and* $C_\Phi = 1/(4\sigma_\zeta)$ *under EV Type I taste shocks (Proposition C.12). Suppose Assumptions C.1–C.7 hold and* $\Xi \equiv \sup_t \Xi_t < \infty$*. Then*

$$\left\|v_{mm,t}\right\|_\infty \leq \beta\mathcal{R}^2\left\|v_{mm,t+1}\right\|_\infty + C_\Phi\,\Xi^2, \tag{254}$$

$$\sup_{t\leq T}\left\|v_{mm,t}\right\|_\infty \leq \max\left\{\left\|v_{mm,T}\right\|_\infty,\ \frac{C_\Phi\,\Xi^2}{1-\beta\mathcal{R}^2}\right\}. \tag{255}$$

Consequently, by Lemma C.5(ii), the EGM step is clean at every t *whenever*

$$S > S^* \equiv \frac{\beta\mathcal{R}^2\,\Xi^2}{(1-\beta\mathcal{R}^2)\,\inf|u''|}, \qquad \textit{respectively} \qquad \sigma_\zeta > \sigma_\zeta^* = S^*/4. \tag{256}$$

Proof. Differentiating the envelope identity $v^j_{m,t} = u'(c^j_t) = \beta\mathcal{R}\mathcal{E}_t$ in m gives $v^j_{mm,t} = \beta\mathcal{R}^2\,\mathbb{E}_t[v_{mm,t+1}]\cdot \partial a_t/\partial m$, and $\partial a_t/\partial m = 1 - \partial c_t/\partial m \in (0,1)$ by Assumption C.2, so $\|v^j_{mm,t}\|_\infty \leq \beta\mathcal{R}^2\|v_{mm,t+1}\|_\infty$. Differentiating (236) twice, $v_{mm,t} = (1-\Phi')v^0_{mm,t} + \Phi' v^1_{mm,t} + \Phi''(\Delta_t)(\partial_m\Delta_t)^2$ with $\Phi' \in [0,1]$, which gives (254). Iterating the affine recursion under Assumption C.7 and summing the geometric series gives (255). Substituting into (242) and solving for the smoothing parameter gives (256). □

Two honest caveats. First, $\Xi < \infty$ uniformly in t is an assumption, provable only on a compact region of the state space; it should be verified numerically. Second, as noted after Lemma C.5, $\inf|u''|$ vanishes at high wealth under CRRA felicity, so (256) is informative only when stated regionally or in consumption-equivalent units. Making the constants in (256) non-vacuous is the substantive remaining work; the structure of the argument is otherwise complete.

C.9 The retirement model has a nested aggregator

The retirement model of Section 2 does not fit the single-layer template of Section C.2, and the difference changes what can be claimed. Collecting terms in the worker's problem, and using $P_1(\varepsilon) = (1-\varepsilon)p + \varepsilon$ and $P_2(\varepsilon) = (1-\varepsilon)(1-p)$,

$$\beta\big[P_1(\varepsilon)\mathcal{V}_t^1 + P_2(\varepsilon)\mathcal{V}_t^2\big] = \beta\big[p\,\mathcal{V}_t^1 + (1-p)\,\mathcal{V}_t^2\big] + \varepsilon\cdot\beta(1-p)\big[\mathcal{V}_t^1 - \mathcal{V}_t^2\big], \tag{257}$$

so that maximising over effort yields the reduced form

$$\mathcal{V}_t(A_t) = \log c_t - \delta + \beta\big[p\,\mathcal{V}_t^1(A_t) + (1-p)\,\mathcal{V}_t^2(A_t)\big] + \Phi^S_{\sigma_s}\Big(\beta(1-p)\,\Delta_t^{\mathrm{opt}}(A_t)\Big), \tag{258}$$

$$\Delta_t^{\mathrm{opt}}(A_t) \equiv \mathcal{V}_t^1(A_t) - \mathcal{V}_t^2(A_t) = \mathbb{E}_{y,\xi}\Big[\big(V_{t+1}^{\mathrm{Ret}}(M_{t+1}) - \xi_{t+1}V_{t+1}^W(M_{t+1})\big)^+\Big], \tag{259}$$

where $\Phi^S_{\sigma_s}$ is the conjugate of θ from Lemma C.14 with S replaced by σ_s.

Proposition C.19 (Search and taste shocks act on different margins)**.** *In the retirement model:*

1. *Search does* not *smooth the retire-versus-work maximum. That maximum appears inside* (259) *as the positive part $(\cdot)^+$, and is smoothed only by mollification over (y_{t+1}, ξ_{t+1}), i.e. by the income and taste shocks, via Lemma C.4.*

2. *Search smooths a different margin, namely the arrival of the retirement option, by applying $\Phi^S_{\sigma_s}$ to the option-value gap Δ_t^{opt}.*

3. *With exact integration over (y, ξ), $\Delta_t^{\mathrm{opt}} \in C^\infty$ by Lemma C.4(i), and $\Phi^S_{\sigma_s}$ then reduces $\mathcal{V}_t$ to exactly C^1:* search is the binding constraint on smoothness. *With K-node discretisation of ξ the ranking reverses, by Proposition C.13.*

4. *Without taste or income uncertainty ($J = K = 1$), $\Delta_t^{\mathrm{opt}}(A) = \big(V^{\mathrm{Ret}} - V^W\big)^+$ vanishes identically to the left of the retirement threshold, so $\varepsilon^* \equiv 0$ there and $\mathcal{V}_t$ coincides with the no-search solution; to the right, $\Delta_t^{\mathrm{opt}} \sim \kappa\,(m - m^*)^+$ and $\Phi^S_{\sigma_s}$ contributes $\big(\beta(1-p)\kappa\big)^2\big((m - m^*)^+\big)^2/(2\sigma_s)$, a purely one-sided quadratic.*

Proposition C.19(iv) reproduces Fig. 3 analytically: the smoothing is exactly quadratic on one side of the threshold and exactly absent on the other, with a second-derivative jump of $(\beta(1-p)\kappa)^2/\sigma_s$. Part (iii) explains why taste shocks are indispensable in this application but not in the durable one, where by contrast Φ^S *replaces* the keep–adjust maximum directly and Proposition C.15 applies verbatim. The claim in Section 4 that the durable exercise demonstrates "the same level of smoothness" from search therefore requires qualification: in the durable model search substitutes for the hard maximum, whereas in the retirement model it does not.

C.10 The durable model in two dimensions

Three items require attention beyond the one-dimensional case.

Composite state. With $\Delta_t(m,n) = v_t^{\mathrm{adj}}\big(m + (1-\tau)n\big) - v_t^{\mathrm{keep}}(m,n)$, the reduced form (236) is unchanged and Propositions C.7, C.12, C.15, C.16 and C.17 transfer verbatim with ∇ in place of ∂_m, given Assumptions C.4–C.5. The one-dimensionality of v_t^{adj} in ℓ_t should be exploited: the adjuster's conditional value is a function of a single scalar, so all statements about Φ apply along the ℓ-direction.

The adjuster's own non-convexity is a separate matter. Lemma C.5 must be replaced by a Jacobian condition. Along the adjuster's first-order locus, with $c(d,a)$ from $u_c = \mathfrak{v}'_c$ and d solving $u_d - \mathfrak{v}'_c + \mathfrak{v}'_d = 0$, the endogenous liquidity map is $\ell(d,a) = c(d,a) + d + a$, and refinement by FUES is required exactly when $\partial\ell/\partial d$ changes sign. A sufficient condition for a clean adjuster step is negative definiteness of

$$\nabla^2_{(c,d)} u \;+\; \beta\, \nabla^2\, \mathbb{E}_t\big[v_{t+1}\big] \tag{260}$$

on the relevant region, which by Proposition C.18 holds once $\|\nabla^2 v_{t+1}\|$ is small relative to $\big|\nabla^2 u\big|$. This is where the claim in the abstract and conclusion must be weakened: smoothing the keep–adjust aggregator bounds $\nabla^2 v_{t+1}$ and therefore makes the adjuster's program locally concave *once S exceeds a threshold*, so "the upper-envelope step becomes unnecessary" is provable, but only as a conditional statement with an explicit threshold. It is not unconditional, and Section 4 already observes that residual discontinuities survive in the adjuster policy.

Concavity is not restored.

Proposition C.20 (No smoother restores concavity)**.** *Let Φ be convex, nondecreasing and non-affine, as is the case for both Φ^σ and Φ^S. Then $v_t = v_t^0 + \Phi(\Delta_t)$ is not concave in general even when v_t^0 and v_t^1 are concave, since Φ composed with the difference of two concave functions need not be concave. Concavity is recovered only in the degenerate limits $\sigma_\zeta \to \infty$ or $S \to \infty$, in which the discrete choice disappears.*

Proposition C.20 should be stated explicitly rather than left open, because it is the first question a reader will ask. It implies that the Euler equation remains merely necessary in principle under either smoother, and therefore that the correct target of the analysis is the bounded-curvature criterion (242) rather than sufficiency of the Euler equation. Proving the weaker but operative result is more informative than leaving the concavity question unaddressed.

C.11 Summary of smoothness classes

Table 5 collects the results. Throughout, "policy" refers to the *conditional* policy c_{t-1} of Remark C.11(ii), income integration is exact, and $\gamma \in \mathbb{N}$ indexes the power cost (253).

The last row requires a comment. For $\gamma > 1$ the power cost has $\sigma''(0^+) = +\infty$ and therefore Φ'' unbounded at the origin, so the uniform bound (255) does not apply directly; the correct statement replaces the sup-norm bound on Φ'' by a bound on $\Phi''(\Delta_t)(\partial_m \Delta_t)^2$ over the relevant range of Δ_t, which is finite away from an arbitrarily small neighbourhood of $\mathcal{K}_t$ and integrable across it. Establishing the analogue of Proposition C.18 for $\gamma > 1$ is the one genuinely open technical item in this appendix.

C.12 Numerical tests implied by the theory

The propositions above make sharp, falsifiable predictions that can replace the visual evidence currently offered in Sections 2 and 4. Each is computable with the existing solver.

1. **Conservation of total jump** (Lemma C.4(iii)). The total variation of c'_t should be invariant in the number of income nodes J while the largest individual jump falls at rate $O(1/J)$. This is the correct diagnostic for Section **??**, and it is strictly more informative than the maximum absolute deviation reported in Table **??**, which by Proposition C.10 is dominated by threshold placement.

Table 5: Smoothness class delivered by each aggregator. "Clean EGM" refers to strict monotonicity of the endogenous grid map; "bias" is the per-period sup-norm deviation of the aggregator from the hard maximum.

Aggregator	Φ class	v_t	$v_{m,t}$	c_{t-1}	MPC
Hard max	$C^0 \setminus C^1$	C^0	discontinuous	discontinuous	unbounded
EV Type I, $\sigma_\zeta > 0$	analytic	C^∞	C^∞	C^∞	continuous
Log-normal, K nodes	—	C^0	discontinuous	discontinuous	unbounded
Log-normal, exact	—	C^∞	C^∞	C^∞	continuous
Search, log cost	$C^1 \setminus C^2$	C^1	C^0	C^0	jumps
Search, power cost γ	$C^\gamma \setminus C^{\gamma+1}$	C^γ	$C^{\gamma-1}$	$C^{\gamma-1}$	continuous if $\gamma \geq 2$
	$\|\Phi''\|_\infty$	bias	clean EGM	one-sided	
Hard max	$+\infty$	0	never	yes	
EV Type I, $\sigma_\zeta > 0$	$1/(4\sigma_\zeta)$	$\leq \sigma_\zeta \log 2$	if $\sigma_\zeta > \sigma_\zeta^*$	no	
Search, log cost	$1/S$	$\leq S$	if $S > S^*$	yes	
Search, power cost γ	$+\infty$	$O(A)$	see text	yes	

2. **The monotonicity threshold** (Lemma C.5, Proposition C.18). Report $\min_a m_t'(a)$ as a function of S and of σ_ζ, and locate the value at which it crosses zero. Compare the empirical threshold with the analytic S^* of (256). This replaces "the figures look smooth" with a measured statement, and identifies the friction required to justify dropping the envelope step.

3. **The second-difference signature** (Propositions C.12 and C.15). Across $\mathcal{K}_t$, the second difference of v_t should be continuous under EV Type I taste shocks and should jump by exactly $\left(\partial_m \Delta_t\right)^2 / S$ under directed search. This is a quantitative validation of the theory against the solver, and it distinguishes the two mechanisms unambiguously.

4. **Uniformity in the horizon** (Proposition C.18). Report $\|v_{mm,t}\|_\infty$ and $|\mathcal{N}_t|$ by t for a sweep over S, and check that both stabilise rather than growing. This is the direct test of the accumulation claim, and it requires a horizon considerably longer than the five periods currently used for the durable figures.

5. **The share of dominated-branch adjusters** (Proposition C.16). At the (S, σ_ξ) or (S, σ_ζ) combination that first delivers a clean EGM step, report the fraction of adjustments occurring where the adjust branch is strictly dominated. Under pure search this is exactly zero; under any admixture of taste shocks it is positive. This is the decisive quantitative test of the paper's economic argument against a taste-shock formulation and it is currently absent.

Remark C.21 (Implementation prerequisite)**.** The results of Section C.7 concern the exact conjugate (251). The accompanying implementation selects search intensity by maximisation over a finite grid on $[0, 1]$ rather than by the closed form (201). A grid-selected ε makes the realised aggregator piecewise affine and reintroduces kinks of order the grid spacing, which would contaminate test 3 above in particular. The closed form should be used, or the two should be verified to agree well below the magnitudes being measured.

C.13 Open items

Ranked by difficulty. (i) Uniform boundedness of Ξ_t and of $|\mathcal{N}_t|$ across the backward recursion, and non-vacuous constants in (256): the one-step and invariant-set arguments are complete, but the constants should be calibrated numerically rather than bounded analytically. (ii) The analogue of Proposition C.18 for the power cost (253) with $\gamma > 1$, where Φ'' is unbounded at the origin. (iii) Assumption C.5 in the durable model, which should be verified on the solved model rather than proved from primitives. Everything else — Lemmas C.3–C.14 and Propositions C.7–C.20 — is self-contained and proved above.

C.14 Glossary of mathematical properties

This subsection collects, for reference, precise statements of every property invoked above, together with a pointer to where each is used. Throughout, $U \subseteq \mathbb{R}^n$ is open, $f : U \to \mathbb{R}$, and $|\cdot|$ is the Euclidean norm.

C.14.1 Smoothness and singularity classes

Differentiability class C^k. $f \in C^k(U)$ if all partial derivatives of f of order $\leq k$ exist and are continuous on U. By convention $C^0(U)$ is the continuous functions and $C^\infty(U) = \bigcap_{k\geq 0} C^k(U)$. We write $f \in C^k \setminus C^{k+1}$ to mean that f is *exactly* of class C^k, i.e. $f \in C^k$ but some $k+1$-st partial derivative fails to exist or fails to be continuous somewhere in U. All statements of the form "$v_t \in C^1 \setminus C^2$" (Proposition C.15) are exactness claims in this sense, and are what distinguishes the mechanisms from one another.

Real analyticity. f is real analytic on U if in a neighbourhood of each point of U it equals a convergent power series. Real analyticity is strictly stronger than C^∞ and is the class of the Extreme Value Type I log-sum Φ^σ (Proposition C.12).

Lipschitz continuity. f is Lipschitz on U with modulus L if $|f(x) - f(y)| \leq L|x - y|$ for all $x, y \in U$. A pointwise maximum of finitely many C^1 functions with locally bounded gradients is Lipschitz but generally not C^1; this is the exact sense in which the hard-max value function of Proposition C.7 is "well behaved but not smooth," and it is why the failure is invisible to a plot of v_t itself and visible only in $v_{m,t}$.

One-sided derivatives and the jump bracket. Let $\gamma : (-\epsilon, \epsilon) \to U$ be C^1 with $\gamma(0) = x^*$. The one-sided directional derivatives of f at x^* along γ are

$$\partial^{\pm} f(x^*; \gamma) \equiv \lim_{s \to 0^{\pm}} \frac{d}{ds} f(\gamma(s)), \tag{261}$$

when they exist, and the *jump bracket* is $[\![f']\!](x^*; \gamma) \equiv \partial^+ f - \partial^- f$. The jump bracket of the second derivative, $[\![f'']\!]$, is defined analogously. Equation (244) is a statement about $[\![v_t']\!]$, and Lemma C.14 is a statement about $[\![(\Phi^S)'']\!] = 1/S$.

Kink; convex and concave kinks. f has a *kink* at x^* along γ if it is continuous there and $[\![f']\!] \neq 0$. The kink is *convex* if $[\![f']\!] > 0$ and *concave* if $[\![f']\!] < 0$. The distinction is not cosmetic: only a convex kink contributes a *positive* singular part to f'', and only a positive contribution violates (242). Proposition C.7 shows that a hard maximum always produces a convex kink, which is why the failure of the EGM step under a hard maximum is systematic rather than parameter dependent.

Second-order kink. f has a second-order kink at x^* if $f \in C^1$ near x^* but $[\![f'']\!] \neq 0$. This is the exact defect of directed search under the log-type cost (Proposition C.15(i)): the level and the gradient are continuous, the curvature is not.

Singular set. $\mathcal{N}(f) \equiv \{x \in U : f \text{ is not } C^1 \text{ at } x\}$, and $\mathcal{N}_t \equiv \mathcal{N}(v_t)$. Proposition C.9 is a statement about the growth of $\mathcal{N}_t$ as t recedes from T, and is the formal content of the claim that kinks propagate backwards.

Piecewise class $\mathcal{PC}^k(\mathcal{N})$. For a closed set $\mathcal{N} \subset U$ of Lebesgue measure zero, $f \in \mathcal{PC}^k(\mathcal{N})$ if $f \in C^k(U \setminus \mathcal{N})$ and all partial derivatives of order $\leq k$ admit finite one-sided limits at every point of $\mathcal{N}$. This scale is needed because the hard-max case produces objects that are not in any C^k with $k \geq 1$, so the C^k scale alone cannot distinguish "one kink" from "a thousand kinks" — a distinction on which Proposition C.9 turns.

Distributional derivative and Dirac mass. If f is C^1 away from x^* and f' has jump κ at x^*, then the second derivative of f in the sense of distributions is the measure

$$f'' \;=\; \{f''\} \;+\; \kappa\,\delta_{x^*}, \tag{262}$$

where $\{f''\}$ is the classical part and δ_{x^*} is the Dirac measure at x^*. This is the precise sense in which $\mathbb{E}_t[v_{mm,t+1}]$ is unbounded at a kink and in which condition (242) is "violated infinitely badly" under a hard maximum with quadrature (Proposition C.8). Under exact income integration the Dirac is convolved into the bounded spike of Proposition C.10.

C.14.2 Geometry of the indifference set

Regular value and transversality. Let $\Delta \in C^1(U)$. The value 0 is a *regular value* of Δ if $\nabla\Delta(x) \neq 0$ for every $x \in \Delta^{-1}(0)$. A C^1 curve γ *crosses* $\mathcal{K} = \Delta^{-1}(0)$ *transversally* at $\gamma(0) = x^*$ if $\frac{d}{ds}\Delta(\gamma(s))\big|_{s=0} \neq 0$, i.e. if $\dot\gamma(0)$ is not tangent to $\mathcal{K}$. Assumption C.4 is exactly the requirement that 0 be a regular value of Δ_t for every t.

Codimension-one C^1 submanifold. $\mathcal{K} \subset \mathbb{R}^n$ is a C^1 embedded submanifold of codimension one if each of its points has a neighbourhood in which $\mathcal{K}$ is the graph of a C^1 function of $n-1$ of the coordinates. By the regular value theorem (a corollary of the implicit function theorem), if 0 is a regular value of Δ then $\Delta^{-1}(0)$ is such a submanifold. In the retirement model this reduces to a finite set of isolated threshold points M^*; in the durable model it is the (S, s) boundary curve in (m, n) space, and the theorem is what guarantees that the boundary can be written locally as $m = \bar m(n)$ with $\bar m \in C^1$.

Jump magnitude κ_t and its supremum Ξ_t. $\kappa_t = \inf_{\mathcal{K}_t}|\nabla\Delta_t|$ and $\Xi_t = \sup_{\mathcal{K}_t}|\nabla\Delta_t|$. By Proposition C.7 the marginal value function jumps by exactly $|\nabla\Delta_t|$ across $\mathcal{K}_t$ in the crossing direction, so $\kappa_t > 0$ is what makes the negative result non-vacuous and $\Xi_t < \infty$ is what makes the positive results in Propositions C.12, C.15 and C.18 quantitative. In the retirement model $\kappa_t = |\lambda_t^W - \lambda_t^{\text{Ret}}|$ evaluated at the threshold.

Single crossing. $\Delta_t(\cdot, n)$ satisfies single crossing in m if it changes sign at most once and, where it does, from negative to positive; equivalently, $\{m : \Delta_t(m, n) > 0\}$ is an interval unbounded above. Single crossing at a regular value implies $\partial_m\Delta_t > 0$ at the zero, which is the form used in Assumption C.5. Economically this is what delivers an (S, s) *band*: adjustment occurs on an interval of m rather than on an arbitrary set.

C.14.3 Convexity, duality, and control costs

Strict concavity; concave program. f is strictly concave on a convex U if $f(\lambda x + (1-\lambda)y) > \lambda f(x) + (1-\lambda)f(y)$ for all $x \neq y$ and $\lambda \in (0,1)$. A *concave program* maximises a concave objective over a convex feasible set; in such a program the Karush–Kuhn–Tucker conditions are necessary *and* sufficient, and the maximiser is unique if the objective is strictly concave. The entire difficulty addressed in this paper is that the conditional problems are concave programs only if the continuation value is concave, and by Proposition C.20 no smoother restores that concavity, which is why the operative target is (242) rather than sufficiency of the Euler equation.

Inada conditions. $u'(0^+) = +\infty$ and $u'(+\infty) = 0$ in each argument. These guarantee that the optimal continuous choice is interior, so that the kinks studied here are attributable to the discrete choice rather than to a corner (Assumption C.2).

Translation equivariance. $\Lambda : \mathbb{R}^2 \to \mathbb{R}$ is translation equivariant if $\Lambda(v^0+k, v^1+k) = \Lambda(v^0, v^1)+k$ for every $k \in \mathbb{R}$. All three aggregators in (237) have this property, and it is exactly what permits the one-dimensional reduced form (236). Without it, regularity would have to be analysed in two arguments rather than one and the appendix would not simplify.

Convex conjugate (Legendre–Fenchel transform). For $\sigma : [0,1] \to \mathbb{R}$, the convex conjugate is

$$\sigma^*(\delta) \;\equiv\; \sup_{\varepsilon \in [0,1]} \big\{ \varepsilon\delta - \sigma(\varepsilon) \big\}. \tag{263}$$

σ^* is always convex and lower semicontinuous. If σ is strictly convex and C^2 on $(0,1)$ then σ^* is differentiable with $(\sigma^*)' = (\sigma')^{-1}$ and $(\sigma^*)''(\delta) = 1/\sigma''\big((\sigma')^{-1}(\delta)\big)$ [**?**]. Lemma C.6 is a direct application, and the inversion formula for the second derivative is the entire content of Proposition C.17: *curvature of the smoother is the reciprocal of curvature of the cost.*

Danskin's theorem (envelope theorem for maxima). If $g(\varepsilon;\delta)$ is C^1 in δ, concave in ε, and the maximiser $\varepsilon^*(\delta)$ is unique, then $\frac{d}{d\delta}\max_\varepsilon g = \partial_\delta g\big(\varepsilon^*(\delta);\delta\big)$: the indirect effect through ε^* vanishes. This justifies $\Phi' = \varepsilon^*$ in Lemma C.6(i) and the combined marginal-value formulas (202)–(203) of Section B.3.

Control cost; perturbed utility. A *control cost* is a strictly convex function σ on the simplex of randomisations over the discrete alternatives, subtracted from the expected payoff, so that the agent solves $\max_\varepsilon\{\varepsilon\delta - \sigma(\varepsilon)\}$. The resulting choice probability is $\varepsilon^* = (\sigma')^{-1}(\delta)$. When σ is the negative Shannon entropy the solution is logit, which is the sense in which Extreme Value Type I taste shocks and directed search are members of one family [**???**]. Additive random-utility and additive perturbed-utility representations of stochastic choice are observationally equivalent, which is why the "rationalised versus behavioural" contrast cannot by itself distinguish the two mechanisms; see the discussion after Lemma C.6.

One-sidedness of a smoother. Φ is *one-sided* if $\Phi'(\delta) = 0$ for every $\delta \leq 0$, i.e. zero weight is ever placed on a strictly dominated branch. By Lemma C.6(ii) this holds if and only if the control cost has zero marginal cost at zero intensity, $\sigma'(0) = 0$. Directed search is one-sided; logit is not (Proposition C.16).

Inclusive value. The certainty-equivalent value of a discrete choice before the shocks are realised, $\mathbb{E}\big[\max_j\{v^j + \text{shock}_j\}\big]$. For Extreme Value Type I shocks this is the log-sum (60). Note

$\Phi^\sigma \geq \Phi^{\max}$ (the inclusive value exceeds the maximum, an option value) whereas $\Phi^S \leq \Phi^{\max}$ (search absorbs resources); see the remark following Lemma C.14.

C.14.4 Integration and approximation

Mollification (convolution smoothing). Given a density f and scale $y > 0$, the mollification of g is $G(b) = \int g(b + y\eta) f(\eta)\, d\eta = y^{-1} \int g(z) f\big((z-b)/y\big) dz$. The essential property, and the content of Lemma C.4(i), is that differentiation in b falls on f rather than on g, so G inherits the smoothness of the *smoother* of its two factors. Exact integration of a C^∞ density therefore produces a C^∞ function no matter how badly behaved g is.

Quadrature rule; equiprobable discretisation. A quadrature rule replaces $\int h\, dF$ by $\sum_{j \leq J} \omega_j h(\eta_j)$ with $\omega_j > 0$ and $\sum_j \omega_j = 1$; an equiprobable discretisation sets $\omega_j = 1/J$ with η_j the conditional means of J equiprobable bins. A quadrature rule is a finite convex combination of *evaluations* and therefore, by Lemma C.4(ii), has *no* mollifying effect: it inherits every singularity of g, replicated once per node. The distinction between this entry and the previous one is the rigorous content of the phantom-kink discussion in Section **??**.

Total variation. For h on an interval $[\underline{x}, \bar{x}]$, $\mathrm{TV}(h) = \sup \sum_i |h(x_{i+1}) - h(x_i)|$ over all finite partitions; for piecewise C^1 h this equals $\int |h'|$ plus the sum of the absolute jumps. Lemma C.4(iii) states that TV of the derivative is conserved under refinement of a quadrature rule, which makes $\mathrm{TV}(c_t')$ the correct diagnostic for numerically induced roughness — strictly more informative than the maximum absolute deviation reported in Table **??**.

C.14.5 Properties of the EGM step

Endogenous grid map. $m_t(a) = a + c_t(a)$, where $c_t(a) = (u')^{-1}\big(\beta \mathcal{RE}_t(a)\big)$ is obtained by inverting the Euler equation at an exogenous end-of-period asset value a. The method is well posed precisely when this map can be inverted to recover a, and hence c_t, as a function of m_t.

Clean EGM step; suboptimal point. The EGM step at t is *clean* if $a \mapsto m_t(a)$ is continuous and strictly increasing on the unconstrained region (Definition C.2). A *suboptimal point* is a candidate pair $\big(m_t(a), c_t(a)\big)$ satisfying the Euler equation but not attaining the maximum; such points exist exactly when m_t fails to be injective, since then two or more values of a deliver the same m_t and at most one can be optimal. Lemma C.5 characterises cleanliness by the curvature condition (242), and this — not differentiability — is what the abstract's claim about eliminating suboptimal points should be understood to assert.

Upper envelope refinement. Any procedure that, given the candidate cloud $\{(m_t(a), c_t(a), v_t(a))\}$, discards the dominated candidates and returns the upper envelope of the implied value correspondence; DC-EGM [**?**] and FUES [**?**] are instances. A clean EGM step is exactly the condition under which this refinement is unnecessary.

C.14.6 Dynamic-programming properties

Envelope bootstrap. The observation, formalised in Lemma C.3(ii), that the envelope condition $v^j_{m,t} = u'(c^j_t)$ recovers one order of differentiability relative to the Euler equation, so that conditional value functions *preserve* their smoothness class across a backward step. It follows that only the aggregator can degrade smoothness, and this is what reduces the entire induction to a statement about Φ near the origin.

Invariant smoothness class. A class $\mathcal{C}$ is invariant for the backward operator if $v_{t+1} \in \mathcal{C}$ implies $v_t \in \mathcal{C}$. The three mechanisms have invariant classes C^0 (hard maximum, with a growing singular set), C^1 (directed search, log-type cost), and C^∞ (Extreme Value Type I). Invariance is what upgrades a one-step regularity result into a statement about accumulation over the horizon.

A priori estimate; invariant set. An inequality bounding a solution in terms of the data before the solution is known. Proposition C.18 establishes one by showing that the backward operator maps $\{v : \|v_{mm}\|_\infty \leq B\}$ into itself for $B = C_\Phi \Xi^2/(1 - \beta\mathcal{R}^2)$; the set is then invariant and the bound holds uniformly in t.

Contraction modulus. The factor $\beta\mathcal{R}^2 < 1$ (Assumption C.7) by which curvature is discounted per backward step, arising as $\beta\mathcal{R}$ from discounting the marginal value and a further $\mathcal{R}$ from the chain rule through $m_{t+1} = \mathcal{R}a_t + \theta_{t+1}$. It is what makes the geometric series in (255) converge, and hence what rules out unbounded accumulation of curvature over long horizons.